

Deep Software Stack Optimization for AI-Enabled Embedded Systems

Sep. 28, 2025

Prof. Seongsoo Hong, Namcheol Lee, Geonha Park, and Taehyun Kim

{sshong, nclee, ghpark, thkim}@redwood.snu.ac.kr

Dept. of Electrical and Computer Engineering,
Seoul National University

Seoul National University

RTOS Lab

Special Thanks to

- ❖ Sapphire Stream Technology
 - For sponsoring this tutorial and providing RUBIK Pi 3 boards to attendees



Tutorial Outline

Lecture 1: *Exercise Overview and Setup* (1 h 30 min)

Lecturer Seongsoo Hong

Topics Motivating Example

Development Environment Setup

Coffee Break (30 min)

Lecture 2: *From Inference Driver to Inference Runtime* (1 h 30 min)

Lecturer Seongsoo Hong and Namcheol Lee

Topics Step-by-Step Inference Driver Walkthrough

Internals of LiteRT

Lunch Break (1 h)

Lecture 3: *Model Slicer* (1 h 30 min)

Lecturer Seongsoo Hong

Topic Model Slicer: Slicing and Conversion Tool for LiteRT

Coffee Break (30 min)

Lecture 4: *Throughput Enhancement on Heterogeneous Accelerators* (1 h 30 min)

Lecturer Namcheol Lee

Topic Implementing a Pipelined Inference Driver for Heterogeneous Processors

Contents

- I. Our Exercise
- II. Tutorial Execution Environment
- III. Hands-On Exercise: Connect to RUBIK Pi

On-Device AI

❖ Refers to “*executing AI models directly on edge devices*”

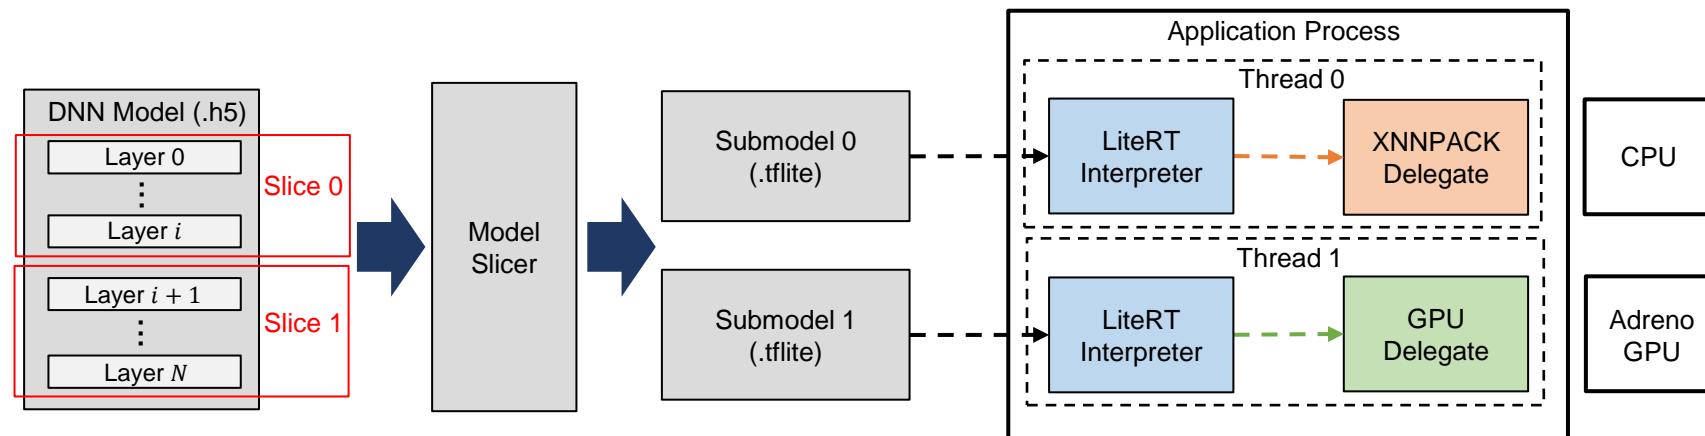
- Benefits
 - Enables *real-time*, *privacy-preserving*, and *offline inference*
- Challenge
 - Operates under limited resources available on edge devices
- Solutions
 - Static model optimizations
 - Quantization, pruning, knowledge distillation, low-rank adaptation, weight streaming, and sparse inference
 - Dynamic optimizations
 - Pipelined inference, operator selection/scheduling, and operator fusion via capabilities
 - Dynamic resource management
 - Accelerator allocation, memory management, QoS regulation, etc.

Our Exercise Problem

- ❖ Lossless stream processing in an edge device
 - Real-world relevance
 - Surveillance cameras
 - Automotive perception applications
 - A transient surge in input data stream may result in data loss
 - When input stream's data rate exceeds device inference throughput
 - Need to scale up inference throughput dynamically

DNN Pipelining (1)

- ❖ Scale up inference throughput via pipelining on multiple accelerators
 1. Partition DNN model into submodels
 2. Create worker threads that perform inference for submodels
 3. Allocate the worker threads to available, heterogeneous accelerators

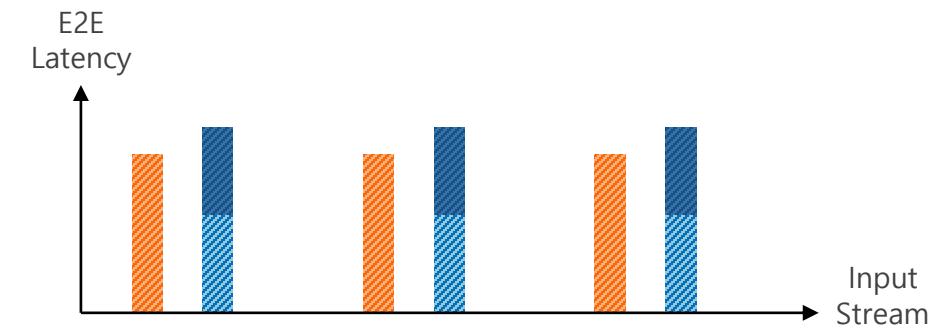
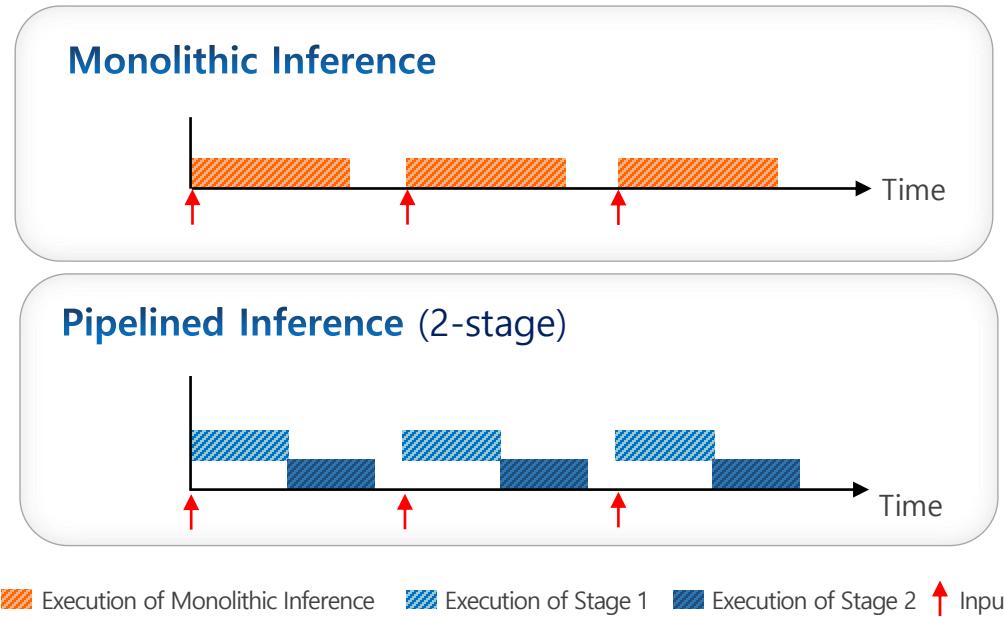


DNN Pipelining (2)

- ❖ What is the problem? And why is pipelining beneficial?
- ❖ Case analysis
 - 1. Normal case: *If input stream's data rate \leq device's inference throughput*
 - *Pipelining can be detrimental*
 - May increase end-to-end (E2E) latency for DNN inference
 - Some stages may run on relatively slower processor (e.g., CPU)
 - Communication delays could be added
 - 2. Exceptional case: *If input stream's data rate $>$ device's inference throughput*
 - *Pipelining is beneficial*
 - The legacy system may incur unbounded E2E latency due to back logs
 - Pipelining can increase throughput, enabling bounded E2E latency

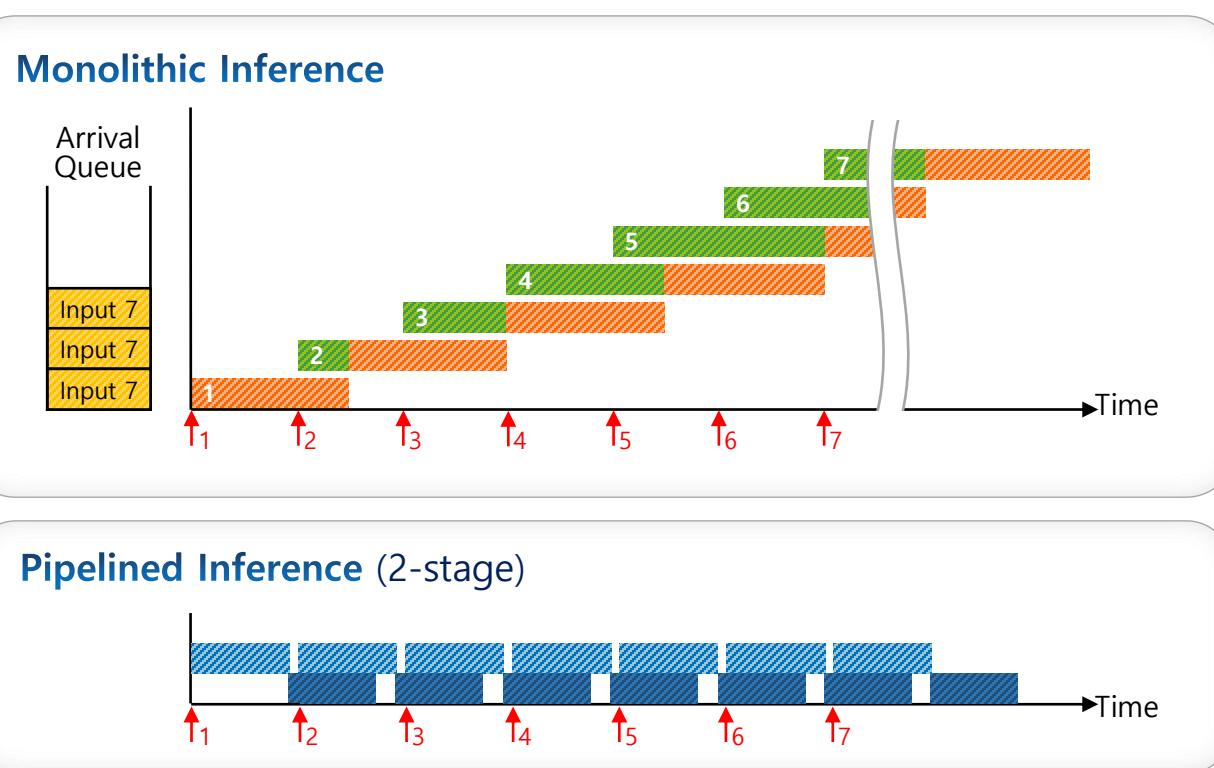
DNN Pipelining (3)

1. Normal case: *If input stream's data rate \leq device's inference throughput*



DNN Pipelining (4)

2. If input stream's data rate > device's inference throughput

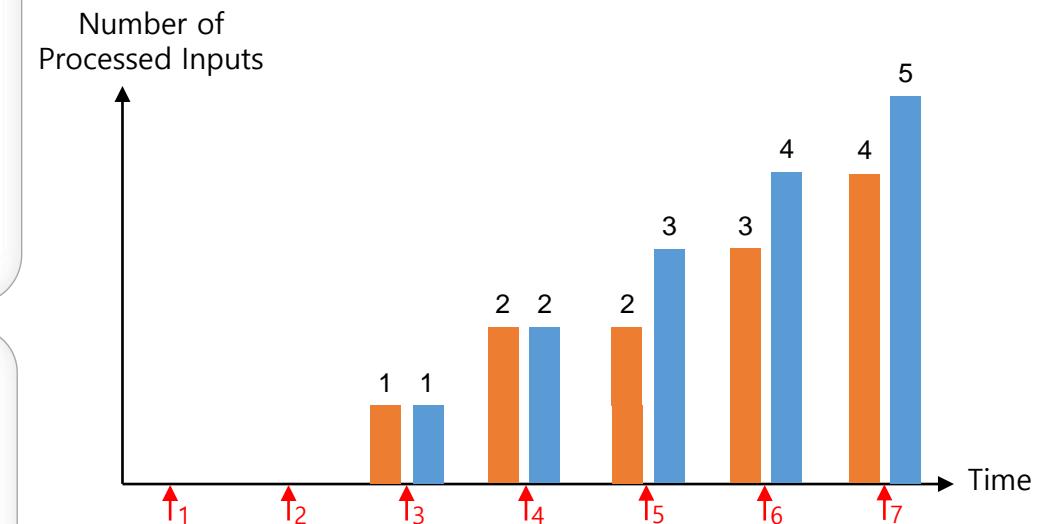


Execution of Monolithic Inference

Execution of Stage 1
Input i

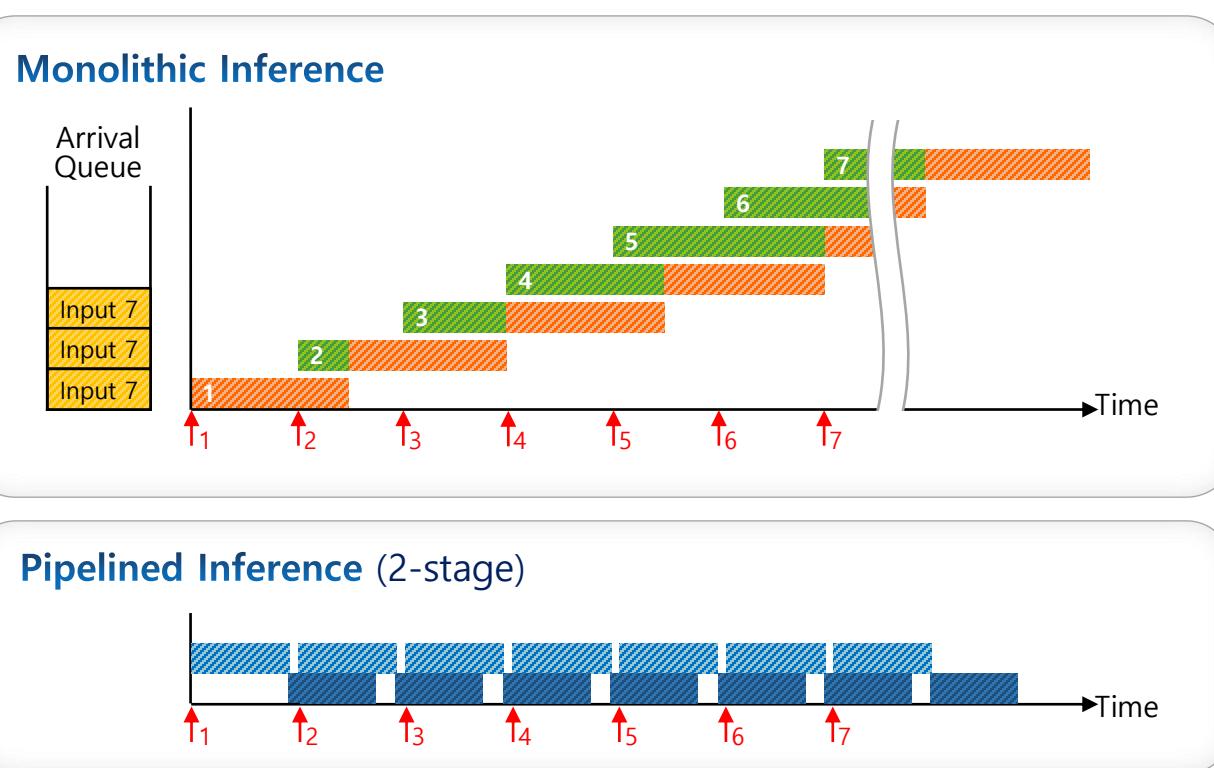
Queuing Delay

Execution of Stage 2



DNN Pipelining (5)

2. If input stream's data rate > device's inference throughput



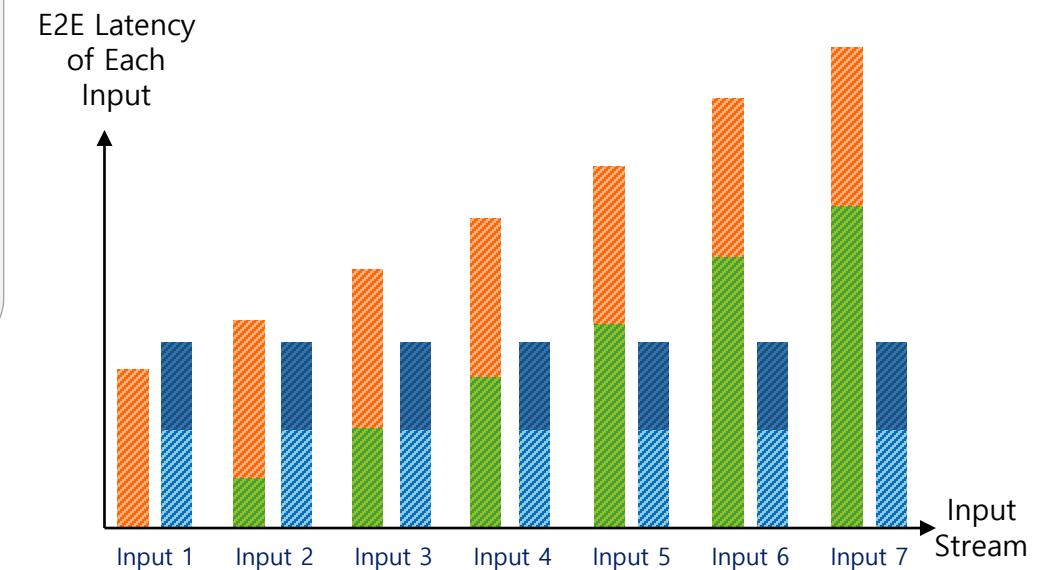
Execution of Monolithic Inference

Queuing Delay

Execution of Stage 1

Execution of Stage 2

Input i



Target DNN Model: ResNet50 (1)

- ❖ We will use ResNet50 FP32 throughout the tutorial which is
 - A 50-layer deep convolutional neural network for image classification
 - Introduced in 2015 by He et al.¹ with the concept of skip connections
- ❖ Why ResNet50?
 - Commonly used for frame-by-frame classification
 - Moderate computational demand
 - Widely used as a benchmark model in inference optimization

¹He, Kaiming, et al. "Deep residual learning for image recognition." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016.

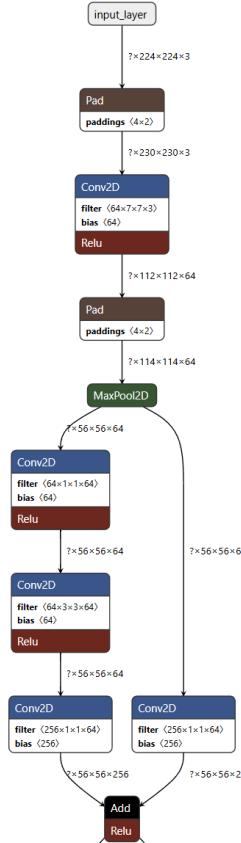
Target DNN Model: ResNet50 (2)

❖ Key details

- Input
 - 224x224 RGB image (224x224x3)
- Output
 - 1x1000 tensor

❖ Observe model structure using Netron²

- `netron ./models/resnet50.tflite`



²<https://netron.app/>

Contents

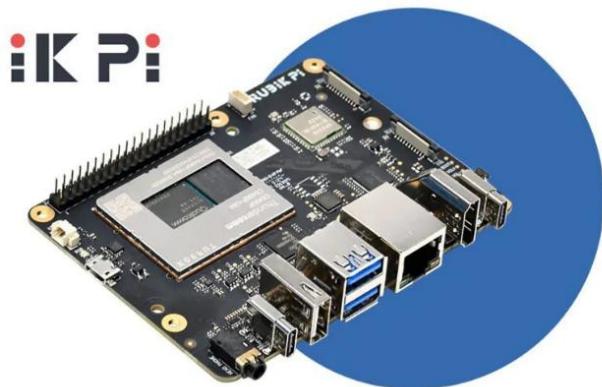
- I. Our Exercise
- II. Tutorial Execution Environment**
- III. Hands-On : Connect to RUBIK Pi

Hardware: RUBIK Pi 3 (1)

❖ Overview

- A Single Board Computer (SBC) built on Qualcomm AI platforms for developers
 - Based on QCS6490 SoC
- Supports multiple operating systems such as
 - Android 13, Qualcomm Linux, Debian 13
- Supports SDKs with Qualcomm AI stack
- Designed and manufactured by Thundercomm

RUBIK Pi



Hardware: RUBIK Pi 3 (2)

❖ Specification

- RUBIK Pi 3 with QCS6490 SoC
 - CPU: 8 x Kryo 670 cores
 - GPU: Adreno 643L
 - NPU: Hexagon 770
 - RAM: 8 GB LPDDR4X
 - Storage: 128 GB UFS 2.2

RUBIK Pi



Hardware: RUBIK Pi 3 (3)

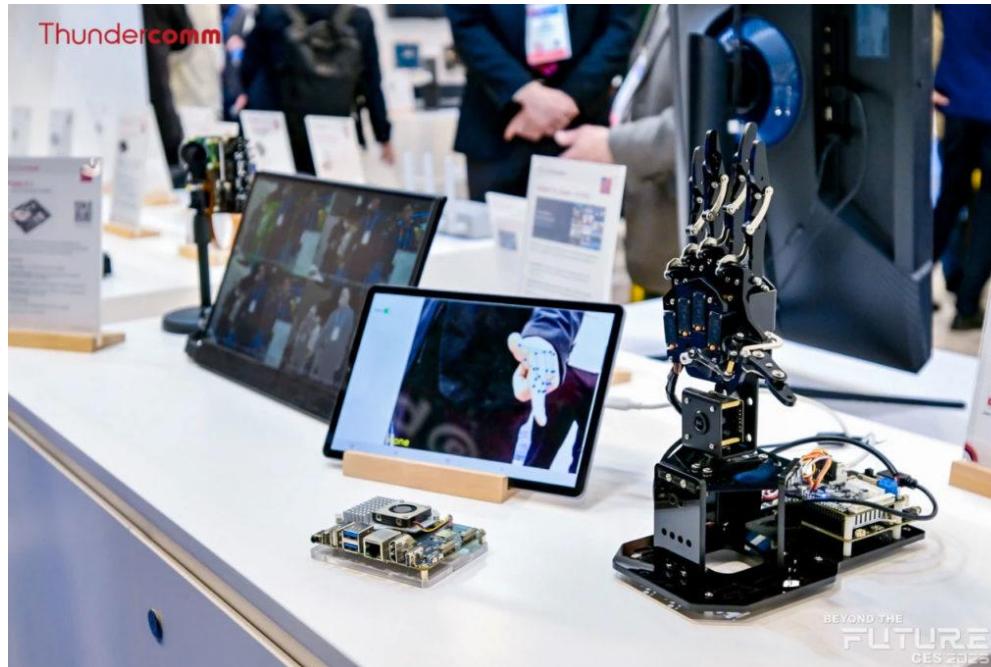
❖ Comparison with Raspberry Pi 5

- RUBIK Pi offers superior AI performance and faster storage, making it ideal for edge AI education and prototyping

	RUBIK Pi 3	Raspberry Pi 5
CPU	8 x Kryo 670 (1 x Cortex-A78 @ 2.7 GHz, 3 x Cortex-A78 @ 2.4 GHz, 4 x Cortex-A55 @ 1.9 GHz)	4 x Cortex-A76 @ 2.4GHz
GPU	Adreno 643L @ Up to 812 MHz	VideoCore VII @ Up to 800MHz
NPU	Hexagon 770 (12 TOPS)	N/A
RAM	8GB	Up to 16GB
ROM	128GB UFS 2.2	N/A (Need to purchase external SD card separately)
Price	\$179	\$120 (16 GB Model)

Hardware: RUBIK Pi 3 (4)

- ❖ Ideal for Hands-On learning of embedded AI software stack
 - RUBIK Pi 3 with QCS6490 SoC enables realistic, practice-oriented exercises aligned with real-world industrial applications



Source: Thundercomm (<https://www.thundercomm.com/rubik-pi-3-embedded-world-2025/>)



Seoul National University

Taxonomy

❖ Inference driver

- Program that calls the inference runtime's APIs to drive inference

❖ Inference runtime

- A software environment that executes DNN inference
 - Such as LiteRT (Google), QNN (Qualcomm), TensorRT (Nvidia)

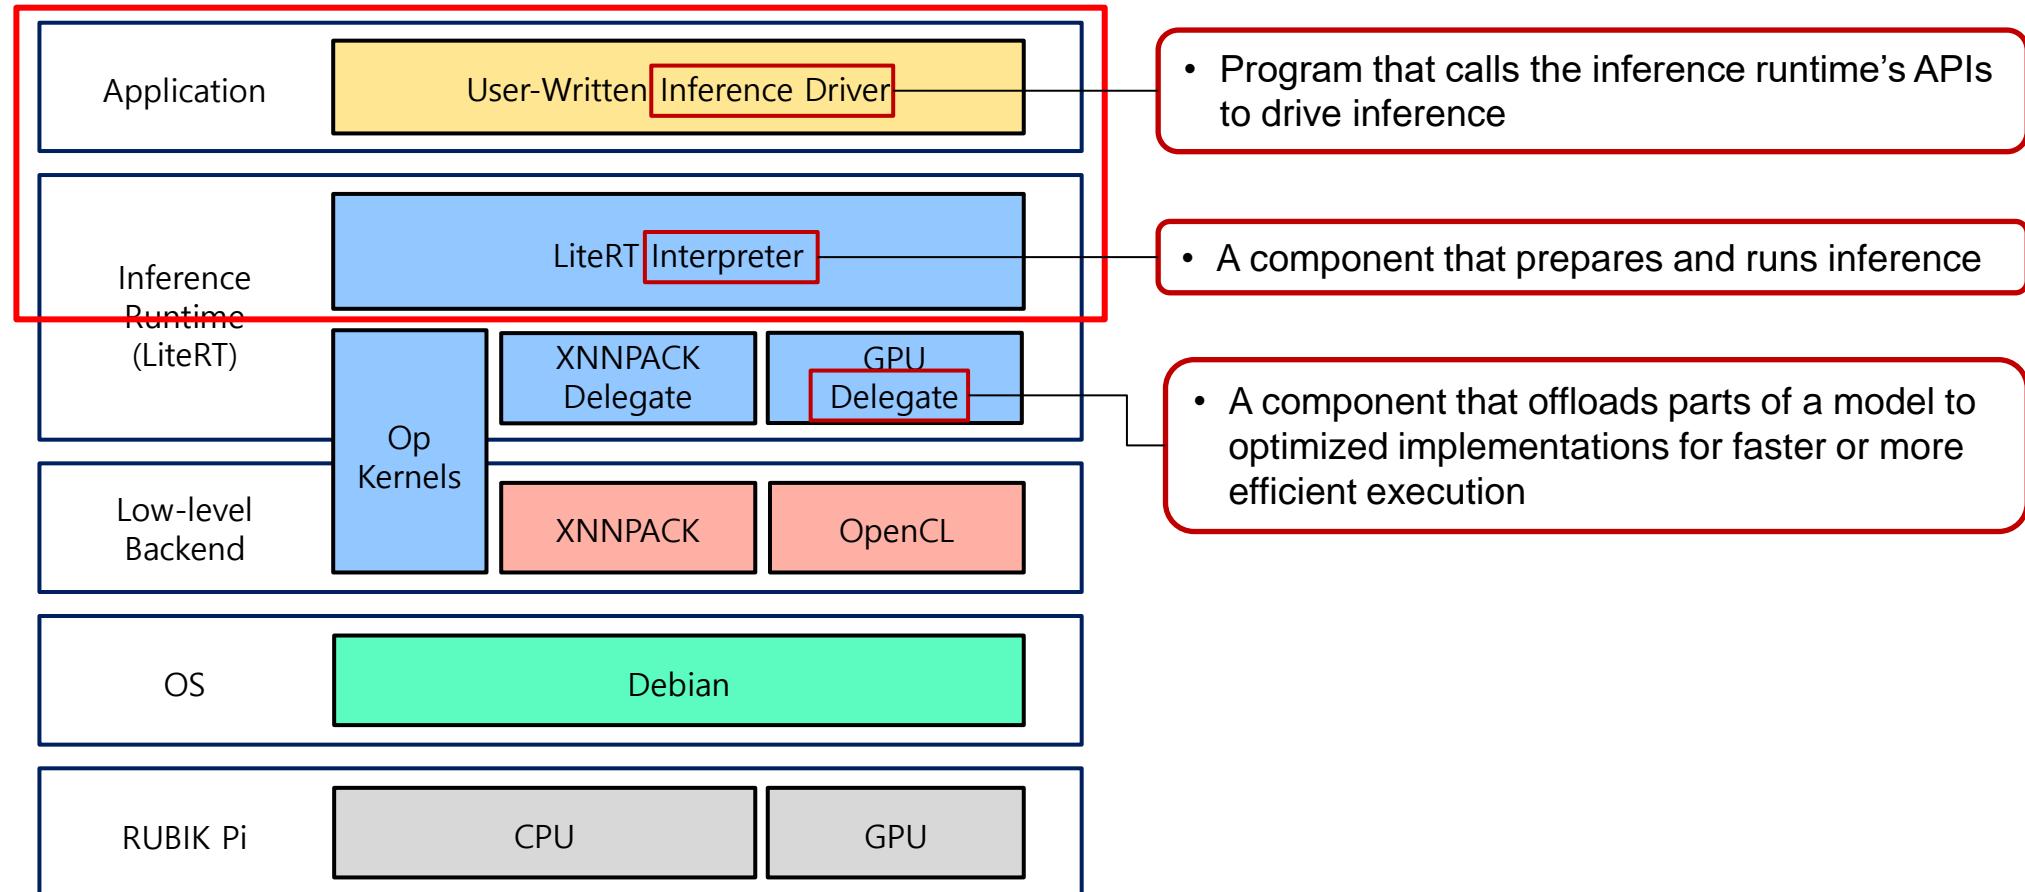
❖ LiteRT interpreter

- A component that prepares and runs inference

❖ Delegate

- A component that offloads parts of a model to optimized implementations for faster or more efficient execution

Software Stack on RUBIK Pi 3



Contents

- I. Our Exercise
- II. Tutorial Execution Environment
- III. Hands-On Exercise: Connect to RUBIK Pi**

Connect to RUBIK Pi

❖ Objective

- Establish an SSH connection from your laptop to RUBIK Pi

❖ Do

- Follow the instructions in this section

❖ Verify

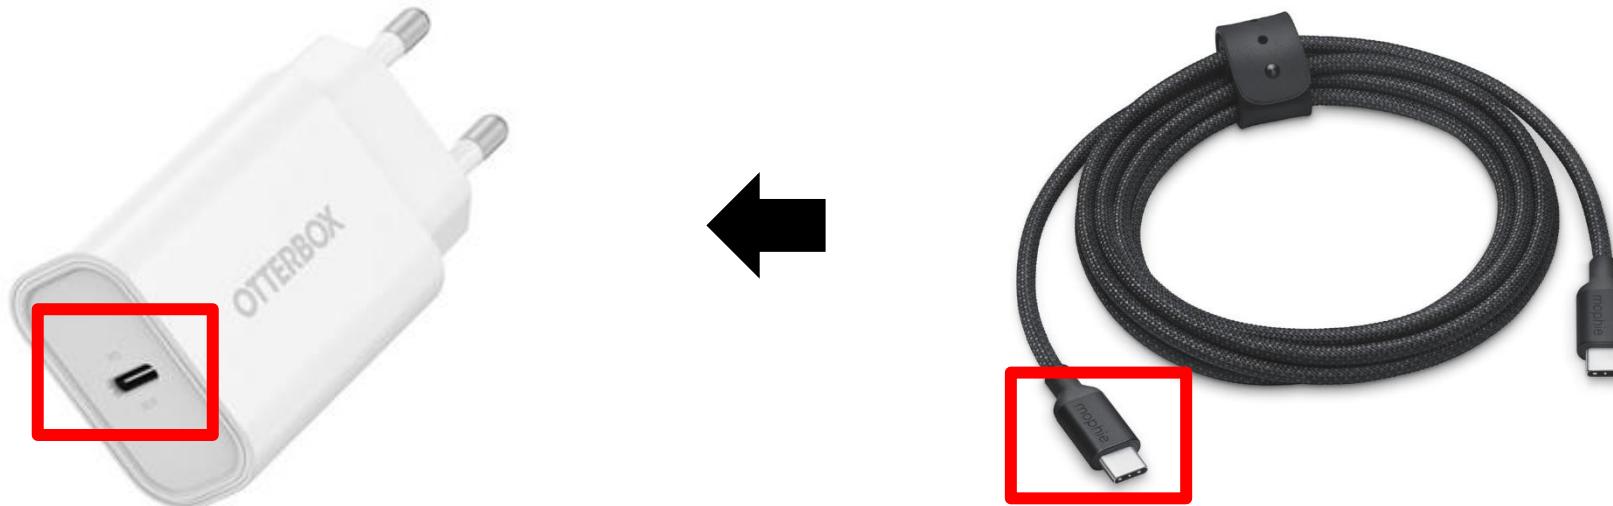
- Your VS Code should show a SSH session connected to RUBIK Pi

❖ Time

- 60 minutes

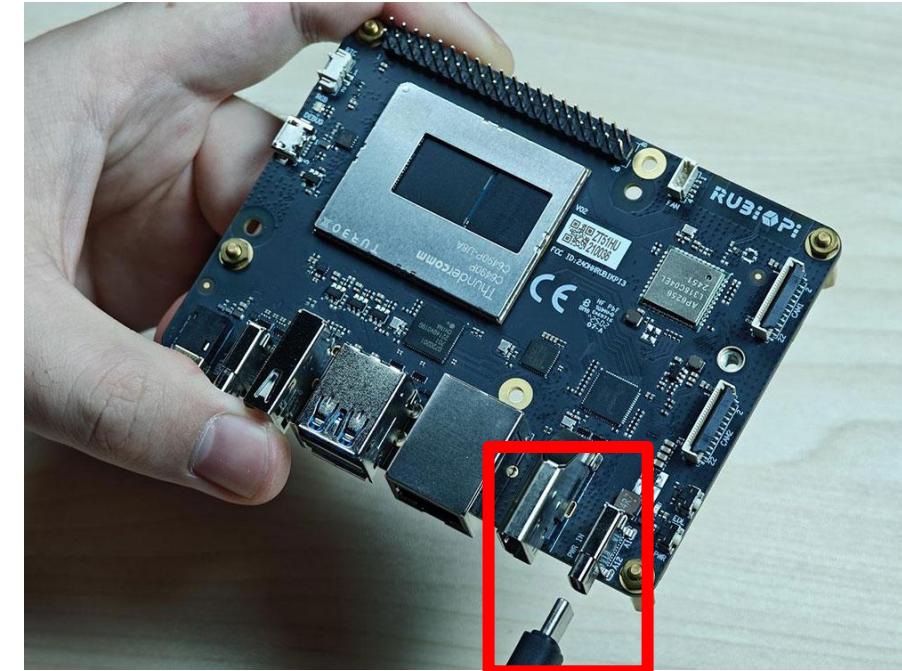
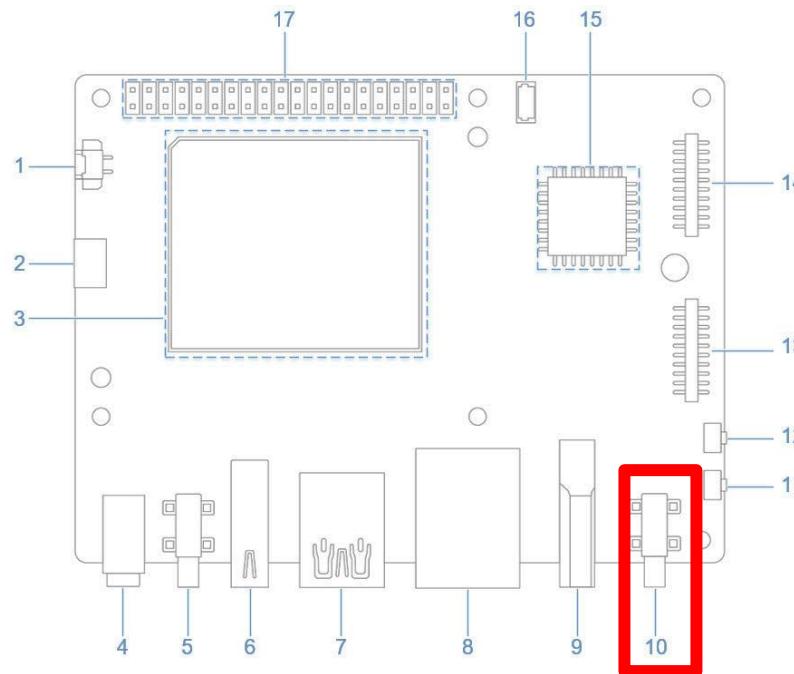
Supply Power (1)

- ❖ Connect power cable to the external power supply, and make sure it is switched on



Supply Power (2)

- ❖ Connect Type-C power cable to port 10 on RUBIK Pi



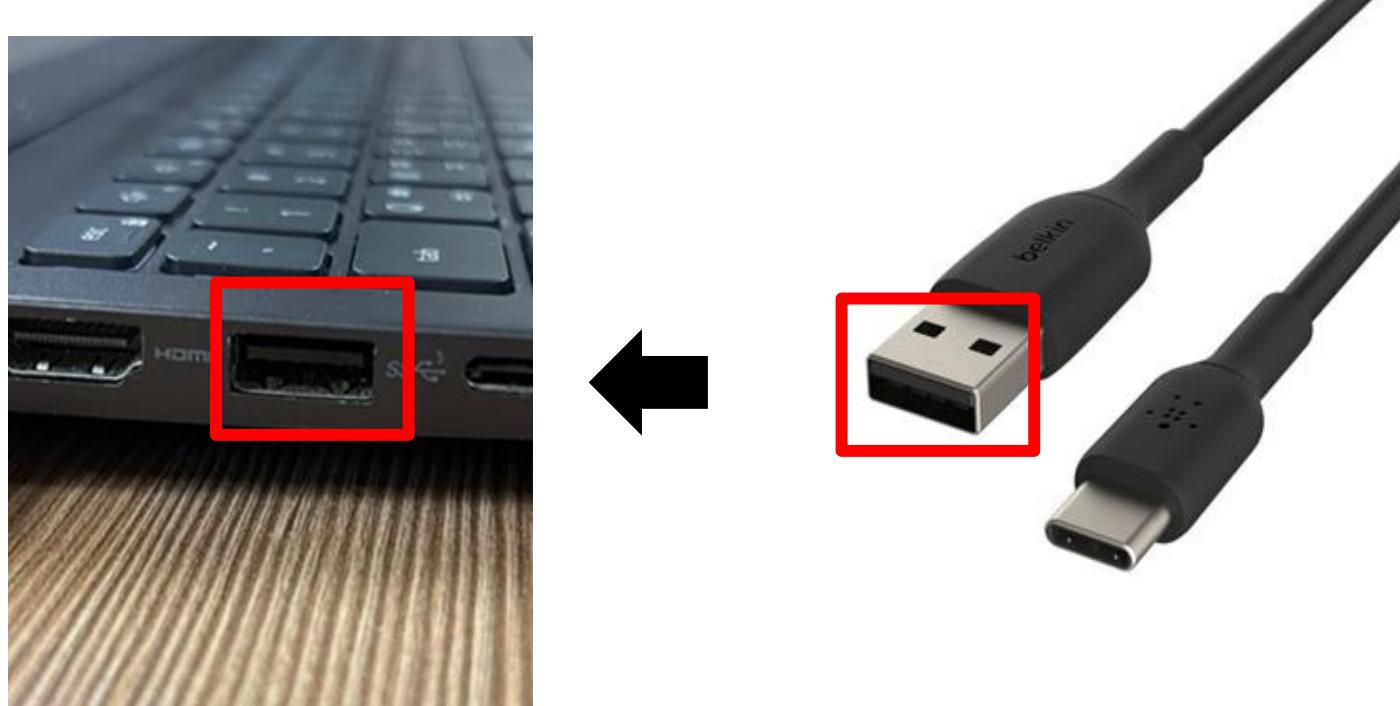
Connect RUBIK Pi to Host PC

- ❖ We will access the RUBIK Pi using SSH with Visual Studio Code for development ease
 - To do so
 1. Connect the device to the host PC with a USB-A to USB-C cable
 2. Establish the device's Wi-Fi connection using ADB (Android Debug Bridge)
 3. Make an SSH connection to the device

III. Hands-On Exercise: Connect to RUBIK Pi

1. Connect Device to Host PC with USB-A to USB-C Cable (1)

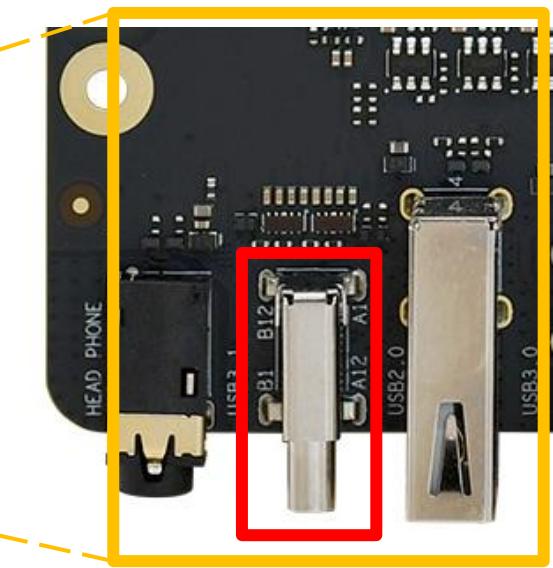
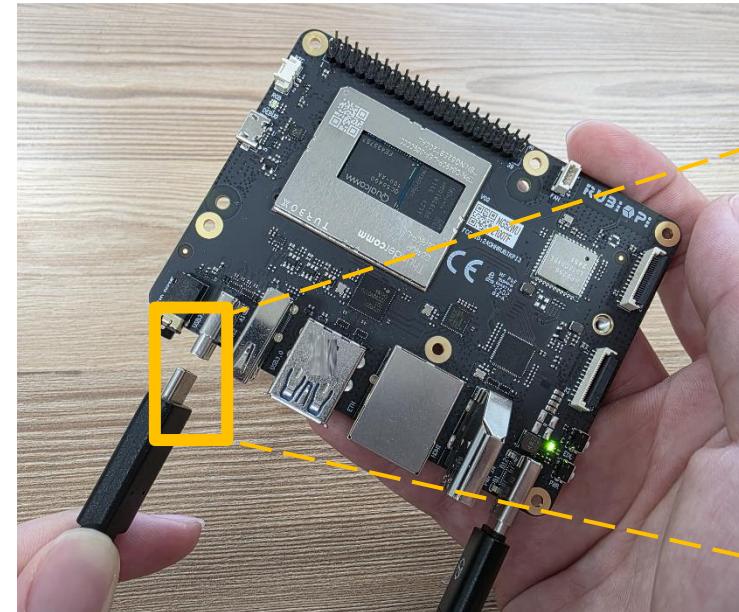
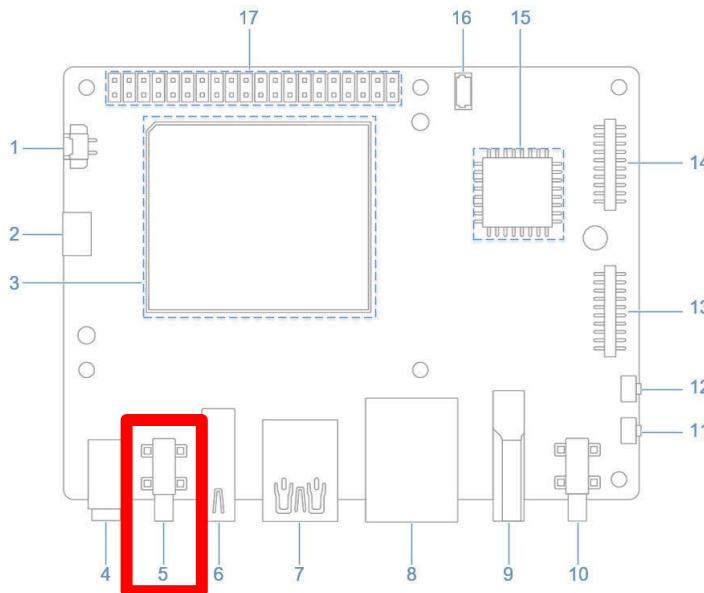
- ❖ Connect the USB-A end of the cable to host PC



III. Hands-On Exercise: Connect to RUBIK Pi

1. Connect Device to Host PC with USB-A to USB-C Cable (2)

- ❖ Connect the USB-C end of the cable to port 5



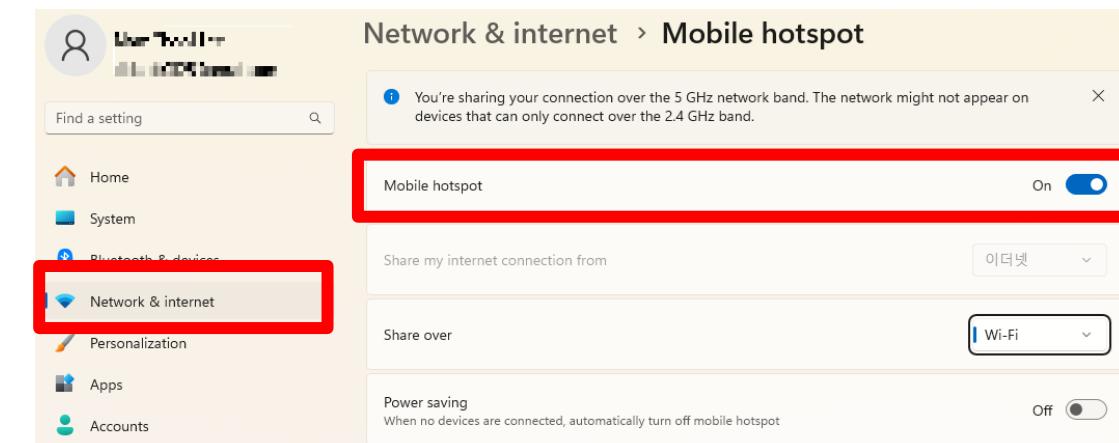
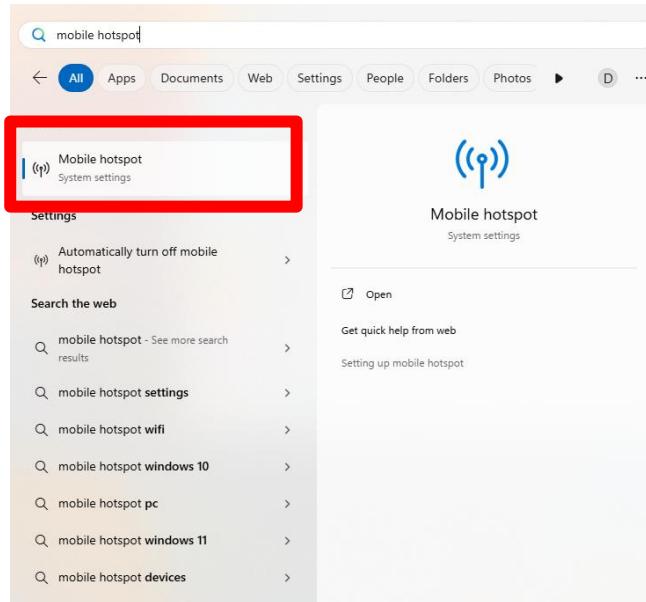
1. Connect Device to Host PC with USB-A to USB-C Cable (3)

- ❖ Check the connection using ADB
 - It should show one device as below
 - The identifier of the device varies across devices

```
C:\Users\Lee>adb devices
List of devices attached
f10e86f9        device
```

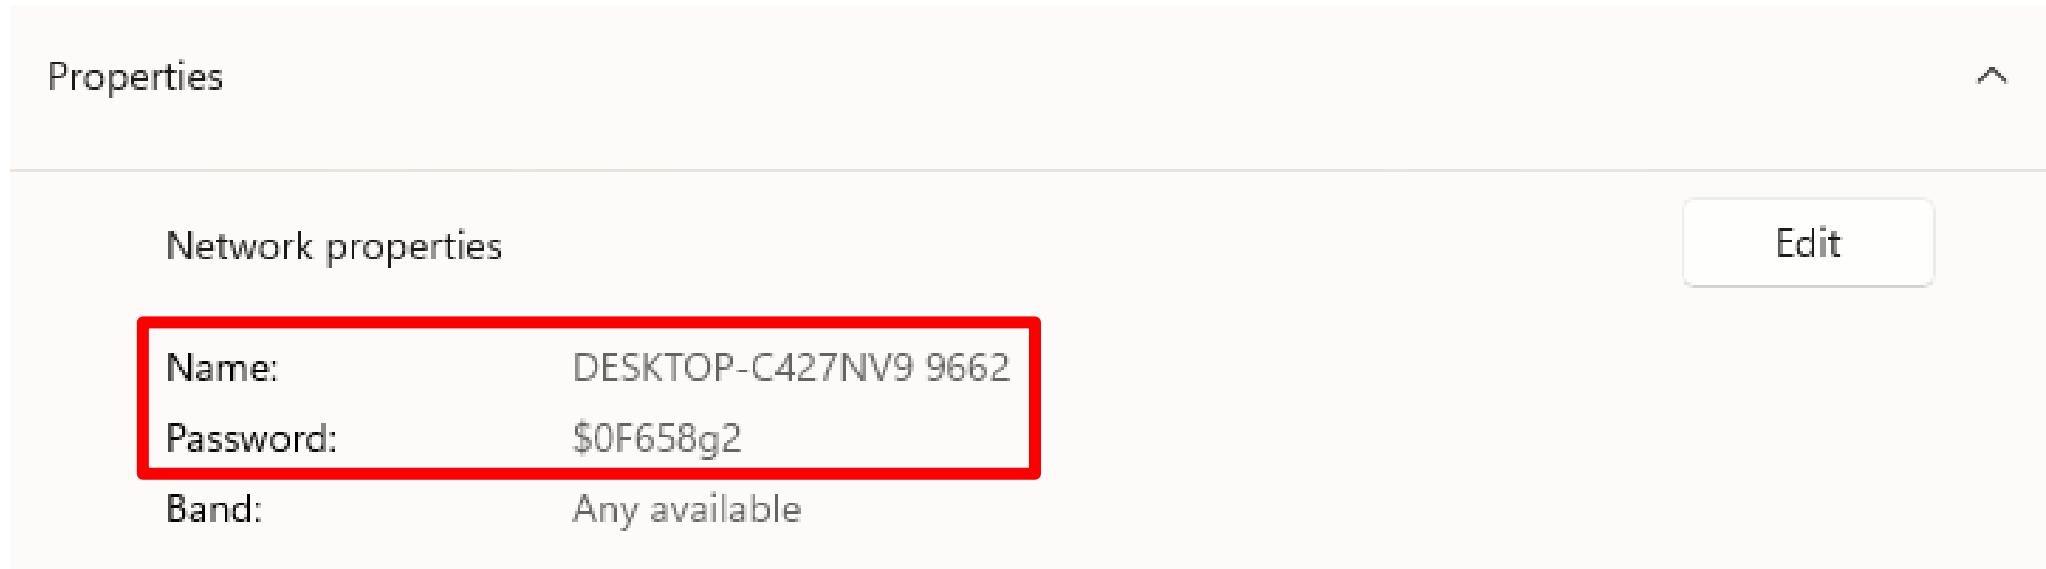
2. Establish Wi-Fi Connection (1)

- ❖ Enable hotspot on host PC
 - Go to **Settings > Network & Internet**



2. Establish Wi-Fi Connection (2)

- ❖ Enable hotspot on host PC (cont'd)
 - Check the name and password of the activated network



2. Establish Wi-Fi Connection (3)

❖ Access the device in a command line tool with the following commands

1. Get root access
 - `adb root`
2. Enter adb shell
 - `adb shell`
3. Switch to root user
 - `su`

```
C:\Users\Lee>adb root  
adb is already running as root  
  
C:\Users\Lee>adb shell  
# su  
root@RUBIKPi:/# |
```

2. Establish Wi-Fi Connection (4)

- ❖ Connect RUBIK Pi to the mobile hotspot
 1. Check the available Wi-Fis using network manager CLI (`nmcli`) of Linux
 - `nmcli device wifi list`

2. Establish Wi-Fi Connection (5)

❖ Connect RUBIK Pi to the mobile hotspot (cont'd)

2. Connect to the mobile hotspot using nmcli

- `nmcli device wifi connect "SSID" --ask`

```
root@RUBIKPi:~# nmcli device wifi connect "DESKTOP-C427NV9 9662" --ask
Push of the WPS button on the router or a password is required to access the wireless
network 'DESKTOP-C427NV9 9662'.
Password (802-11-wireless-security.psk): *****
Device 'wlan0' successfully activated with '0ea65cff-dd17-4aa1-b934-2dedbe33d58f'.
```

3. Check the connection

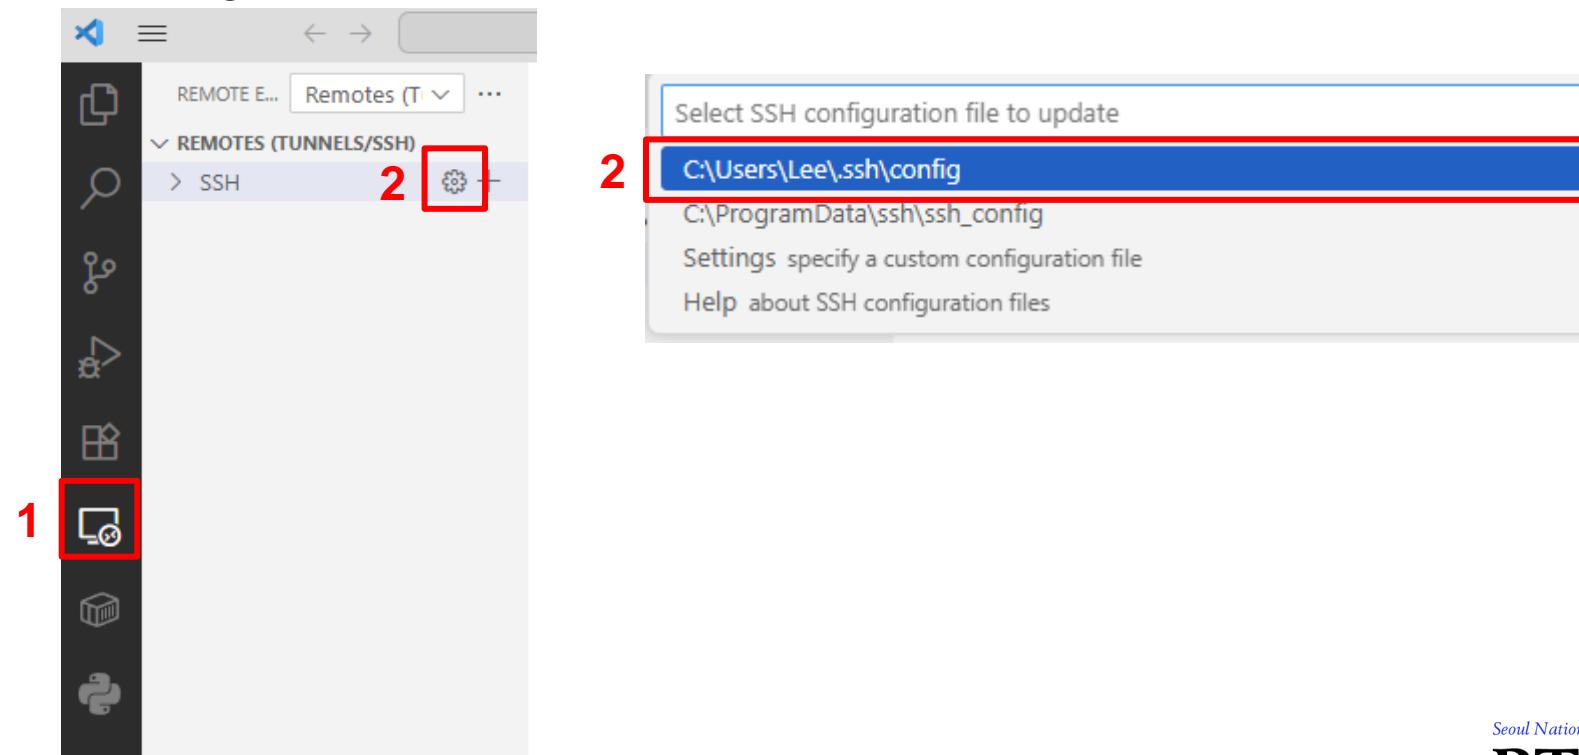
- `ifconfig wlan0`

```
root@RUBIKPi:~# ifconfig wlan0
wlan0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.137.35  netmask 255.255.255.0  broadcast 192.168.137.255
        ether 9c:b8:b4:9f:4f:5c  txqueuelen 1000  (Ethernet)
        RX packets 428199  bytes 627300827 (598.2 MiB)
        RX errors 0  dropped 17738  overruns 0  frame 0
        TX packets 89787  bytes 11357654 (10.8 MiB)
        TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0
```

3. Make SSH Connection (1)

❖ Create SSH configuration in VS Code

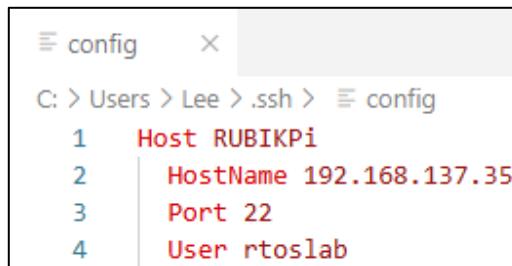
1. Open remote explorer in VS Code in the left side bar
2. Open SSH config file



3. Make SSH Connection (2)

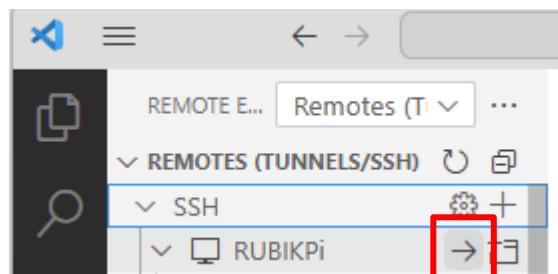
- ❖ Create SSH configuration in VS Code (cont'd)

3. Add the configuration for RUBIK Pi



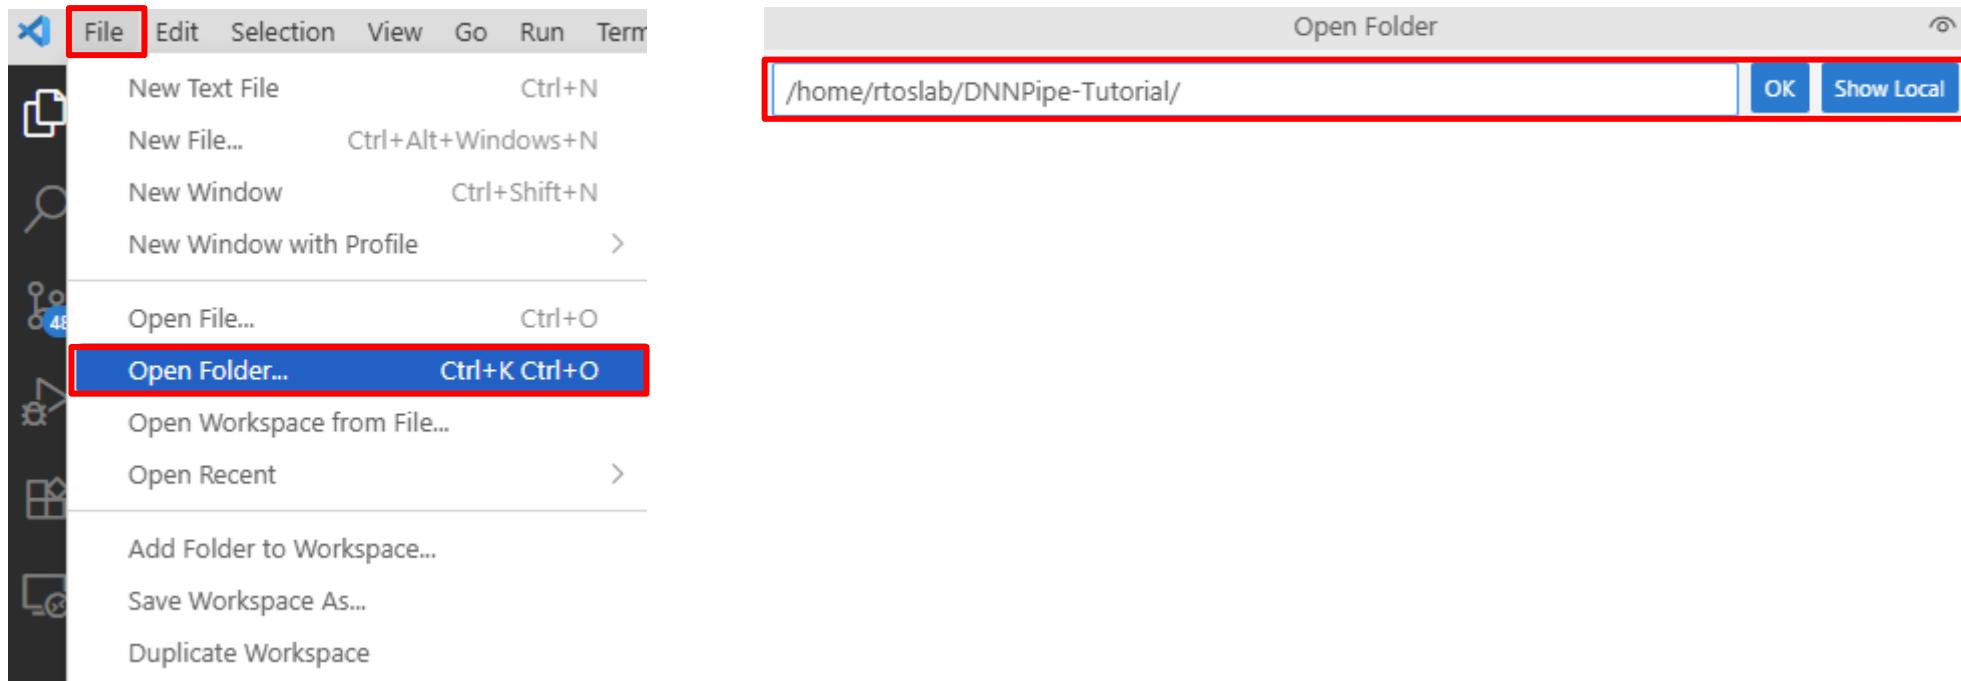
```
config
C: > Users > Lee > .ssh > config
1 Host RUBIKPi
2 HostName 192.168.137.35
3 Port 22
4 User rtoslab
```

4. Connect to the RUBIK Pi with both user name and password being “rtoslab”



3. Make SSH Connection (3)

- ❖ Open `~/DNNPipe-Tutorial` directory



Tutorial Outline

Lecture 1: *Exercise Overview and Setup* (1 h 30 min)

Lecturer Seongsoo Hong

Topics Motivating Example

Development Environment Setup

Coffee Break (30 min)

Lecture 2: *From Inference Driver to Inference Runtime* (1 h 30 min)

Lecturer Seongsoo Hong and Namcheol Lee

Topics Step-by-Step Inference Driver Walkthrough

Internals of LiteRT

Lunch Break (1 h)

Lecture 3: *Model Slicer* (1 h 30 min)

Lecturer Seongsoo Hong

Topic Model Slicer: Slicing and Conversion Tool for LiteRT

Coffee Break (30 min)

Lecture 4: *Throughput Enhancement on Heterogeneous Accelerators* (1 h 30 min)

Lecturer Namcheol Lee

Topic Implementing a Pipelined Inference Driver for Heterogeneous Processors

Contents

-
- I. Directory Structure
 - II. Step-by-Step Inference Driver Walkthrough
 - III. Hands-On Exercise: Build and Run Inference Driver
 - IV. Two Key Internal Entities of LiteRT
 - V. Instrumentation Harness

Directory Structure

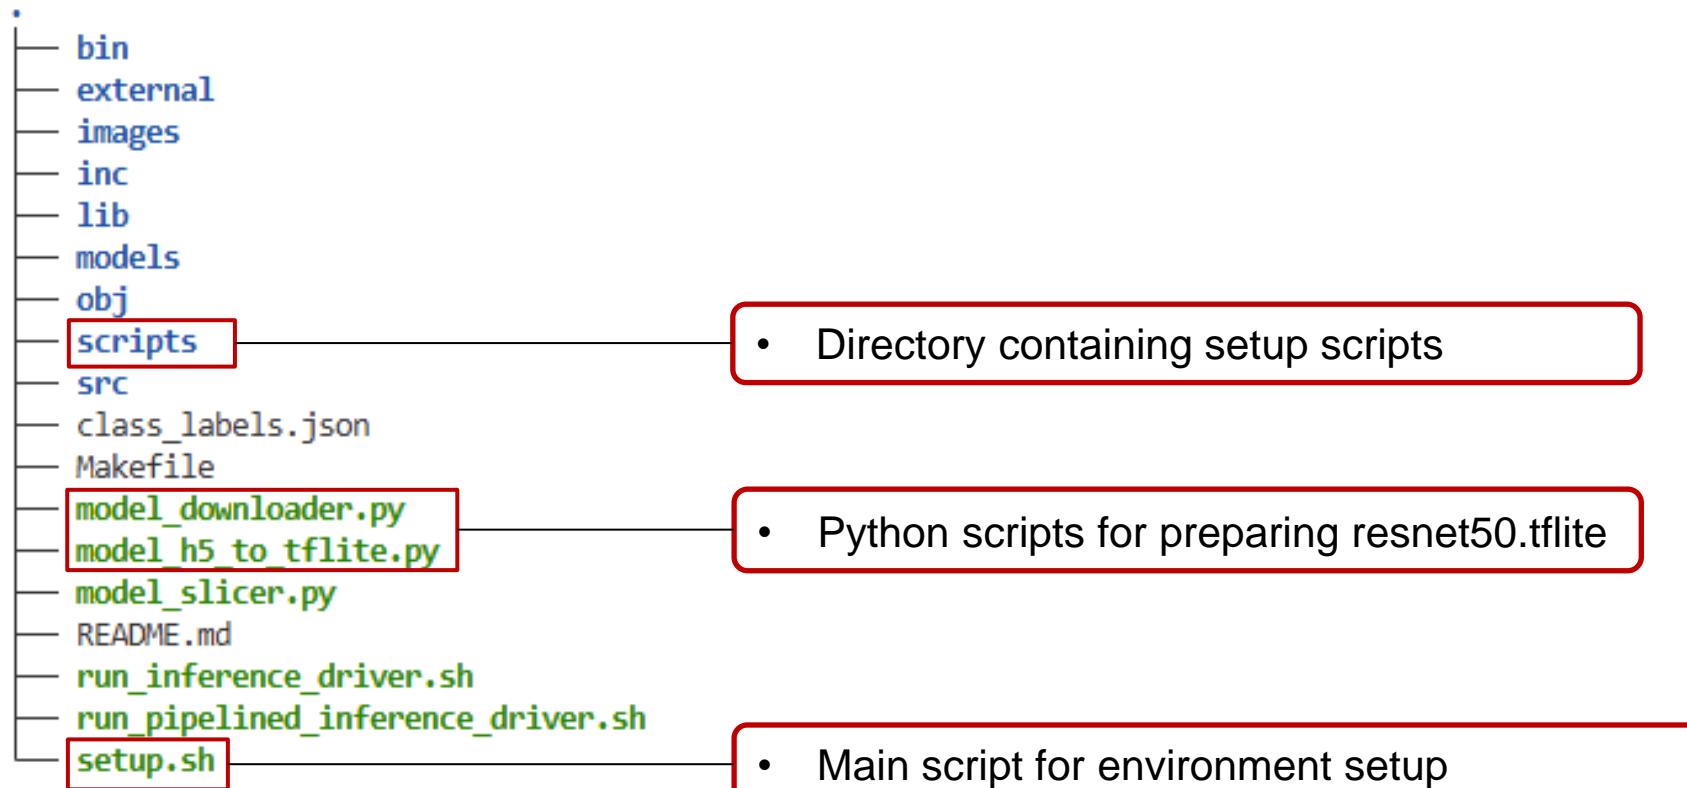
❖ Under `~/DNNPipe-Tutorial`

- Available at <https://github.com/SNU-RTOS/DNNPipe-Tutorial>

```
.
├── bin
├── external
├── images
├── inc
├── lib
├── models
├── obj
├── scripts
└── src
    ├── class_labels.json
    ├── Makefile
    ├── model_downloader.py
    ├── model_h5_to_tflite.py
    ├── model_slicer.py
    ├── README.md
    ├── run_inference_driver.sh
    ├── run_pipelined_inference_driver.sh
    └── setup.sh
```

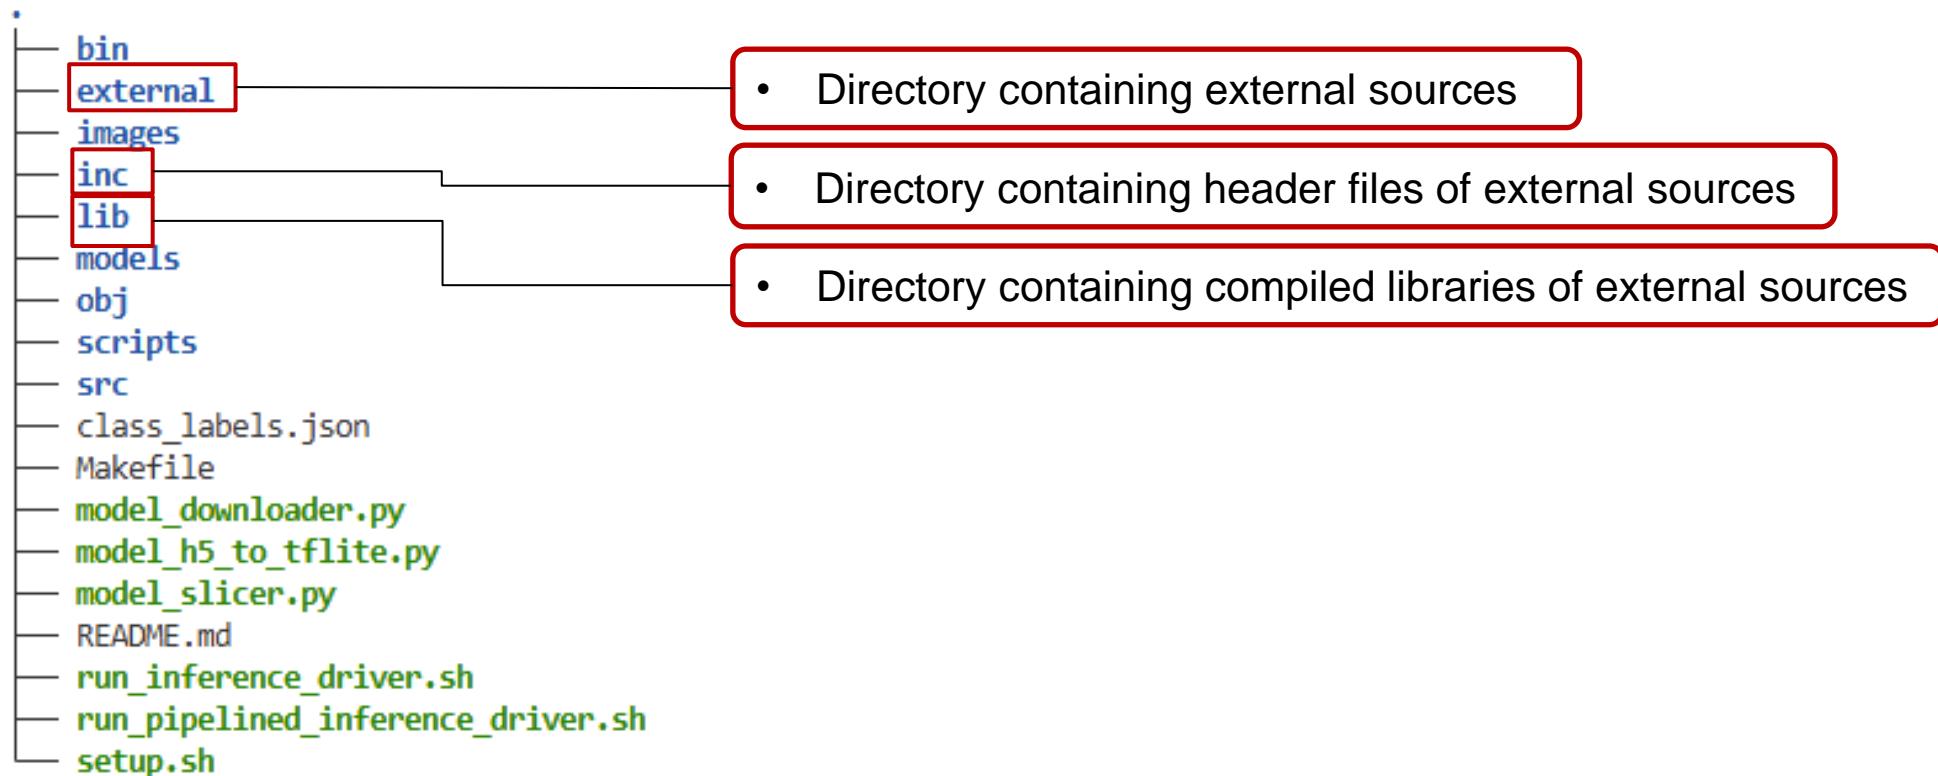
Setup Scripts

- ❖ Scripts to install software packages and prepare required files



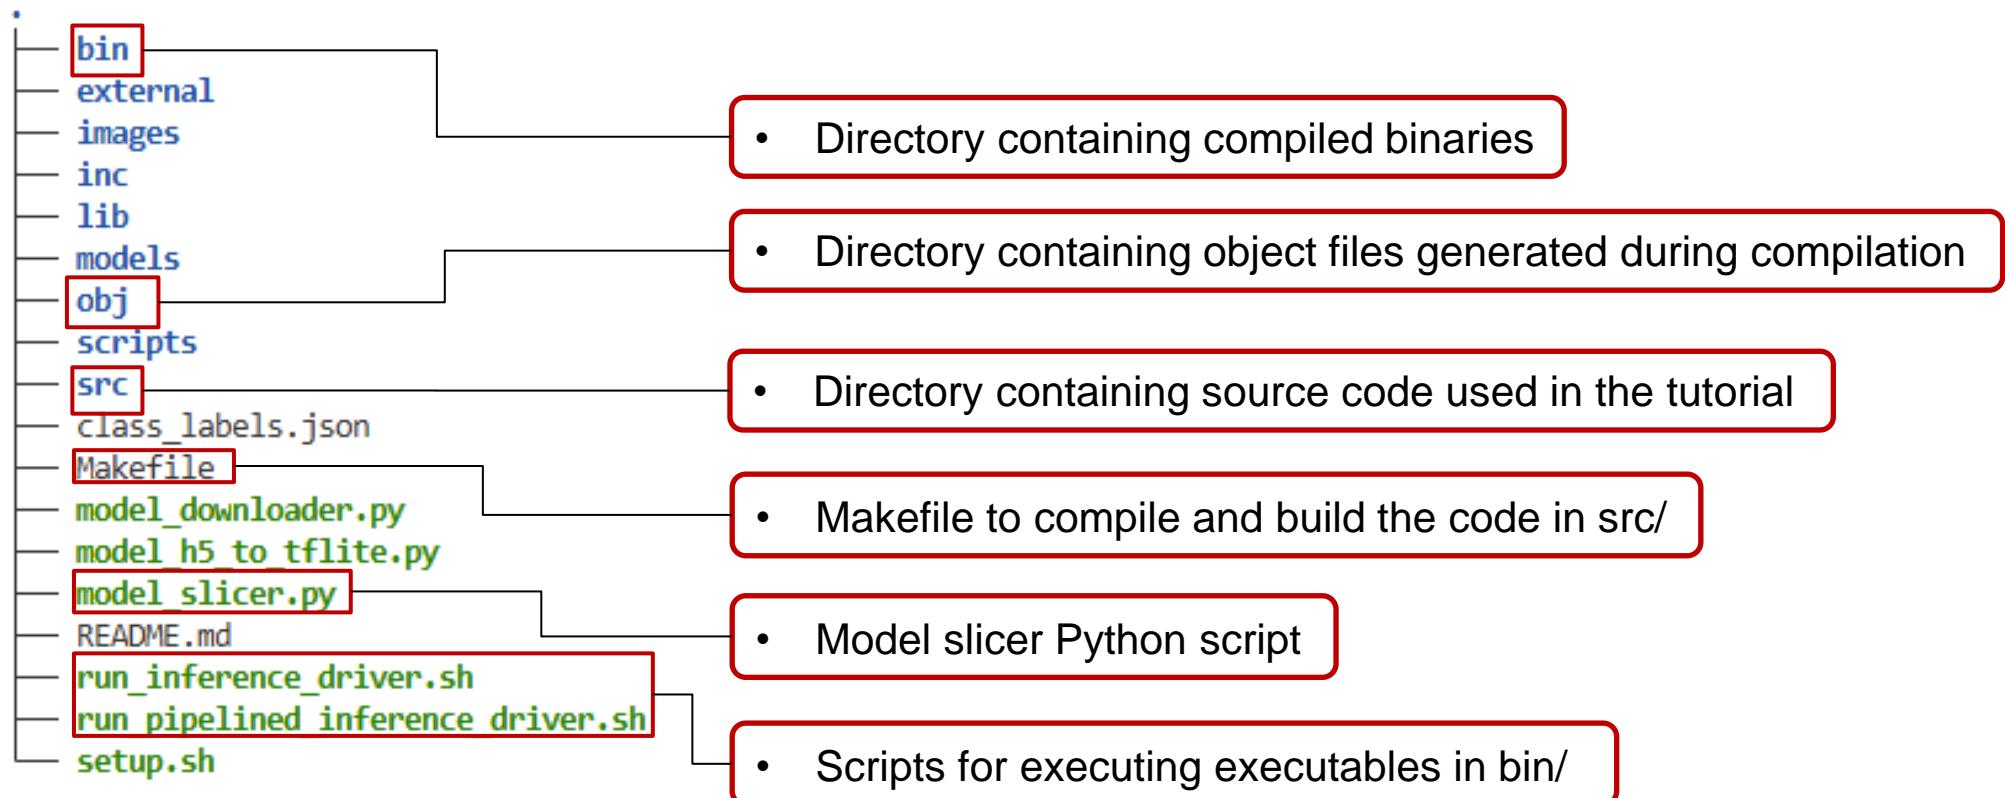
External Libraries

- ❖ Directories for externals libraries such as LiteRT



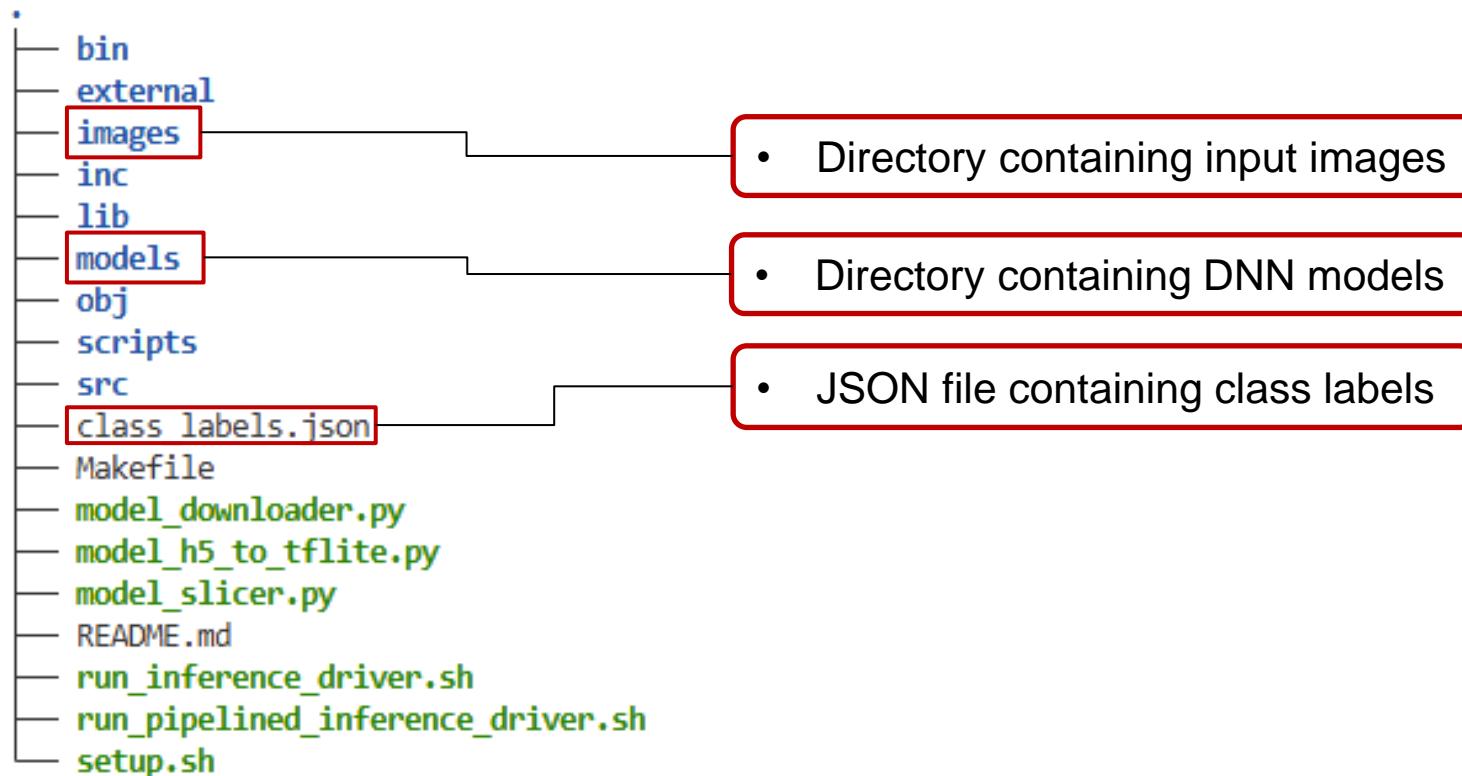
Source Code and Build File

- ❖ Source code and build file used throughout the tutorial



Input Data and Model Files

❖ Input files used by the compiled binaries



Contents

- I. Directory Structure
- II. Step-by-Step Inference Driver Walkthrough**
- III. Hands-On Exercise: Build and Run Inference Driver
- IV. Two Key Internal Entities of LiteRT
- V. Instrumentation Harness

Header Files

❖ Include necessary header files

DNNPipe-Tutorial > src > inference_driver.cpp

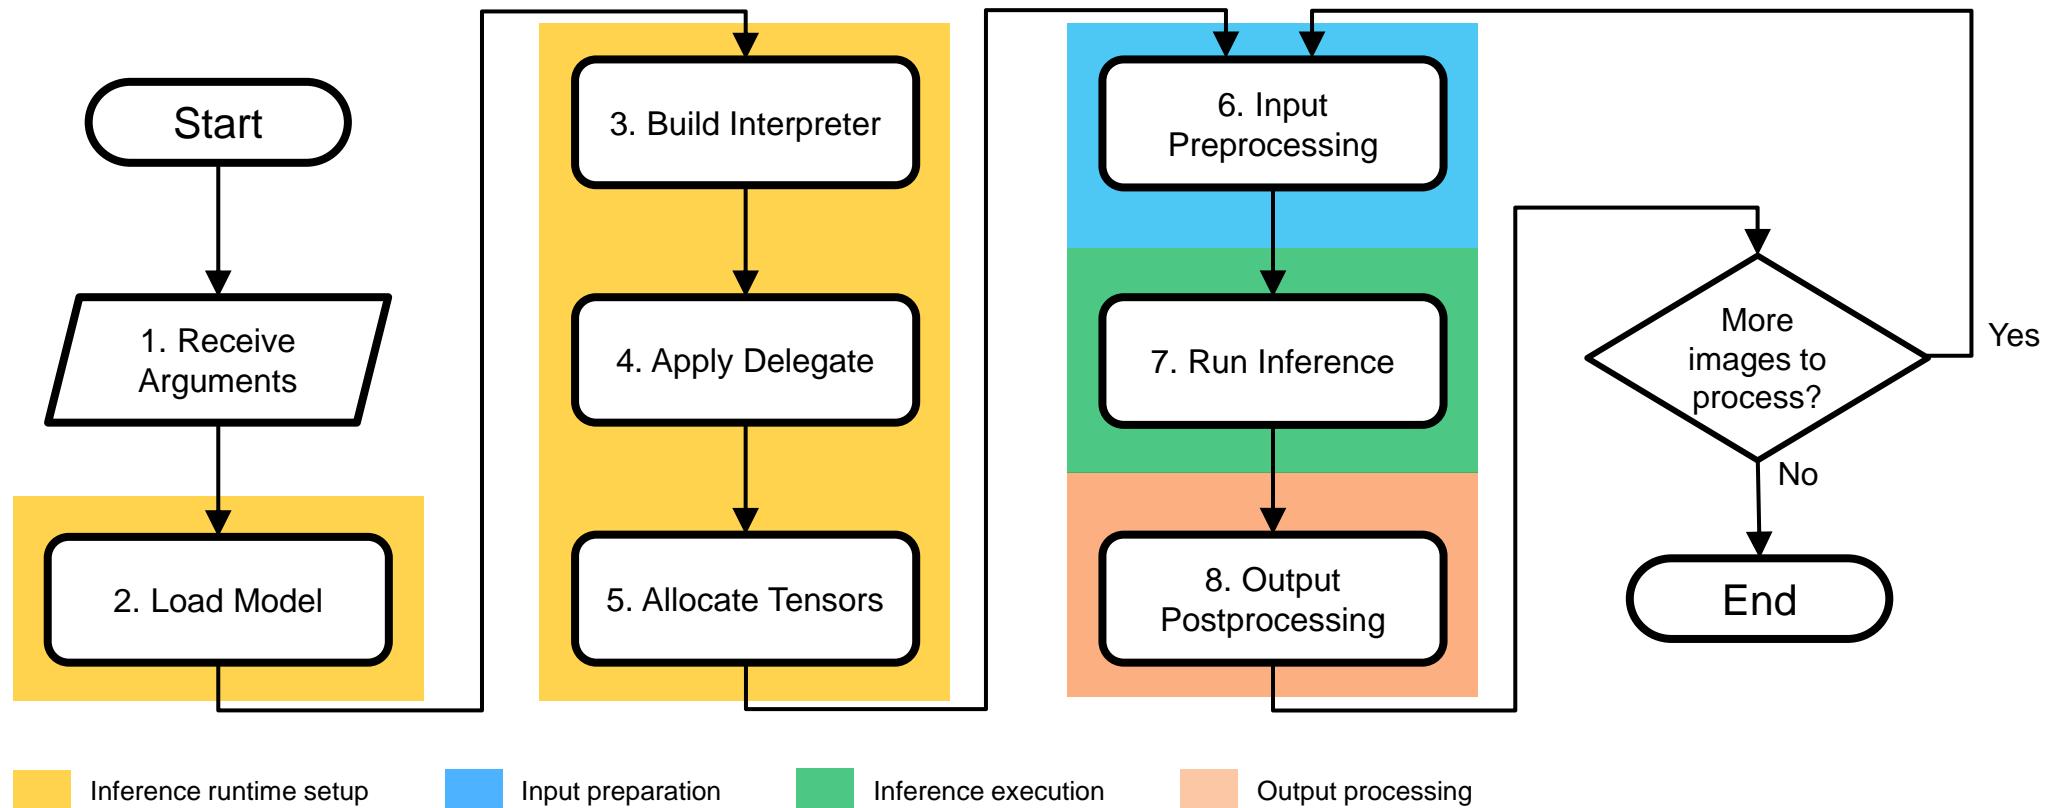
```
15 #include <iostream>
16 #include <thread>
17 #include <vector>
18 #include <opencv2/opencv.hpp>
19 #include "tflite/delegates/xnnpack/xnnpack_delegate.h"
20 #include "tflite/delegates/gpu/delegate.h"
21 #include "tflite/interpreter_builder.h"
22 #include "tflite/interpreter.h"
23 #include "tflite/kernels/register.h"
24 #include "tflite/model_builder.h"
25 #include "util.hpp"
```

C++ standard library headers for I/O stream, thread management, and dynamic arrays
Header file for OpenCV used during input preprocessing

Header files for using LiteRT

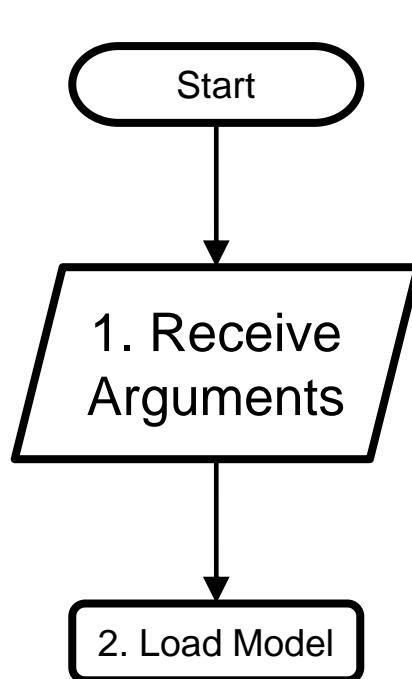
Header file for utility functions used during the tutorial

Inference Driver Code Flow



1. Receive Arguments (1)

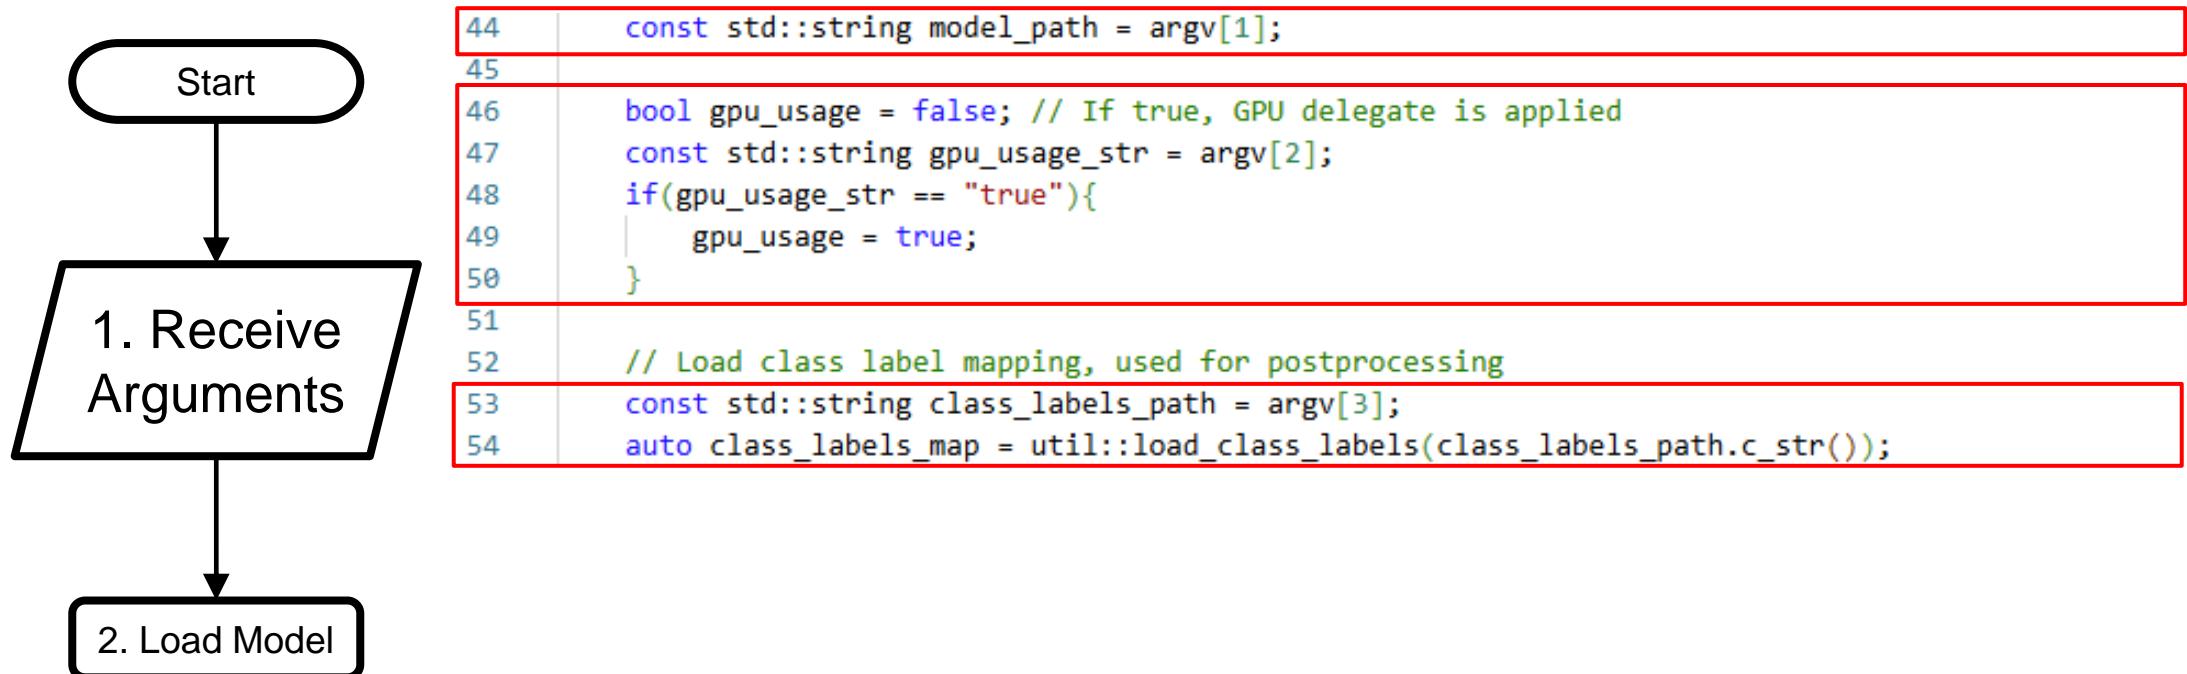
- ❖ Check the number of arguments



```
35     if (argc < 5)
36     {
37         std::cerr << "Usage: " << argv[0]
38         << "<model_path> <gpu_usage> <class_labels_path> <image_path 1> "
39         << "[image_path 2 ... image_path N] [--input-period=milliseconds]"
40         << std::endl;
41     }
42 }
```

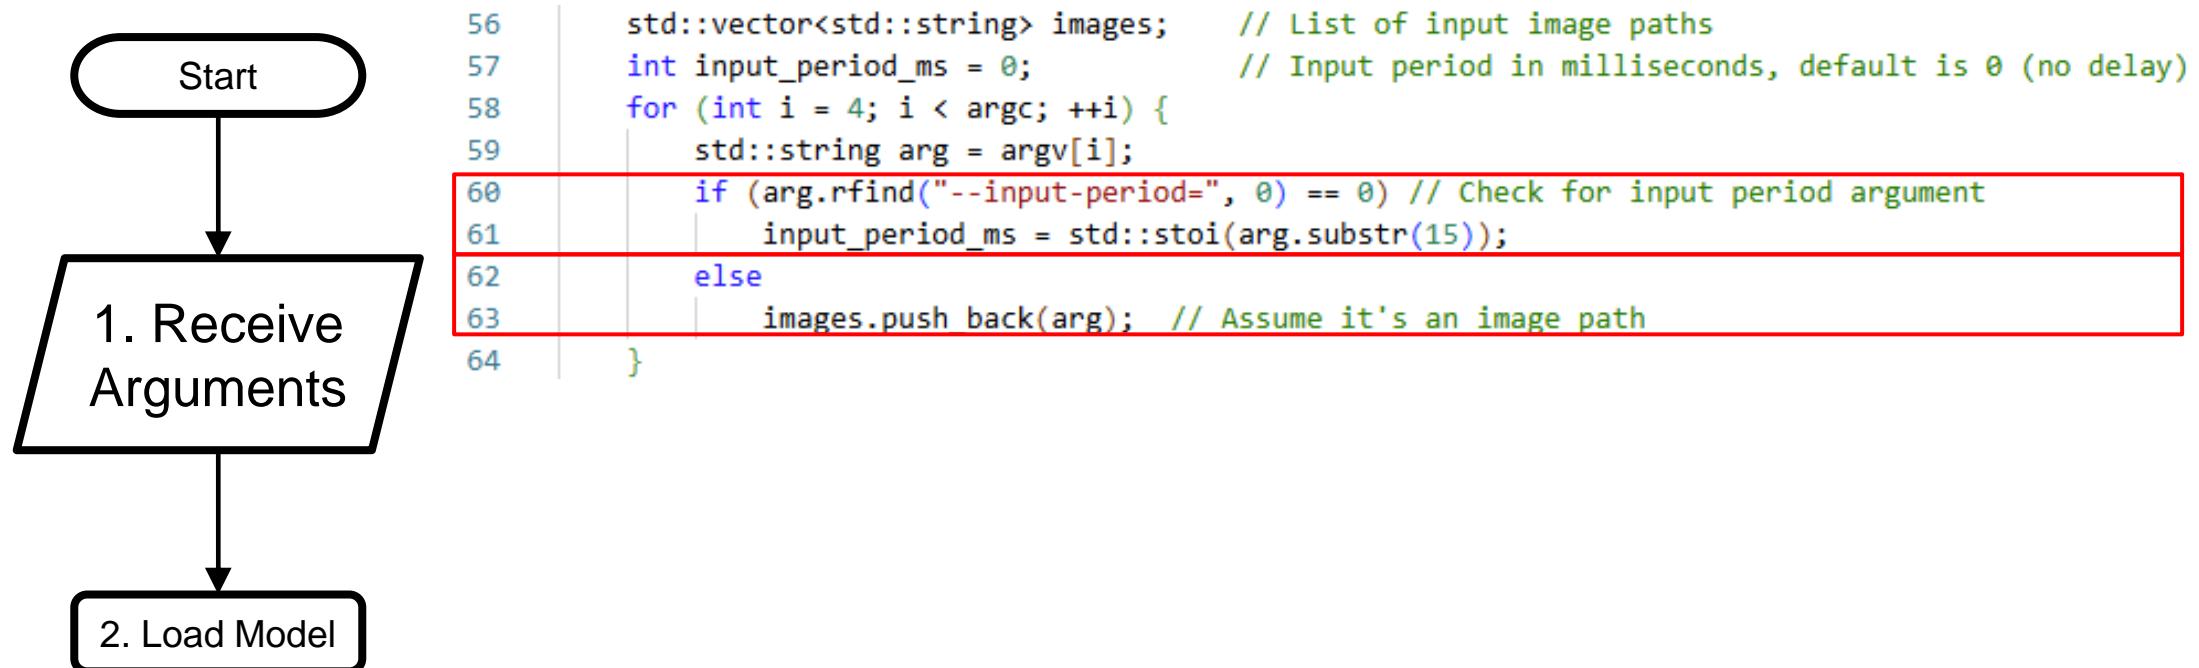
1. Receive Arguments (2)

- ❖ Set variables based on the arguments



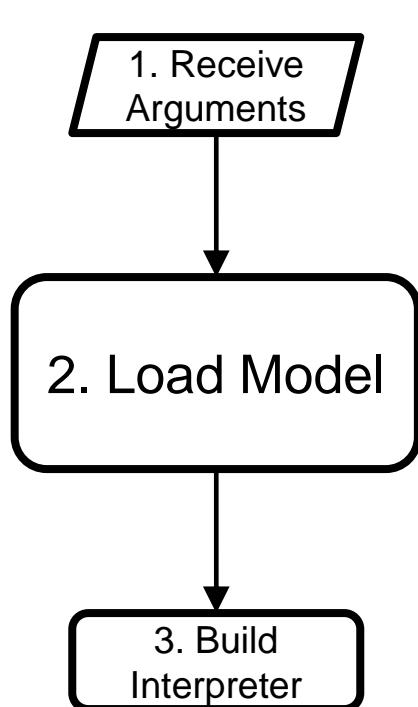
1. Receive Arguments (3)

- ❖ Set variables based on the arguments (cont'd)



2. Load Model

- ❖ Load a model file at `model_path`

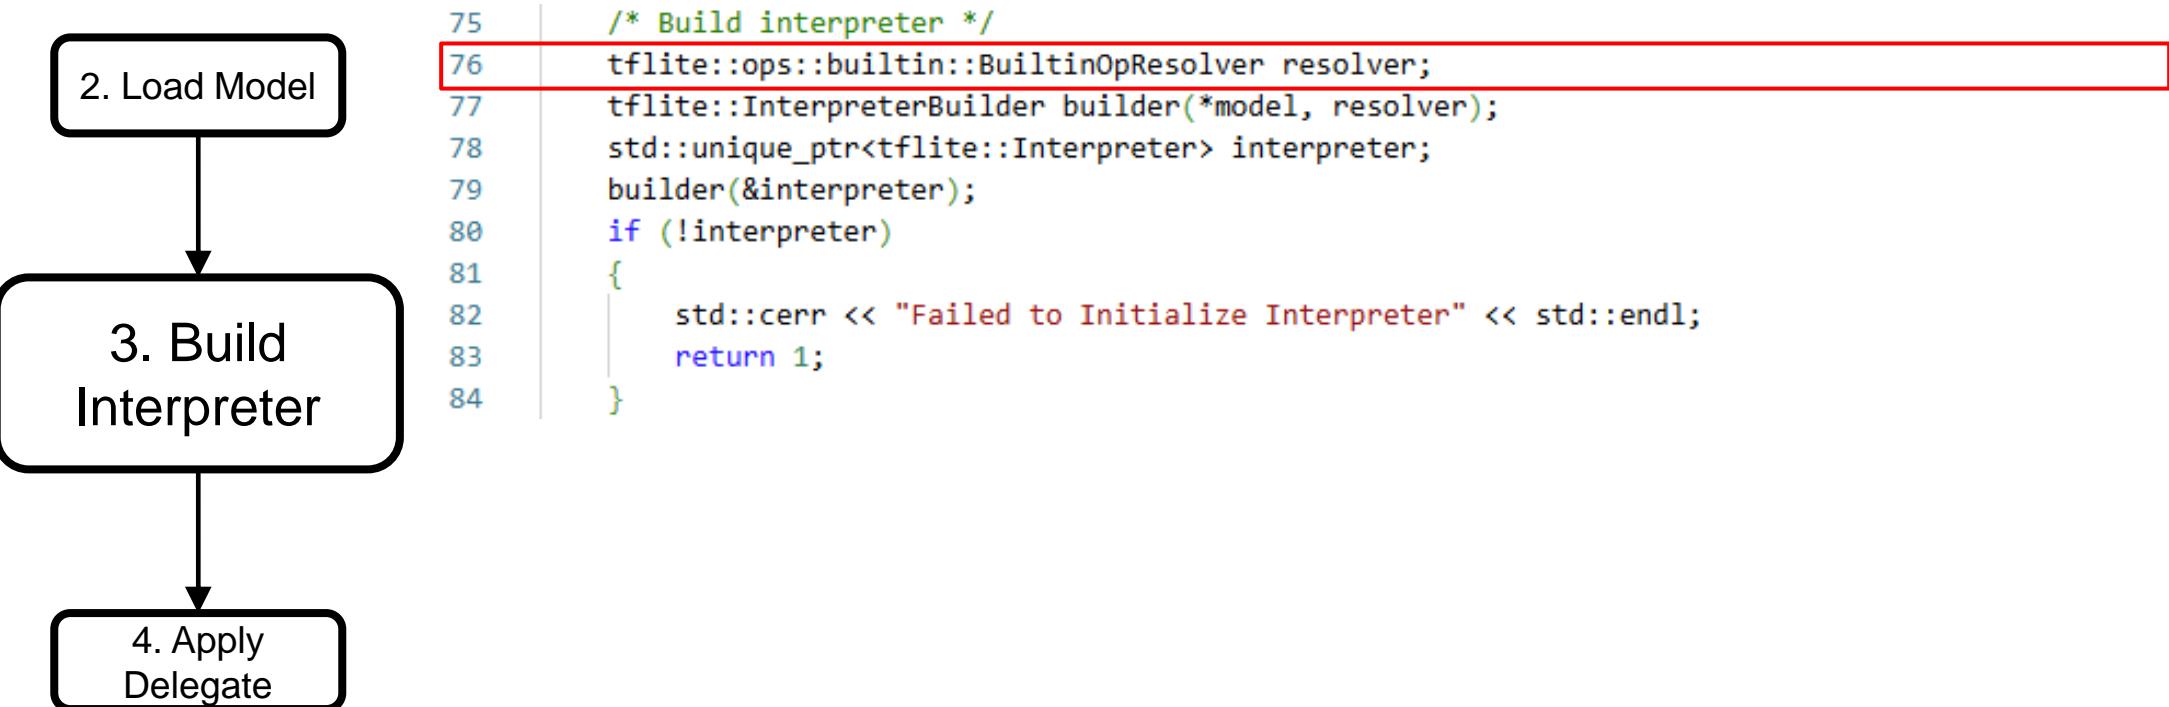


```
66  /* Load model */
67  std::unique_ptr<tflite::FlatBufferModel> model =
68  | tflite::FlatBufferModel::BuildFromFile(model_path.c_str());
69  if (!model)
70  {
71  | std::cerr << "Failed to load model" << std::endl;
72  }
73 }
```

3. Build Interpreter (1)

❖ Create an **OpResolver**

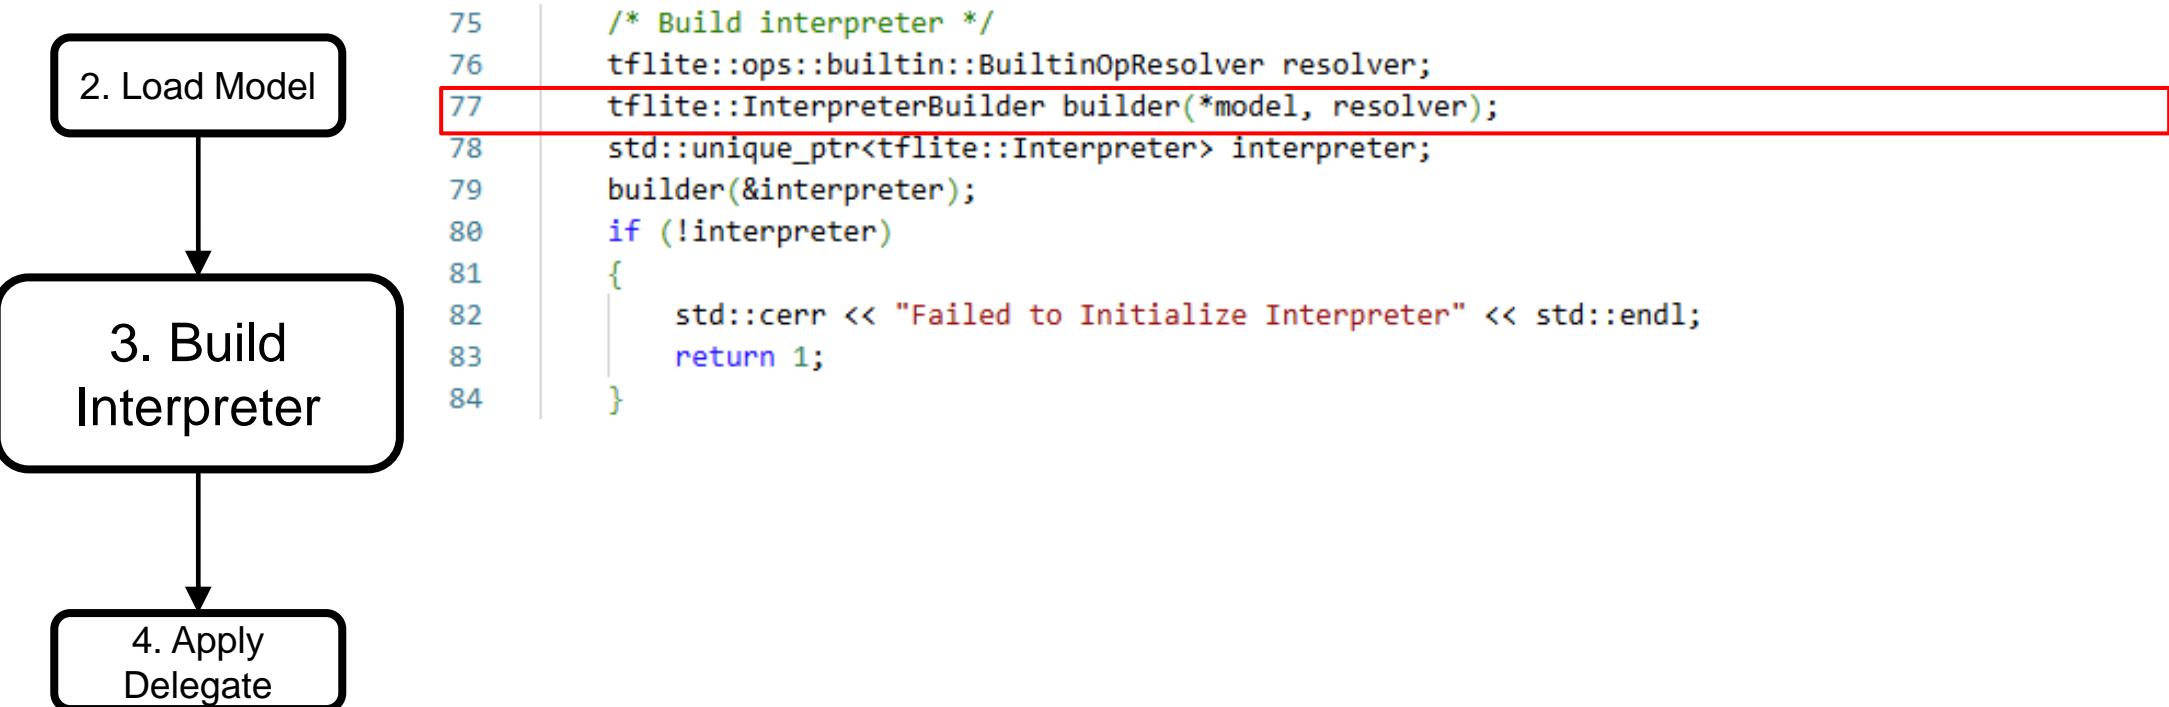
- A registry that maps operator names to their corresponding implementations



3. Build Interpreter (2)

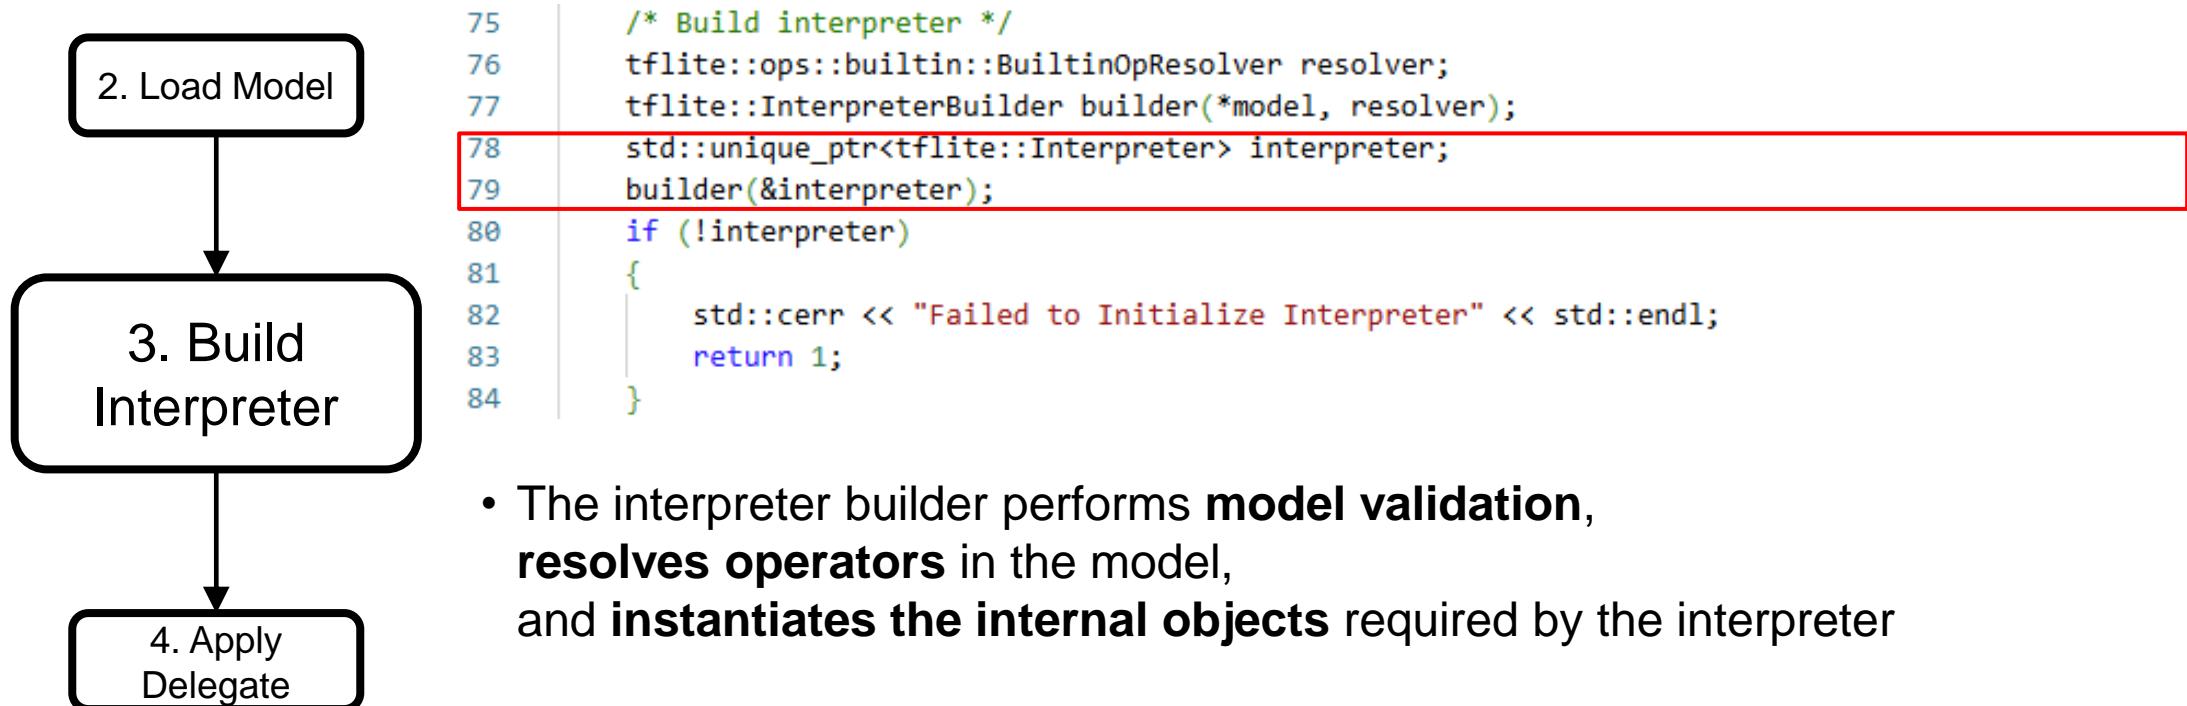
❖ Create an **InterpreterBuilder**

- An object that constructs an interpreter instance using the builder pattern



3. Build Interpreter (3)

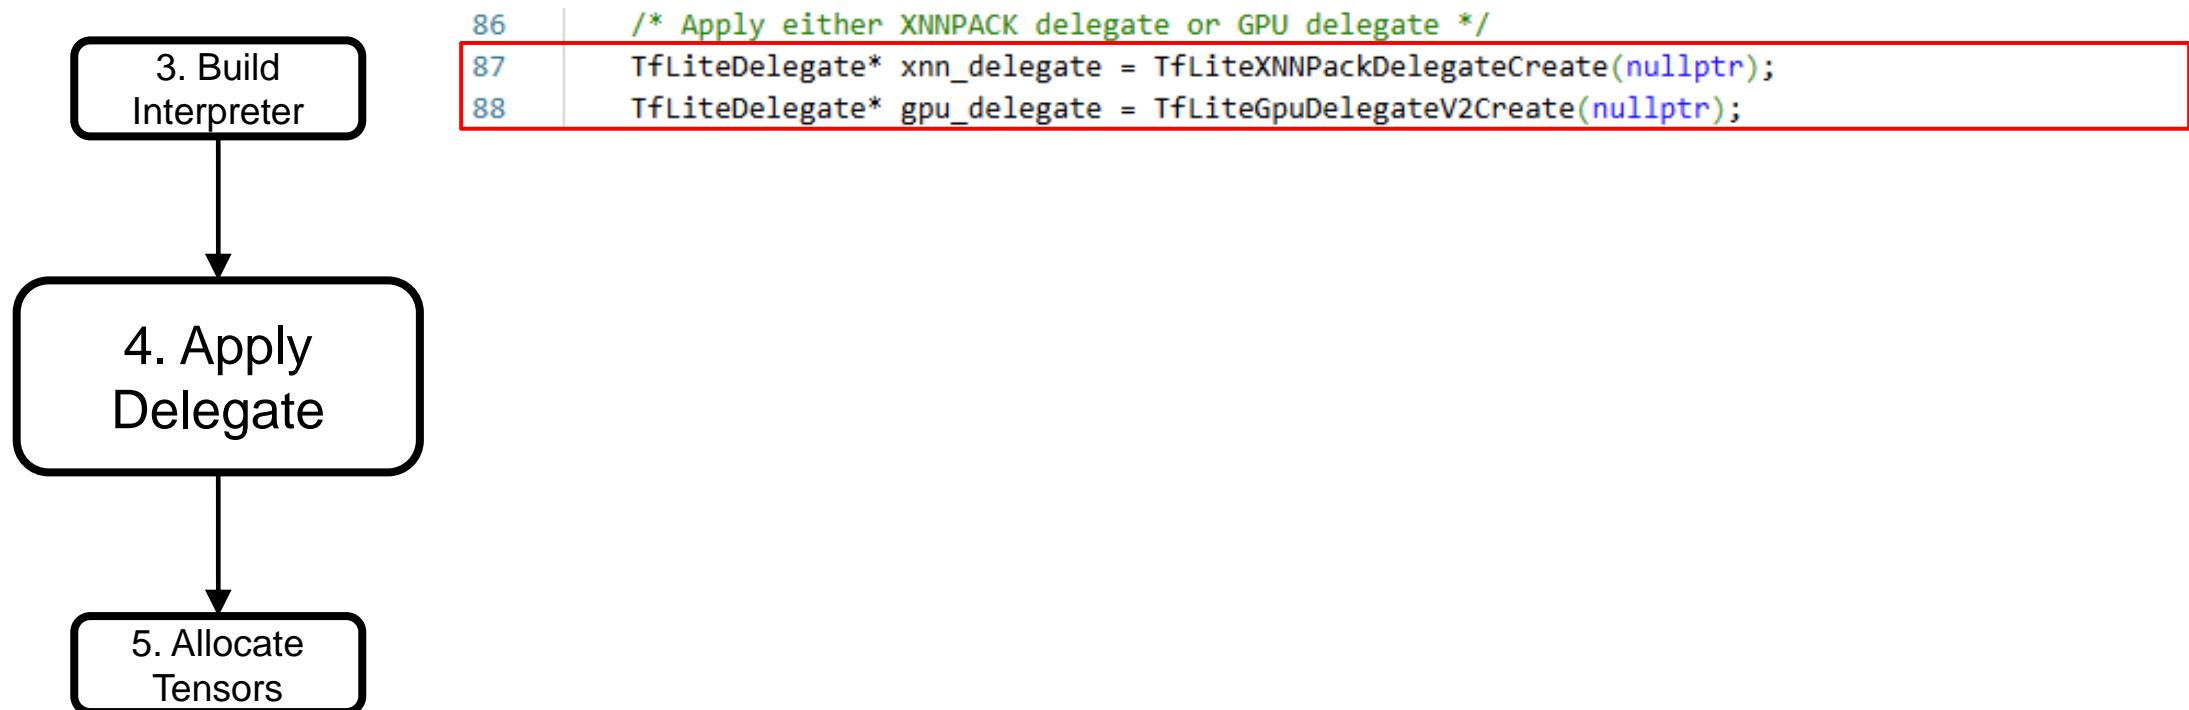
- ❖ Create and build an interpreter
 - If successfully, `_litet_interpreter` will point to a valid interpreter object



4. Apply Delegate (1)

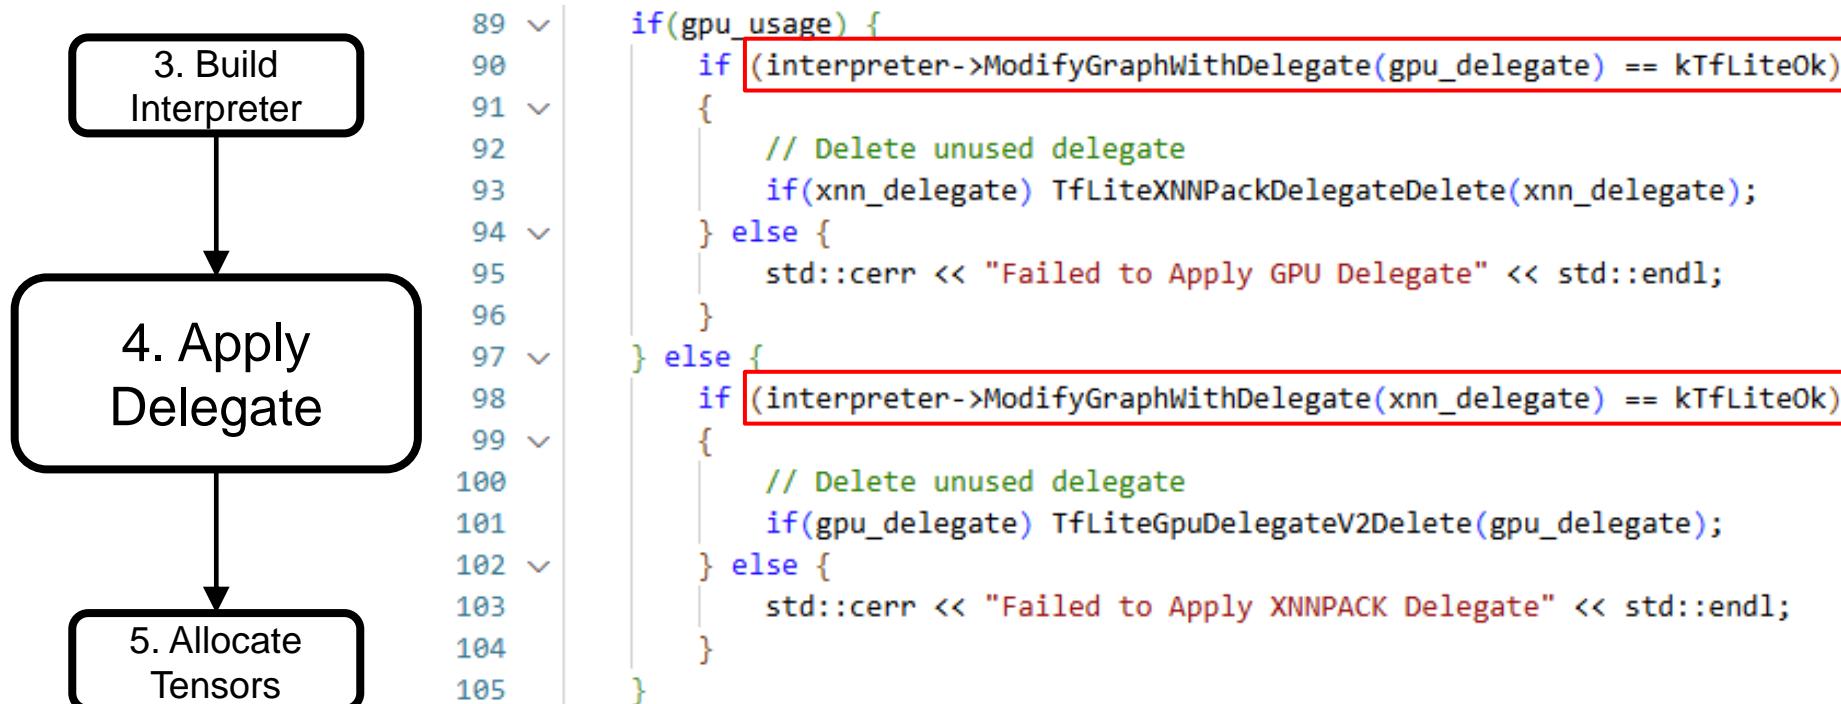
- ❖ Create delegate object(s)

- For CPU execution optimization (XNNPACK) or GPU offloading



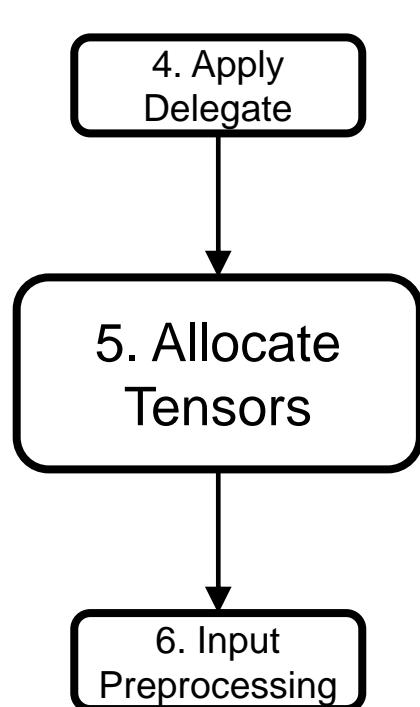
4. Apply Delegate (2)

- ❖ Apply either the GPU delegate or the XNNPACK delegate
 - The delegate takes over supported parts of the model



5. Allocate Tensors

- ❖ Allocate memory for mutable tensors
 - Tensors that store input data, intermediate values, and final results during inference



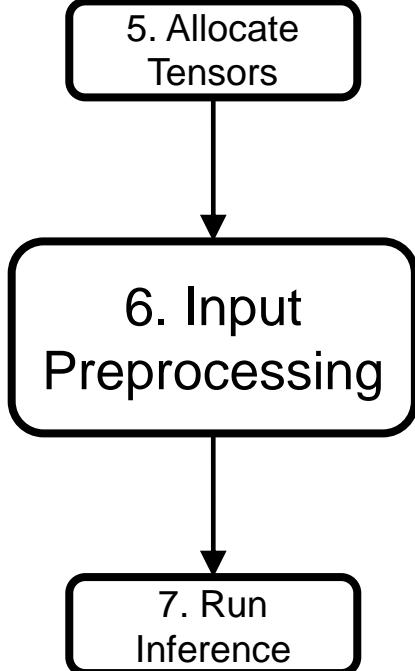
```
107  /* Allocate Tensors */  
108  if (interpreter->AllocateTensors() != kTfLiteOk)  
109  {  
110      std::cerr << "Failed to Allocate Tensors" << std::endl;  
111      return 1;  
112 }
```

- If any tensor is not allocated,
a segmentation fault will occur during inference

6. Input Preprocessing (1)

❖ Enters a do-while loop

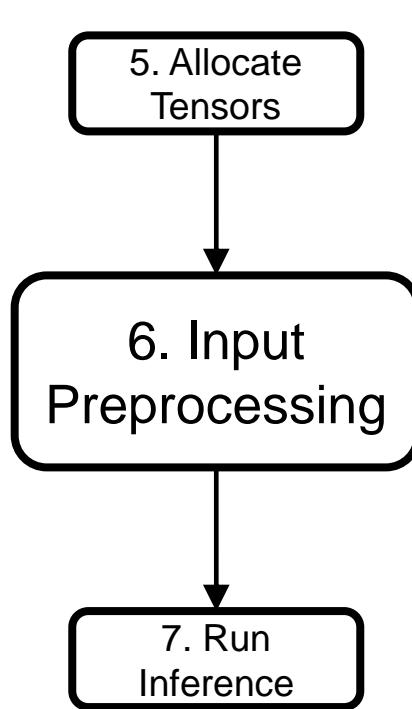
- Record the current time before starting, for use in periodic loop execution later



```
118     auto next_wakeup_time = std::chrono::high_resolution_clock::now();
119     do {
127         /* Preprocessing */
128         // Load input image
129         cv::Mat image = cv::imread(images[count]);
130         if (image.empty()) {
131             std::cerr << "Failed to load image: " << images[count] << "\n";
132             continue;
133         }
134
135         // Preprocess input data
136         cv::Mat preprocessed_image =
137             util::preprocess_image_resnet(image, 224, 224);
138
139         // Copy preprocessed_image to input_tensor
140         float* input_tensor = interpreter->typed_input_tensor<float>(0);
141         std::memcpy(input_tensor, preprocessed_image.ptr<float>(),
142                     preprocessed_image.total() * preprocessed_image.elemSize());
```

6. Input Preprocessing (2)

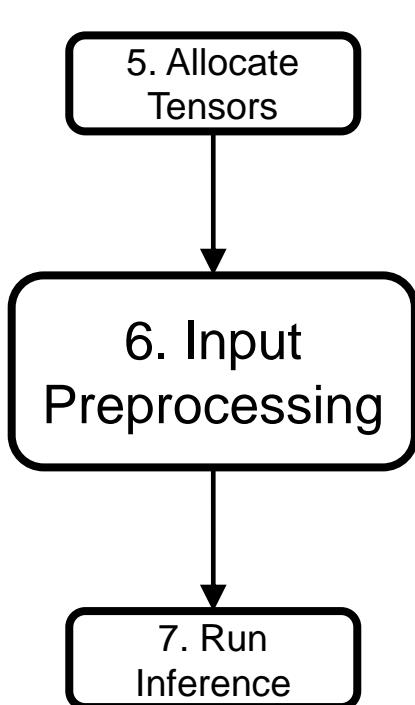
- ❖ Load an image



```
118     auto next_wakeup_time = std::chrono::high_resolution_clock::now();
119     do {
127         /* Preprocessing */
128         // Load input image
129         cv::Mat image = cv::imread(images[count]);
130         if (image.empty()) {
131             std::cerr << "Failed to load image: " << images[count] << "\n";
132             continue;
133         }
134
135         // Preprocess input data
136         cv::Mat preprocessed_image =
137             util::preprocess_image_resnet(image, 224, 224);
138
139         // Copy preprocessed_image to input_tensor
140         float* input_tensor = interpreter->typed_input_tensor<float>(0);
141         std::memcpy(input_tensor, preprocessed_image.ptr<float>(),
142                     preprocessed_image.total() * preprocessed_image.elemSize());
```

6. Input Preprocessing (3)

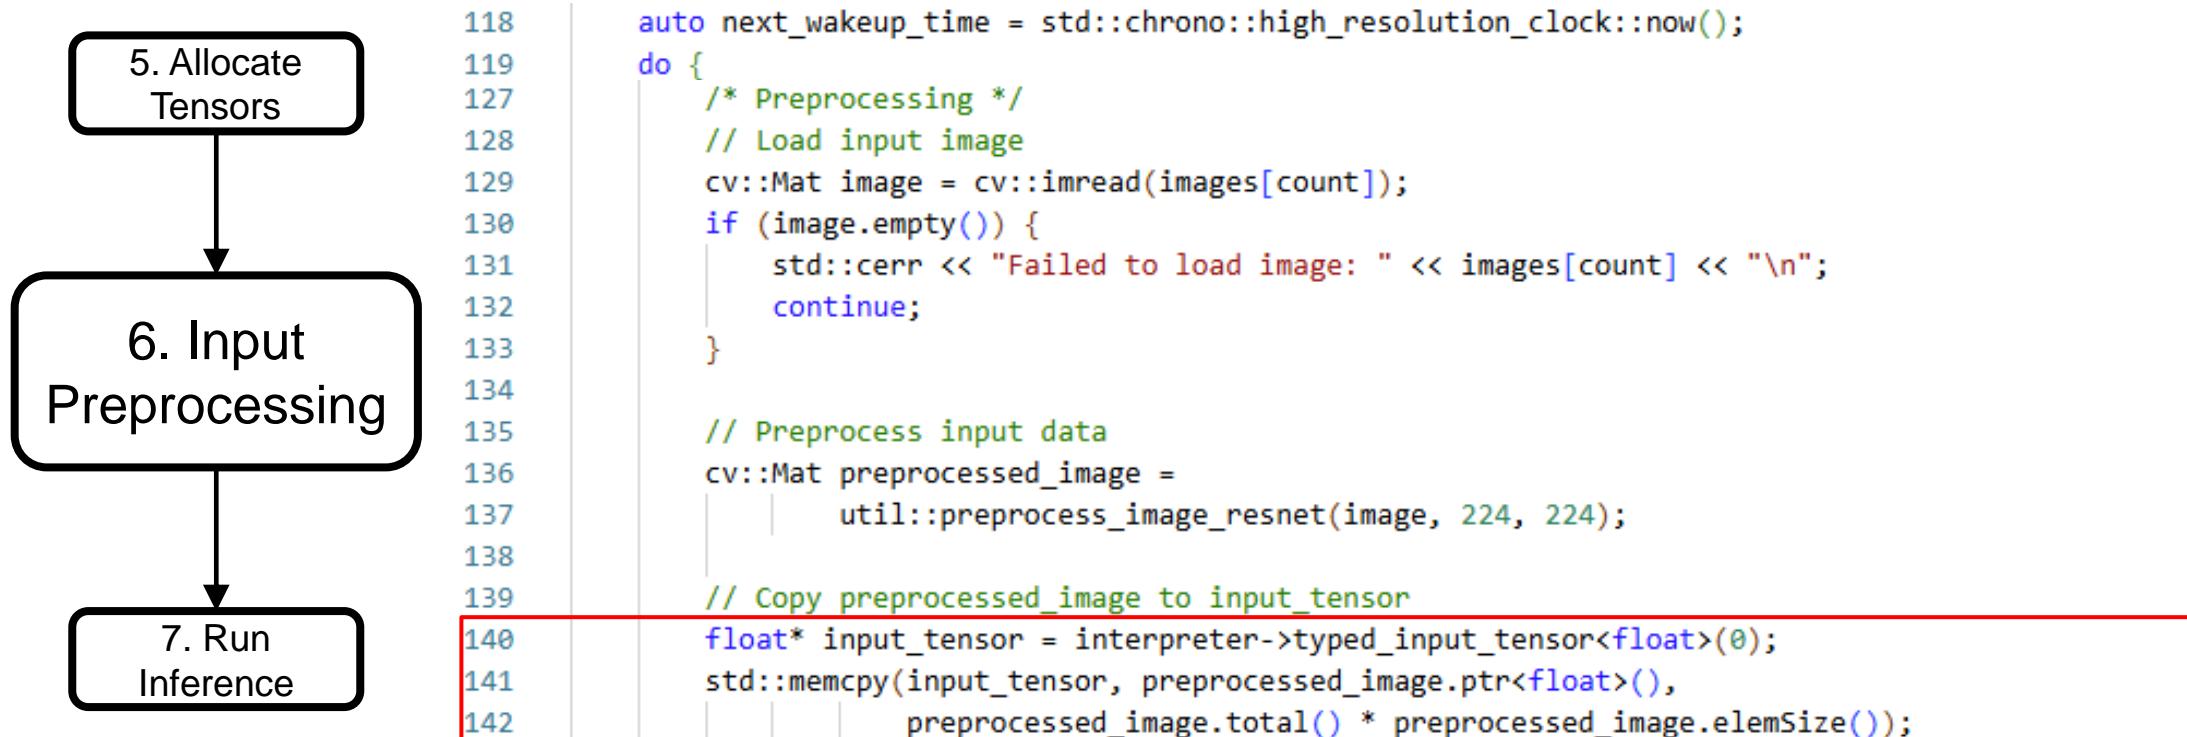
- ❖ Preprocess the loaded image



```
118     auto next_wakeup_time = std::chrono::high_resolution_clock::now();
119     do {
120         /* Preprocessing */
121         // Load input image
122         cv::Mat image = cv::imread(images[count]);
123         if (image.empty()) {
124             std::cerr << "Failed to load image: " << images[count] << "\n";
125             continue;
126         }
127
128         // Preprocess input data
129         cv::Mat preprocessed_image =
130             util::preprocess_image_resnet(image, 224, 224);
131
132         // Copy preprocessed_image to input_tensor
133         float* input_tensor = interpreter->typed_input_tensor<float>(0);
134         std::memcpy(input_tensor, preprocessed_image.ptr<float>(),
135                     preprocessed_image.total() * preprocessed_image.elemSize());
136
137
138
139
140
141
142
```

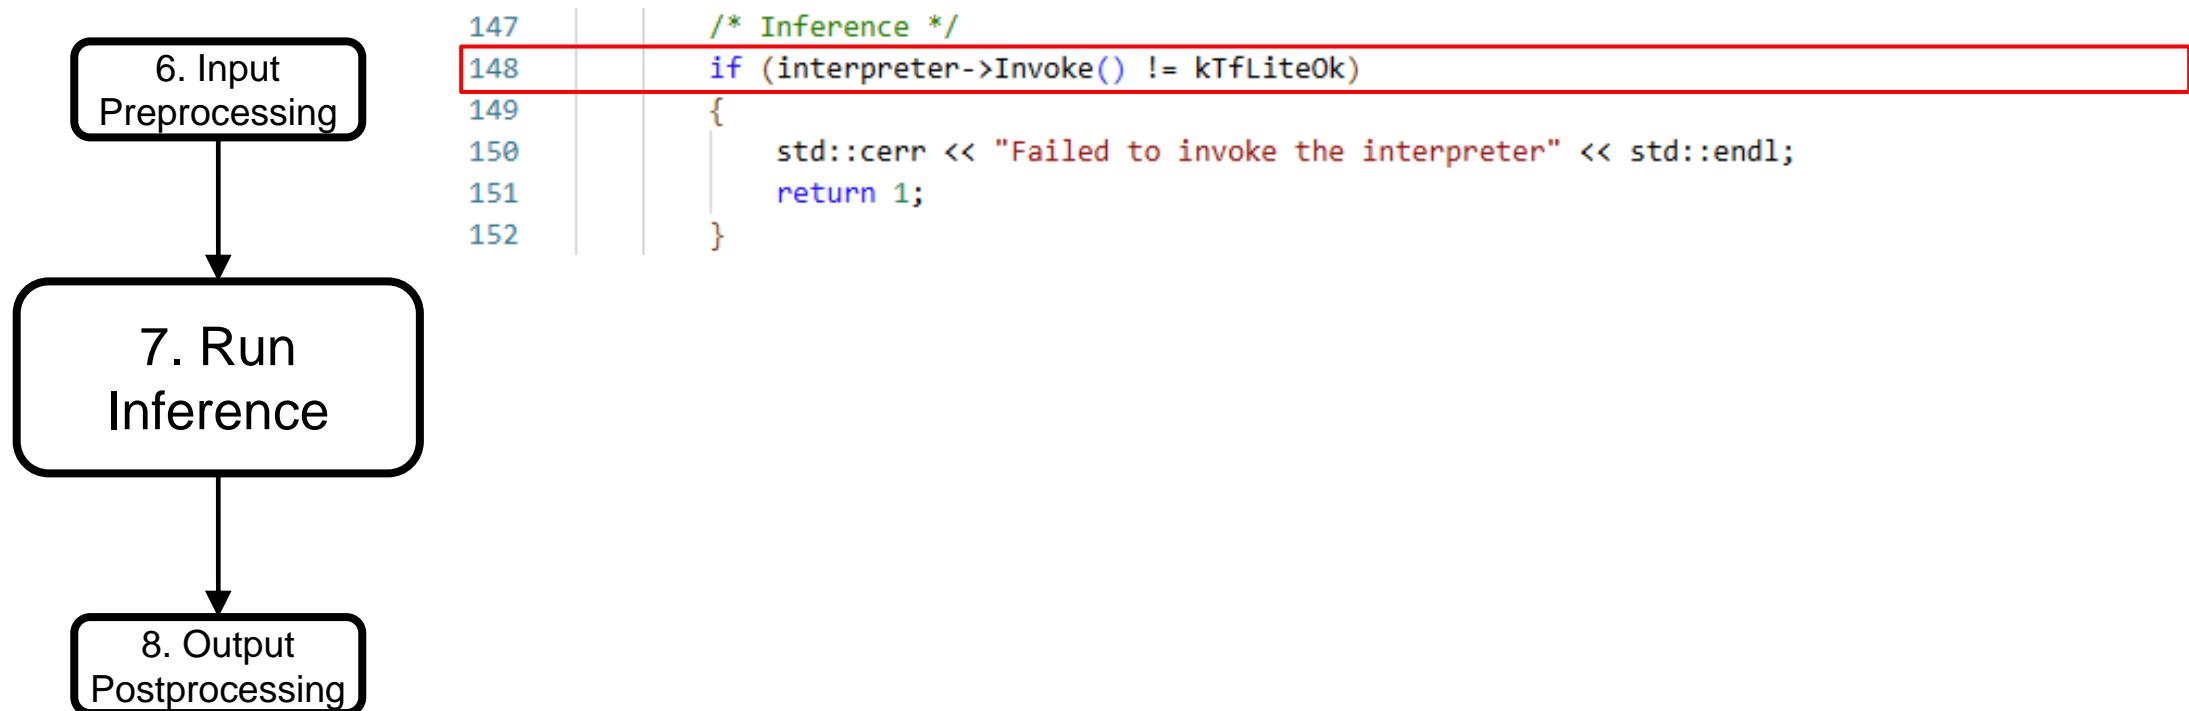
6. Input Preprocessing (4)

- ❖ Access the 0th input tensor as a float pointer and set its value
 - There is only one input tensor in ResNet50 FP32



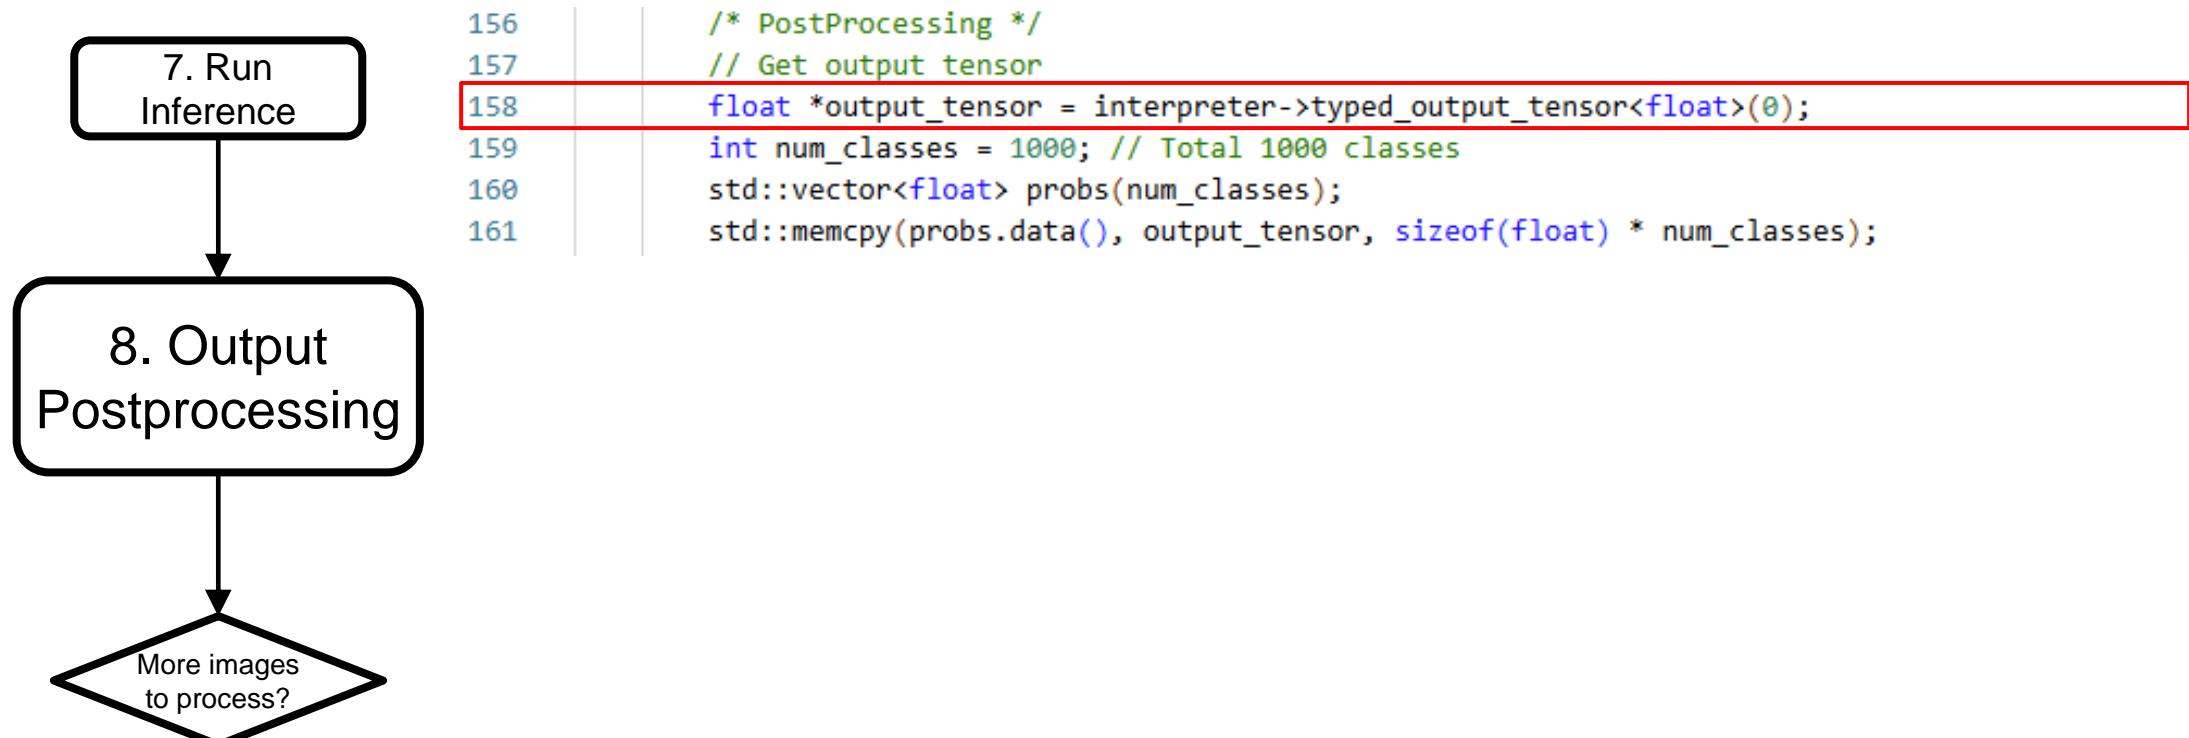
7. Run Inference

- ❖ Invoke the interpreter to run inference



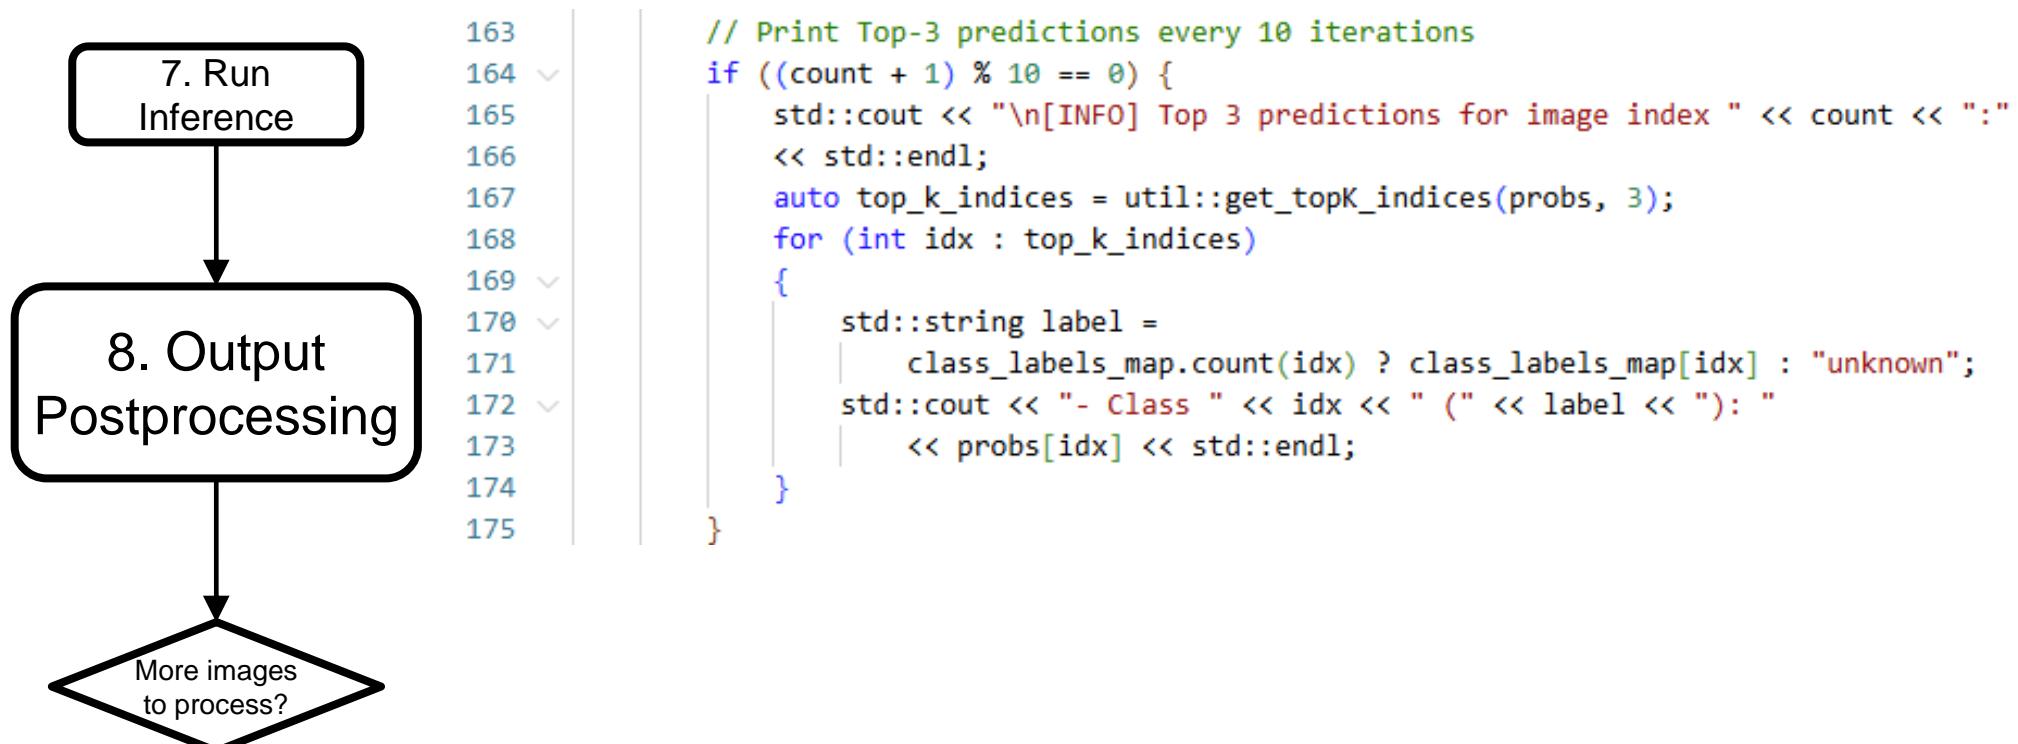
8. Output Postprocessing (1)

- ❖ Access the 0th output tensor as a float pointer
 - ResNet50 has only one output tensor with a shape of 1×1000



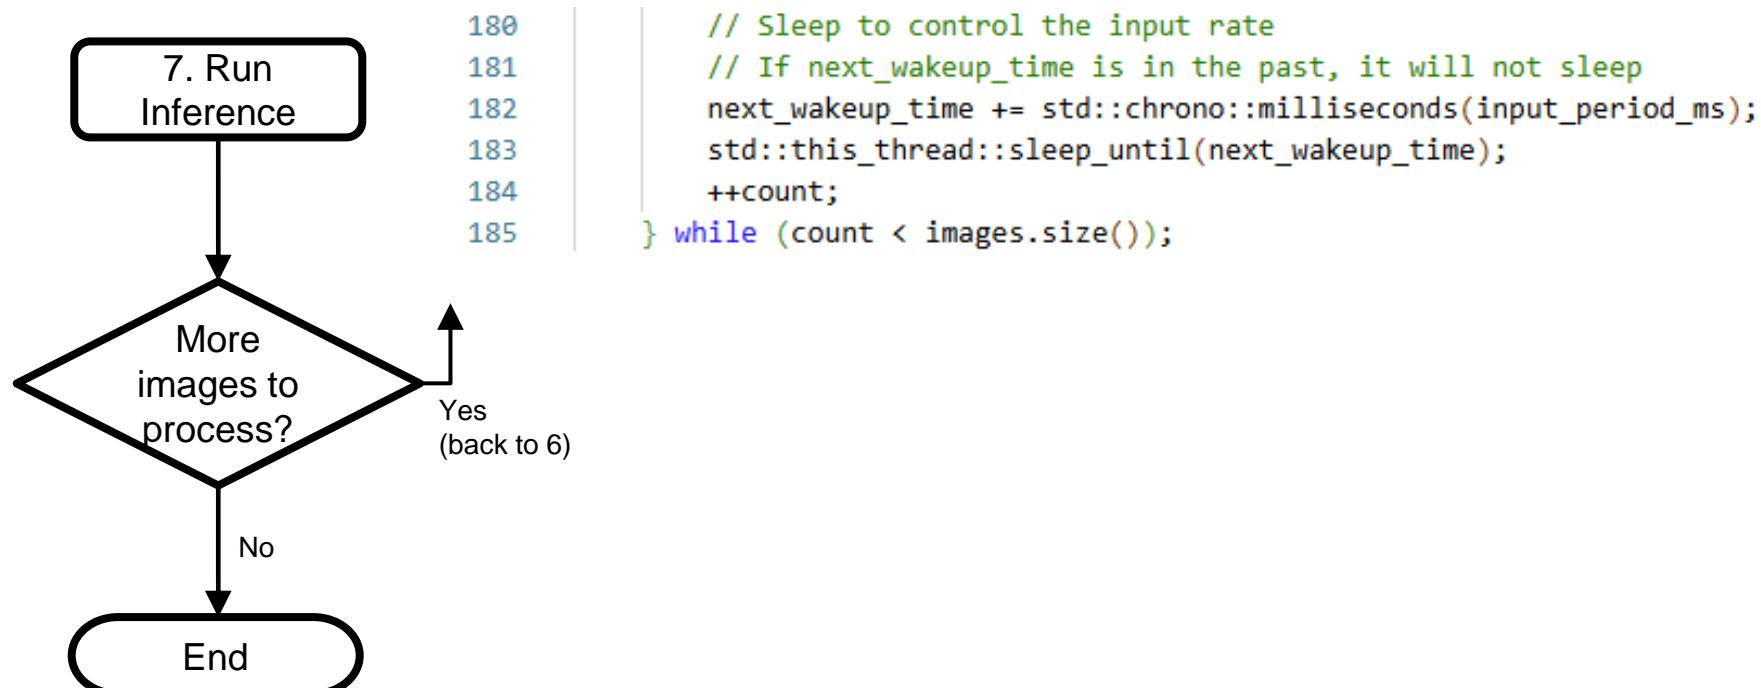
8. Output Postprocessing (2)

- ❖ Print the top 3 predictions with their labels



Periodic Loop Execution

- ❖ Repeat step 6 at each input period until all inputs are processed
 - If the scheduled time has already passed, continue without waiting



Contents

- I. Directory Structure
- II. Step-by-Step Inference Driver Walkthrough
- III. Hands-On Exercise: Build and Run Inference Driver**
- IV. Two Key Internal Entities of LiteRT
- V. Instrumentation Harness

Build and Run the Inference Driver

❖ Objective

- Build and run the inference driver and capture throughput

❖ Do

- Follow the instructions on the next slide

❖ Verify

- Check the terminal output
 - You should see the expected outputs on the next slide

❖ Time

- 2 minutes

Run the Inference Driver

- ❖ In `~/DNNPipe-Tutorial`, run
 - `make`
 - `./run_inference_driver.sh`
- ❖ Expected output

```
[INFO] Top 3 predictions for image index 499:  
- Class 867 (trailer_truck): 0.531833  
- Class 675 (moving范): 0.459434  
- Class 569 (garbage_truck): 0.00453667  
  
[INFO] Average E2E latency (500 runs): 48.72 ms  
  
[INFO] Average Preprocessing latency (500 runs): 4.934 ms  
  
[INFO] Average Inference latency (500 runs): 43.188 ms  
  
[INFO] Average Postprocessing latency (500 runs): 0 ms  
  
[INFO] Throughput: 20.312 items/sec (500 items in 24616 ms)
```

Contents

- I. Directory Structure
- II. Step-by-Step Inference Driver Walkthrough
- III. Hands-On Exercise: Build and Run Inference Driver
- IV. Two Key Internal Entities of LiteRT**
- V. Instrumentation Harness

Two Key Entities of LiteRT

-
1. Model object
 - Static description of the model stored in a .tflite file
 2. Interpreter
 - Runtime instance built from the model object

1. Model Object (1)

1. SubGraph

- Computation graph with its own operators and tensors
 - Multiple subgraphs are used to support advanced features such as control flow and modular function calls

2. Operator

- Basic unit of neural operations
- Each operator references an operator code and tensors it uses

3. Operator code

- Identifier that defines the type of neural operation

1. Model Object (2)

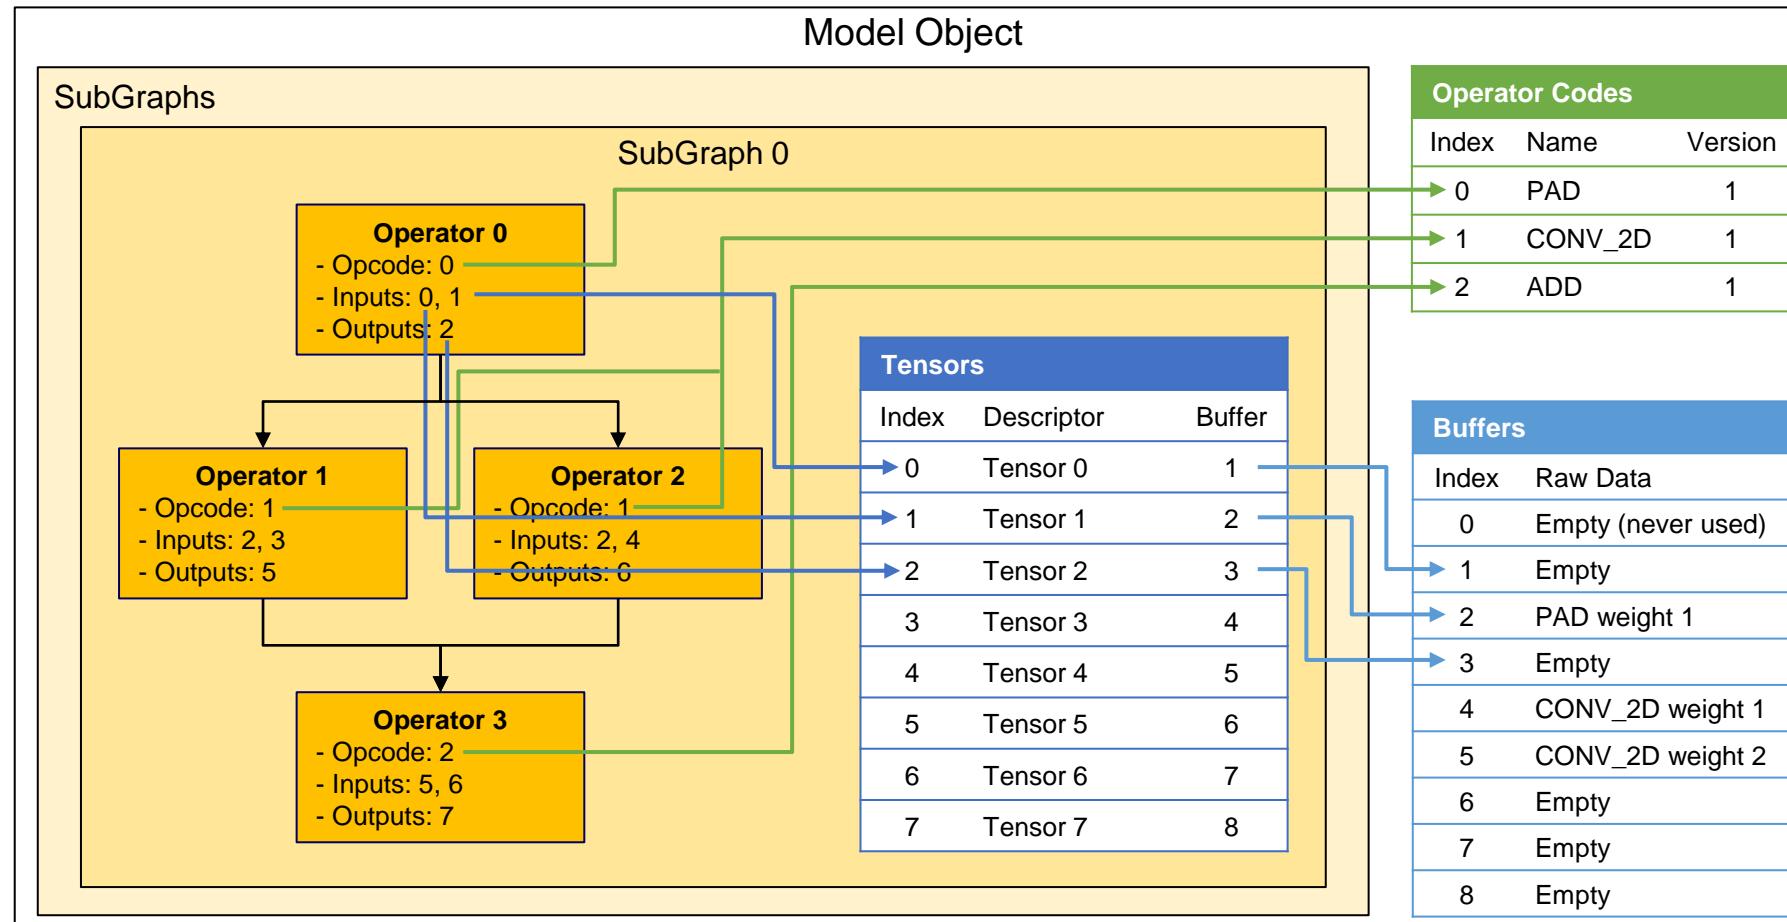
4. Tensor

- Descriptor of a multi-dimensional array instantiated at runtime
 - Defines name, shape, data type, and buffer reference

5. Buffer

- Container of raw constant data referenced by immutable tensors
 - e.g., weights and biases
 - Empty for mutable tensors

1. Model Object (3)



2. Interpreter

1. Subgraph

- Computation graph created by interpreter builder from the SubGraph (nodes, tensors, and execution plan) of the model object

2. Node

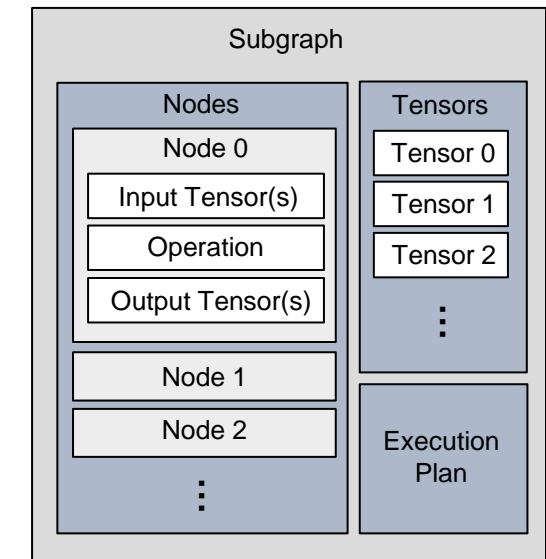
- Runtime unit of computation corresponding to an operator in the model object

3. Tensor

- Multi-dimensional array that holds inputs, outputs, weights, biases, and/or intermediate results

4. Execution plan

- Ordered list of node indices
- Defines the sequence of execution for the subgraph



Contents

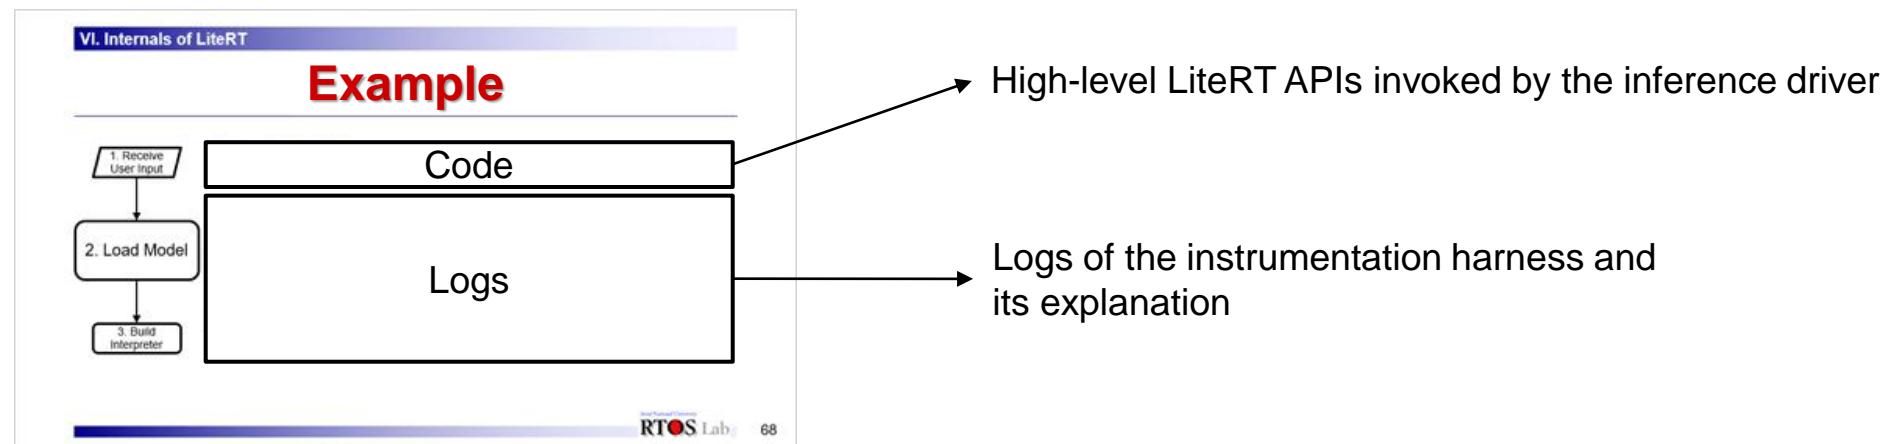
- I. Directory Structure
- II. Step-by-Step Inference Driver Walkthrough
- III. Hands-On Exercise: Build and Run Inference Driver
- IV. Two Key Internal Entities of LiteRT
- V. **Instrumentation Harness**

Instrumentation Harness for LiteRT APIs

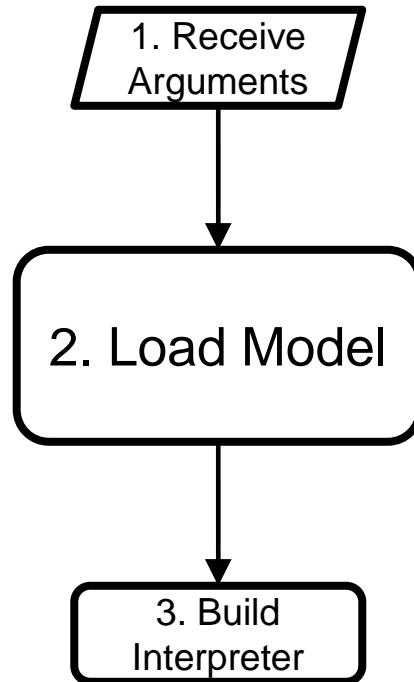
- ❖ What is an instrumentation harness?
 - Separate code that replicates the sequence of low-level LiteRT API calls invoked by high-level LiteRT APIs in an inference driver
 - Logs internal changes after each low-level LiteRT API call
- ❖ To understand the internals of high-level LiteRT APIs,
 - Follow the code flow of the inference driver and
 - Visualize logs that the implementation harness collects

Run the Instrumentation Harness

- ❖ In `~/DNNPipe-Tutorial`, run
 - `./bin/instrumentation_harness ./models/resnet50.tflite`
 - You will see an indicator `Press Enter to continue...` at each step
 - Let's go through the steps together to observe what happens internally



2. Load Model

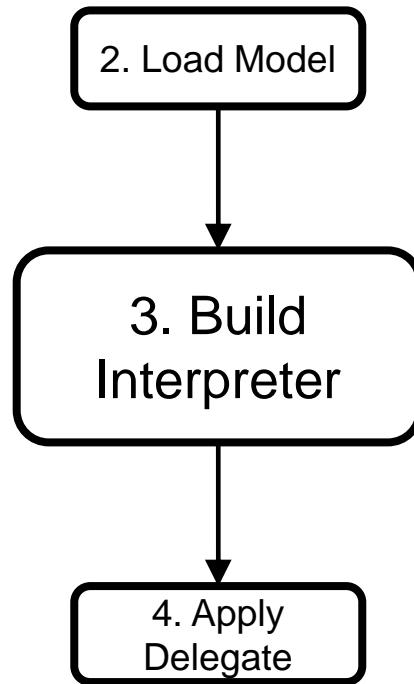


```
66  /* Load model */  
67  std::unique_ptr<tflite::FlatBufferModel> model =  
68  tflite::FlatBufferModel::BuildFromFile(model_path.c_str());
```

VM Address Permission Disk Major:Minor No.
7fa1a00000-7fa7b6d000 r--s 00000000 08:05 264200
/home/rtoslab/DNNPipe-Tutorial/models/resnet50.tflite File Path

- The model file is memory-mapped into the process's virtual address space
- You can verify this by looking at `/proc/<pid>/maps` of the process

3. Build Interpreter (1)

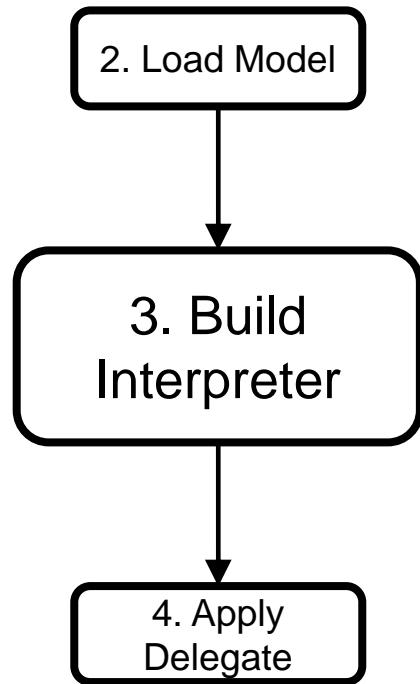


```
75  /* Build interpreter */  
76  tflite::ops::builtin::BuiltinOpResolver resolver;  
77  tflite::InterpreterBuilder builder(*model, resolver);  
78  std::unique_ptr<tflite::Interpreter> interpreter;  
79  builder(&interpreter);
```

Schema version of the model: 3
Supported schema version: 3

- What is a **schema** in LiteRT?
 - Specification that defines of how a LiteRT model is serialized in FlatBuffer format
- Validates that the model's schema version matches the runtime's supported version
- If the versions do not match, an error is returned
 - Because the runtime cannot properly parse the model file
- The schema definition for LiteRT can be found at
`~/DNNPipe-Tutorial/external/litert/bazel-litert/external/org_tensorflow/tensorflow/compiler/mlir/lite/schema/schema.fbs`
- You can deserialize a LiteRT model using FlatBuffers,
but this is beyond the scope of this tutorial

3. Build Interpreter (2)

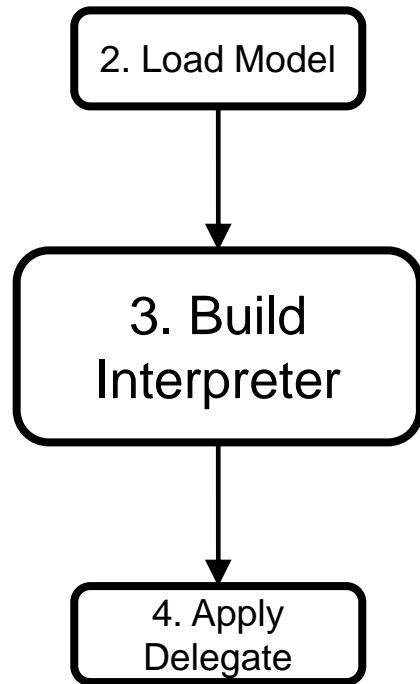


Total 7 operator codes in the model

```
[0] PAD, version: 1, supported (Y/N): Y  
[1] CONV_2D, version: 1, supported (Y/N): Y  
[2] MAX_POOL_2D, version: 1, supported (Y/N): Y  
[3] ADD, version: 1, supported (Y/N): Y  
[4] MEAN, version: 1, supported (Y/N): Y  
[5] FULLY_CONNECTED, version: 1, supported (Y/N): Y  
[6] SOFTMAX, version: 1, supported (Y/N): Y
```

- Parses the operator code in the model and checks if the OpResolver supports it
- If any operator code is not supported, an error is raised
 - Custom operators are handled in a more complicated way

3. Build Interpreter (3)

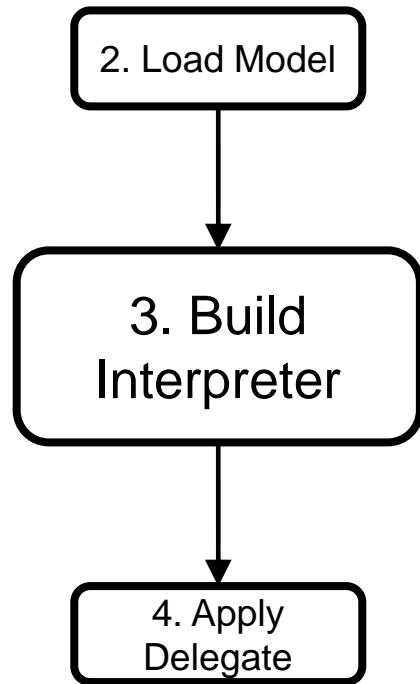


Number of subgraphs: 1

```
SubGraph [0] main
Total 187 tensors in SubGraph [0]
Tensor [0] serving_default_input_1:0, type = FLOAT32, shape = [1, 224, 224, 3], buffer = 1 (empty)
Tensor [1] resnet50/conv1_bn/FusedBatchNormV3, type = FLOAT32, shape = [64], buffer = 2 (has data, size = 256)
Tensor [2] resnet50/conv2_block1_0_bn/FusedBatchNormV3, type = FLOAT32, shape = [256], buffer = 3 (has data, size = 1024)
Tensor [3] resnet50/conv2_block1_1_bn/FusedBatchNormV3, type = FLOAT32, shape = [64], buffer = 4 (has data, size = 256)
Tensor [4] resnet50/conv2_block1_2_bn/FusedBatchNormV3, type = FLOAT32, shape = [64], buffer = 5 (has data, size = 256)
Tensor [5] resnet50/conv2_block1_3_bn/FusedBatchNormV3, type = FLOAT32, shape = [256], buffer = 6 (has data, size = 1024)
Tensor [6] resnet50/conv2_block2_1_bn/FusedBatchNormV3, type = FLOAT32, shape = [64], buffer = 7 (has data, size = 256)
Tensor [7] resnet50/conv2_block2_2_bn/FusedBatchNormV3, type = FLOAT32, shape = [64], buffer = 8 (has data, size = 256)
Tensor [8] resnet50/conv2_block2_3_bn/FusedBatchNormV3, type = FLOAT32, shape = [256], buffer = 9 (has data, size = 1024)
...
Tensor [184] resnet50/avg_pool/Mean, type = FLOAT32, shape = [1, 2048], buffer = 185 (empty)
Tensor [185] resnet50/predictions/MatMul;resnet50/predictions/BiasAdd, type = FLOAT32, shape = [1, 1000], buffer = 186 (empty)
Tensor [186] StatefulPartitionedCall:0, type = FLOAT32, shape = [1, 1000], buffer = 187 (empty)
```

- Parses the number of subgraphs in the model and creates the corresponding subgraph objects
- For each subgraph, tensor information is parsed and corresponding tensor objects are instantiated

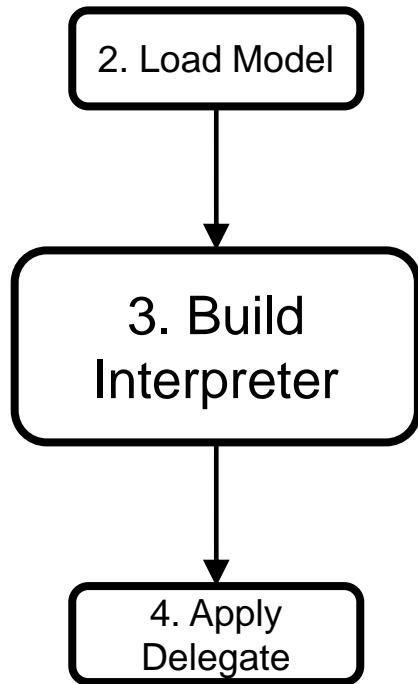
3. Build Interpreter (4)



```
Total 75 nodes in SubGraph [0]
Node [0]: PAD
Input tensors: 0 110
Output tensors: 112
Node [1]: CONV_2D
Input tensors: 112 54 1
Output tensors: 113
Node [2]: PAD
Input tensors: 113 108
Output tensors: 114
Node [3]: MAX_POOL_2D
Input tensors: 114
Output tensors: 115
Node [4]: CONV_2D
Input tensors: 115 55 2
Output tensors: 116
...
Node [73]: FULLY_CONNECTED
Input tensors: 184 111 107
Output tensors: 185
Node [74]: SOFTMAX
Input tensors: 185
Output tensors: 186
```

- For each subgraph, node information is parsed and corresponding node objects are instantiated
- Execution plan is initialized in this step, matching the execution order to the node indices

3. Build Interpreter (5)

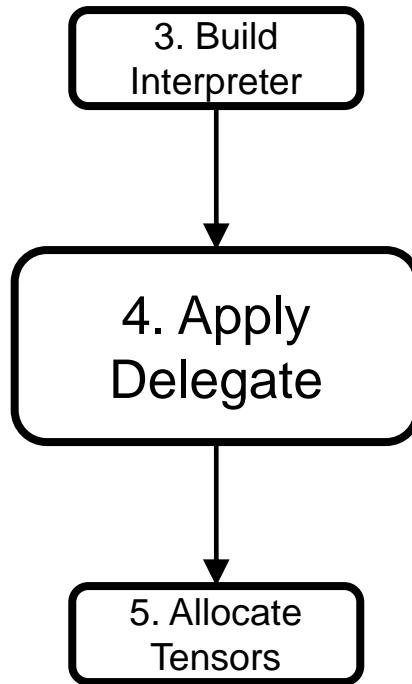


Number of subgraphs: 1
Number of nodes of subgraph 0: 75
Number of tensors in subgraph 0: 187
Execution plan size of subgraph 0: 75

- Checking whether if the interpreter is successfully created

Node 0: PAD	Node 31: CONV_2D	Node 61: CONV_2D
Node 1: CONV_2D	Node 32: CONV_2D	Node 62: CONV_2D
Node 2: PAD	Node 33: ADD	Node 63: ADD
Node 3: MAX_POOL_2D	Node 34: CONV_2D	Node 64: CONV_2D
Node 4: CONV_2D	Node 35: CONV_2D	Node 65: CONV_2D
Node 5: CONV_2D	Node 36: CONV_2D	Node 66: CONV_2D
Node 6: CONV_2D	Node 37: CONV_2D	Node 67: ADD
Node 7: CONV_2D	Node 38: ADD	Node 68: CONV_2D
Node 8: ADD	Node 39: CONV_2D	Node 69: CONV_2D
Node 9: CONV_2D	Node 40: CONV_2D	Node 70: CONV_2D
Node 10: CONV_2D	Node 41: CONV_2D	Node 71: ADD
Node 11: CONV_2D	Node 42: ADD	Node 72: MEAN
Node 12: ADD	Node 43: CONV_2D	Node 73: FULLY_CONNECTED
Node 13: CONV_2D	Node 44: CONV_2D	Node 74: SOFTMAX
Node 14: CONV_2D	Node 45: CONV_2D	
Node 15: CONV_2D	Node 46: ADD	
Node 16: ADD	Node 47: CONV_2D	
Node 17: CONV_2D	Node 48: CONV_2D	
Node 18: CONV_2D	Node 49: CONV_2D	
Node 19: CONV_2D	Node 50: ADD	
Node 20: CONV_2D	Node 51: CONV_2D	
Node 21: ADD	Node 52: CONV_2D	
Node 22: CONV_2D	Node 53: CONV_2D	
Node 23: CONV_2D	Node 54: ADD	
Node 24: CONV_2D	Node 55: CONV_2D	
Node 25: ADD	Node 56: CONV_2D	
Node 26: CONV_2D	Node 57: CONV_2D	
Node 27: CONV_2D	Node 58: ADD	
Node 28: CONV_2D	Node 59: CONV_2D	
Node 29: ADD	Node 60: CONV_2D	
Node 30: CONV_2D		

4. Apply Delegate

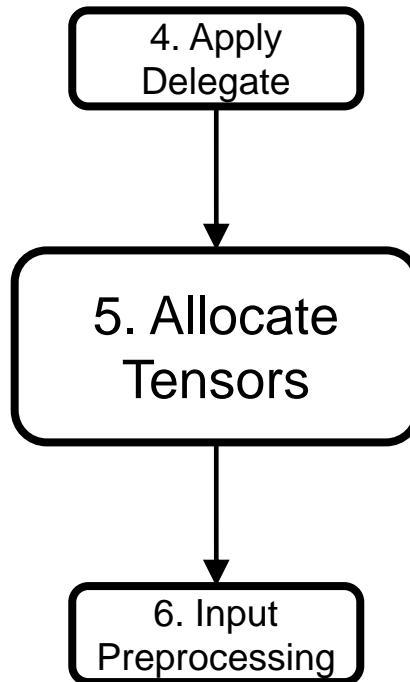


98

```
|     |     if (interpreter->ModifyGraphWithDelegate(xnn_delegate) == kTfLiteOk)
```

- The delegate traverses the **execution plan** and replaces supported nodes
 - Number of nodes of subgraph 0: 76
 - Node 75: DELEGATE
 - Execution plan size of subgraph 0: 1
 - Node 75: DELEGATE
 - Inputs:
 - 0 (type: FLOAT32, dims: [1, 224, 224, 3])
 - 1 (type: FLOAT32, dims: [64])
 - 2 (type: FLOAT32, dims: [256])
 - ...
 - 110 (type: INT32, dims: [4, 2])
 - 111 (type: FLOAT32, dims: [1000, 2048])
 - Outputs:
 - 186 (type: FLOAT32, dims: [1, 1000])
 - Total 113 tensors are used.
- As a result, a new node is created in the subgraph, which indicates the interpreter to call the delegate
- Execution plan is modified
- Tensors for intermediate results are no longer used
- The delegate manages intermediate results by itself

5. Allocate Tensors



```

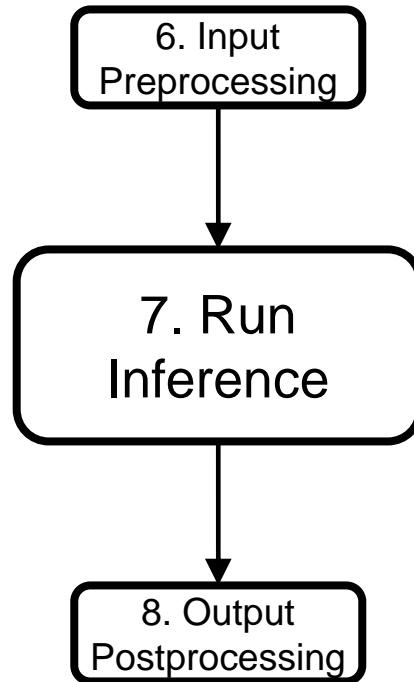
107     /* Allocate Tensors */
108     if (interpreter->AllocateTensors() != kTfLiteOk)
109     {
110         std::cerr << "Failed to Allocate Tensors" << std::endl;
111         return 1;
112     }

===== Before Allocate Tensors =====
Tensor 0: serving_default_input_1:0 | AllocType: kTfLiteArenaRw | Shape: [1, 224, 224, 3] | Bytes: 602112 | Address: 0
Tensor 1: resnet50/conv1_bn/FusedBatchNormV3 | AllocType: kTfLiteMmapRo | Shape: [64] | Bytes: 256 | Address: 0x7f70764bd0
Tensor 2: resnet50/conv2 block1 0 bn/FusedBatchNormV3 | AllocType: kTfLiteMmapRo | Shape: [256] | Bytes: 1024 | Address: 0x7f707647c4
...
Tensor 111: resnet50/predictions/MatMul | AllocType: kTfLiteMmapRo | Shape: [1000, 2048] | Bytes: 8192000 | Address: 0x7f6a6005c0
Tensor 186: StatefulPartitionedCall:0 | AllocType: kTfLiteArenaRw | Shape: [1, 1000] | Bytes: 4000 | Address: 0

===== After Allocate Tensors =====
Tensor 0: serving_default_input_1:0 | AllocType: kTfLiteArenaRw | Shape: [1, 224, 224, 3] | Bytes: 602112 | Address: 0x7f61637040
Tensor 1: resnet50/conv1_bn/FusedBatchNormV3 | AllocType: kTfLiteMmapRo | Shape: [64] | Bytes: 256 | Address: 0x7f70764bd0
Tensor 2: resnet50/conv2 block1 0 bn/FusedBatchNormV3 | AllocType: kTfLiteMmapRo | Shape: [256] | Bytes: 1024 | Address: 0x7f707647c4
...
Tensor 111: resnet50/predictions/MatMul | AllocType: kTfLiteMmapRo | Shape: [1000, 2048] | Bytes: 8192000 | Address: 0x7f6a6005c0
Tensor 186: StatefulPartitionedCall:0 | AllocType: kTfLiteArenaRw | Shape: [1, 1000] | Bytes: 4000 | Address: 0x7f616ca040
  
```

• Only mutable tensors are affected

7. Run Inference



```
147     /* Inference */
148     if (interpreter->Invoke() != kTfLiteOk)
149     {
150         std::cerr << "Failed to invoke the interpreter" << std::endl;
151         return 1;
152     }
```

0: Node 75 -> TfLiteXNNPackDelegate

- The interpreter traverses through the execution plan and executes the nodes in it
- The internal execution logic of each delegate varies

Tutorial Outline

Lecture 1: *Exercise Overview and Setup* (1 h 30 min)

Lecturer Seongsoo Hong

Topics Motivating Example

Development Environment Setup

Coffee Break (30 min)

Lecture 2: *From Inference Driver to Inference Runtime* (1 h 30 min)

Lecturer Seongsoo Hong and Namcheol Lee

Topics Step-by-Step Inference Driver Walkthrough

Internals of LiteRT

Lunch Break (1 h)

Lecture 3: *Model Slicer* (1 h 30 min)

Lecturer Seongsoo Hong

Topic Model Slicer: Slicing and Conversion Tool for LiteRT

Coffee Break (30 min)

Lecture 4: *Throughput Enhancement on Heterogeneous Accelerators* (1 h 30 min)

Lecturer Namcheol Lee

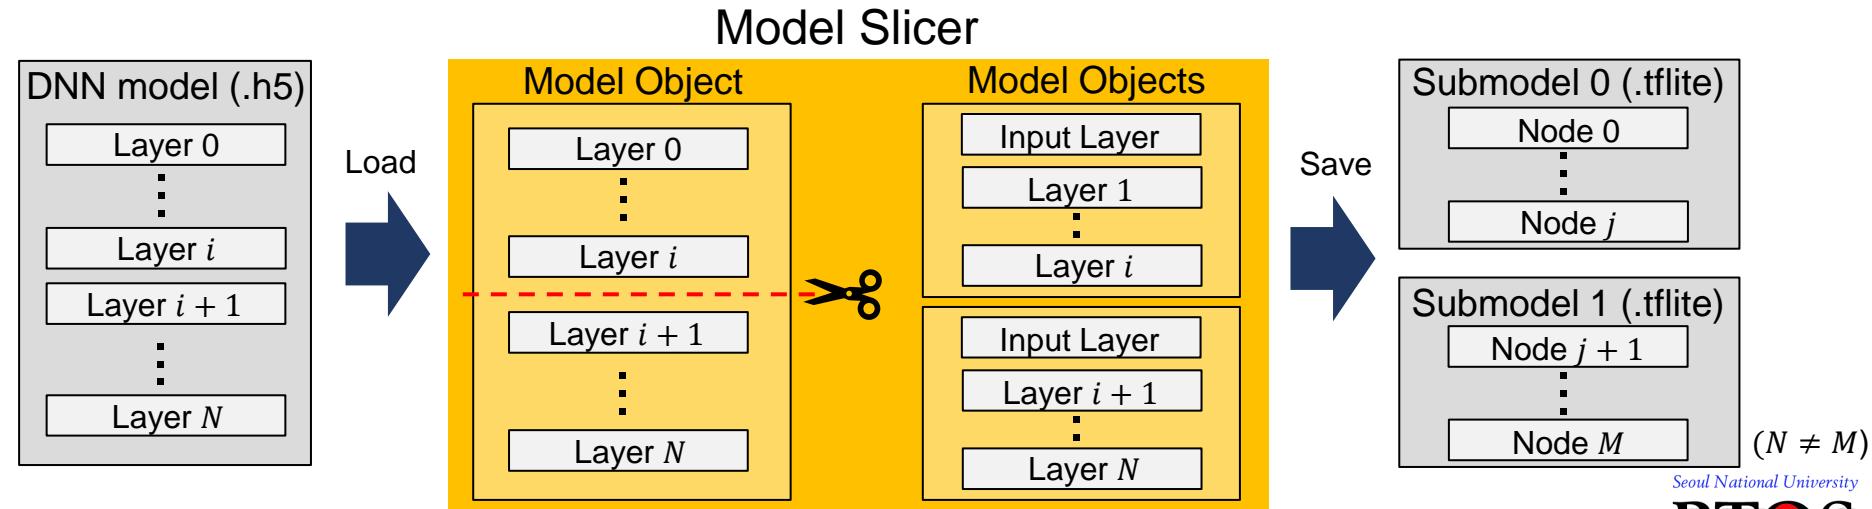
Topic Implementing a Pipelined Inference Driver for Heterogeneous Processors

Contents

-
- I. **Slicing TensorFlow Model in HDF5: Big Picture**
 - II. Prior Knowledge to Understand Model Slicer
 - III. The Slicing Algorithm
 - IV. Step-by-Step Model Slicer Walkthrough
 - V. Hands-On Exercise: Run Model Slicer

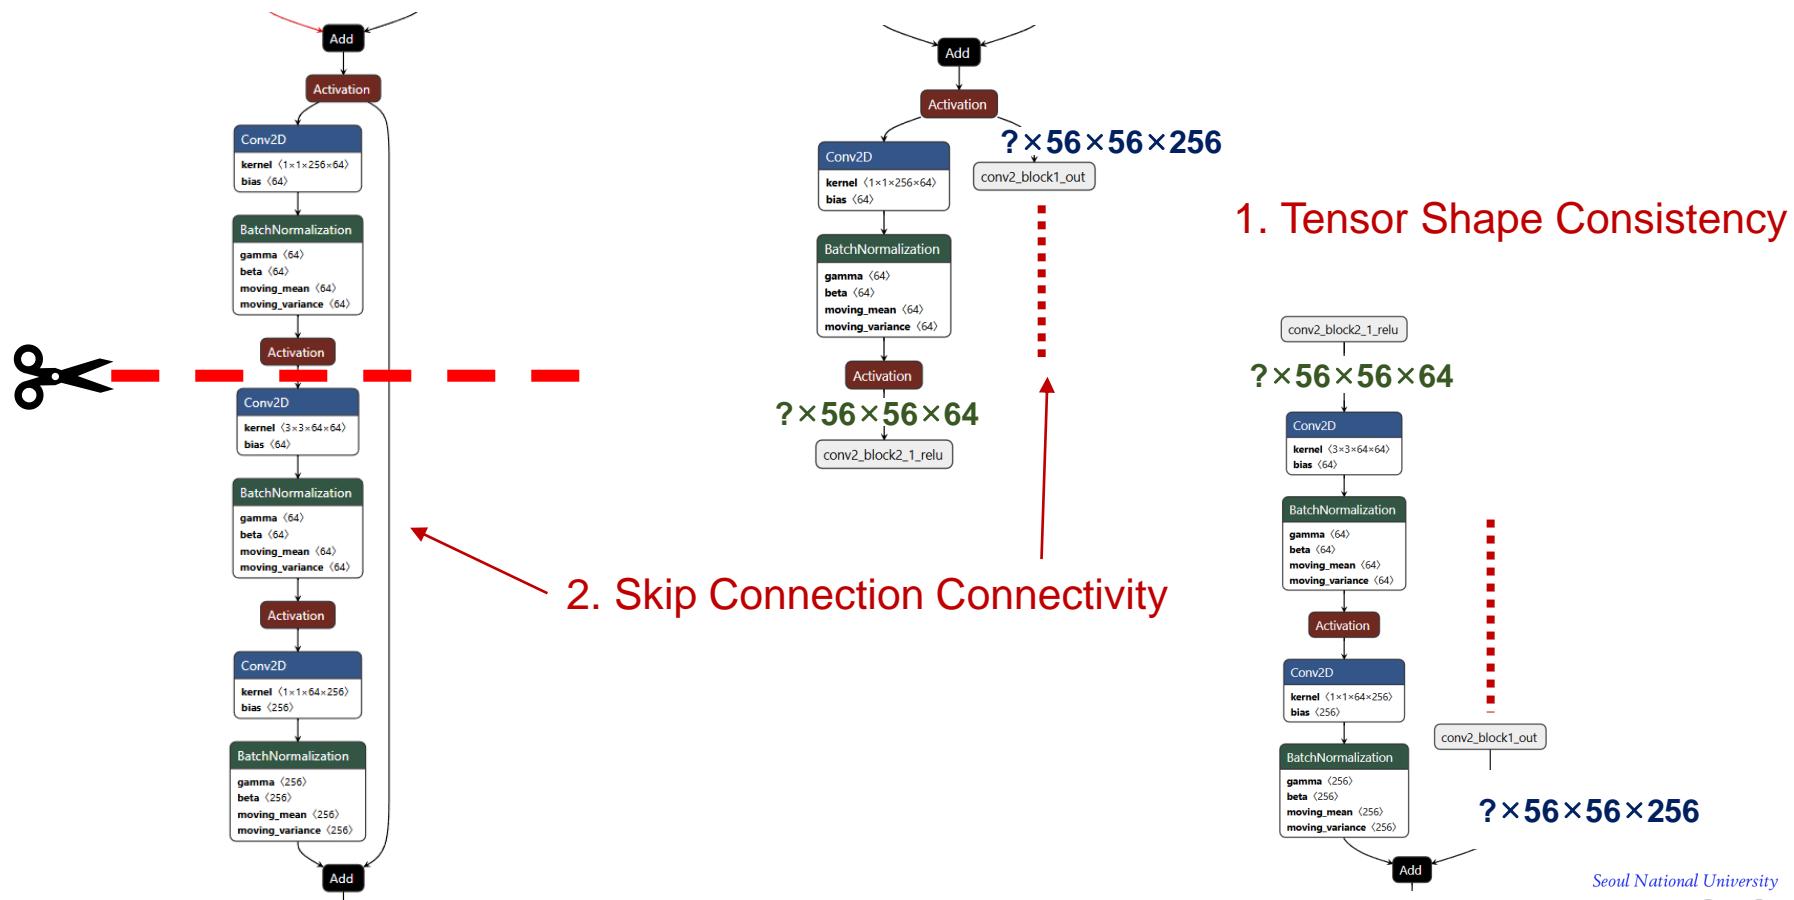
Model Slicing for DNN Pipelining (1)

- ❖ Slice a DNN model into multiple submodels
 - Accept a TensorFlow model in HDF5 (Hierarchical Data Format version 5) (.h5)
 - .h5 is a “*serialized storage format on disk*”
 - Load the model in .h5 format to create a “*model object in memory*”
 - Construct submodels from the model object by hand-crafting their model objects
 - Generate multiple LiteRT models in .tflite format



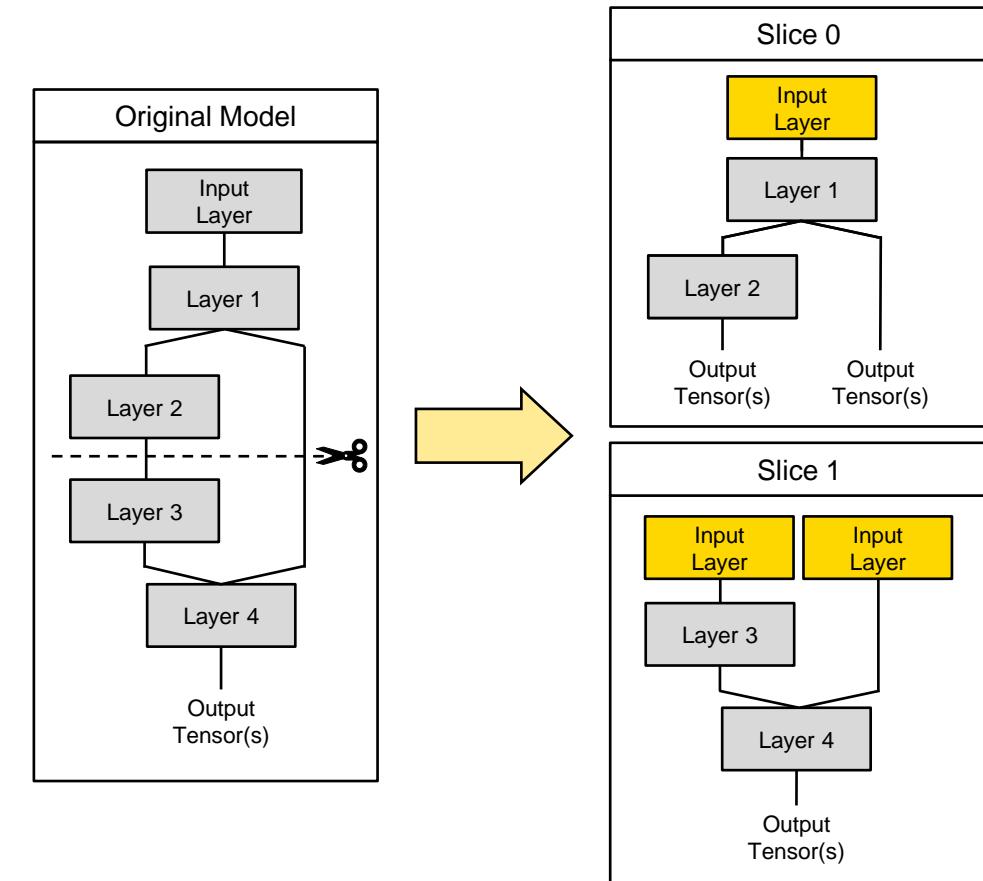
Model Slicing for DNN Pipelining (2)

- ## ❖ Characteristics to be preserved during slicing



Model Slicing for DNN Pipelining (3)

- ❖ Input layers are newly created
 - Per output (tensor(s)) of the previous slice

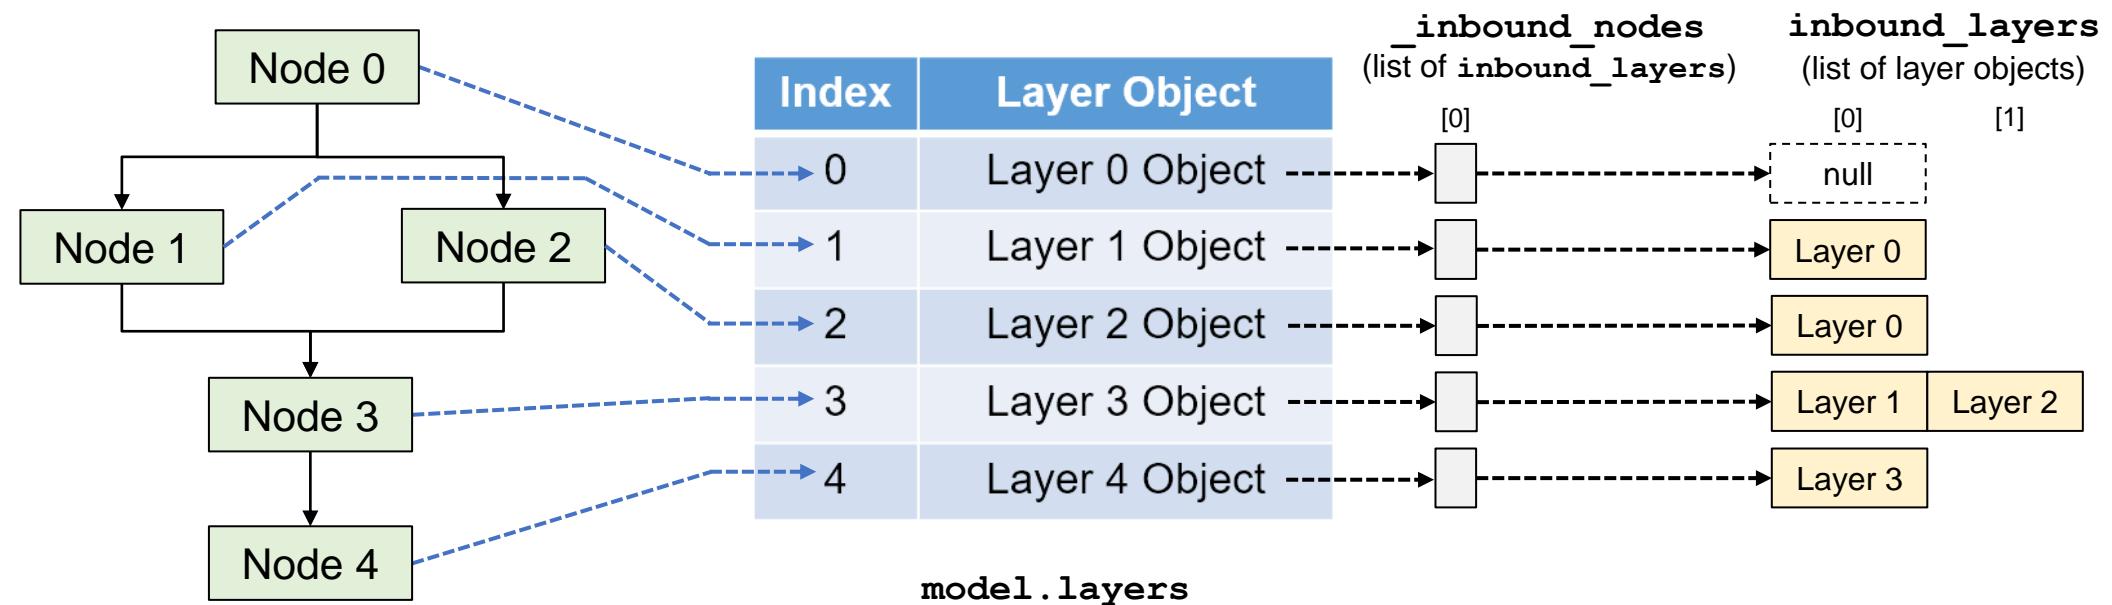


Contents

- I. Slicing TensorFlow Model in HDF5: Big Picture
- II. Prior Knowledge to Understand Model Slicer**
- III. The Slicing Algorithm
- IV. Step-by-Step Model Slicer Walkthrough
- V. Hands-On Exercise: Run Model Slicer

1. Internal Representation of Keras Model Object (1)

- ❖ Computational graph (DAG) defined by “*layer objects*” and “*node objects*”
 - “`model.layers`” works as a hashing table for the DAG



1. Internal Representation of Keras Model Object (2)

- ❖ Important to understand differences between “*layers*” and “*nodes*”
 - Layer objects: similar to *function definitions*
 - Static computational building blocks (units of computation)
 - “`model.layers`” works as a hashing table for the DAG
 - Node objects: similar to *function invocations*
 - Instances of layers in DAG
 - Created dynamically when a node needs to appear in the DAG
 - By calling the `__call__` method of the layer (functor `layer()`)
 - “`model.layers[idx]._inbound_nodes[0]`” is the node created first
 - *One-to-one correspondence between layers and nodes in most Keras models*
 - For the sake of simplicity

1. Internal Representation of Keras Model Object (3)

❖ Keras tensors

- A built-in tensor object that has information for graph connections
- Attributes: **shape**, **dtype**, and **name** properties

```
KerasTensor(type_spec=TensorSpec(shape=(None, 7, 7, 512), dtype=tf.float32, name=None),  
           name='conv5_block3_2_conv/BiasAdd:0', description="created by layer 'conv5_block3_2_conv'")
```

`conv5_block3_2_conv` : Layer name that produced this KerasTensor

`BiasAdd` : TensorFlow op

`0` : First output tensor

- Hidden property “_keras_history”

```
KerasHistory(layer=<keras.layers.convolutional.conv2d.Conv2D object at 0x7f291f63e0>, node_index=1, tensor_index=0)
```

- **layer**: layer whose functor is invoked to generate the tensor(s)
- **node_index**: index of the functor invocation that generate the tensor(s)
- **tensor_index**: index of the tensor in the list of the generated tensors(s)

2. How the Model Slicer Works

❖ Key Idea

1. Creates in-memory representation (a.k.a. model object) by loading .h5 file

For each slice, handles layers one by one

2. Accepts input tensors and calls `layer()` to
 - Create a node object
 - Make inbound connections
 - Execute the layer's operator
 - Create output tensors
 - Make outbound connections
3. Process inside-ending skips and outside-ending skips

2-1. Create Model Object (1)

- ❖ Load HDF5 file

```
import tensorflow as tf
from tensorflow.keras.models import load_model

model = load_model(args.model_path, compile=False)
```

2-1. Create Model Object (2)

- ❖ “`model.layers`” works as a hashing table for the DAG of the model object
 - Correspondence between layers and nodes is represented by
 - `_inbound_nodes` (attribute of layer)
 - Each time a layer is connected to a new input, a node is added to `layer._inbound_nodes`
 - `inbound_layers` (attribute of node)
 - References to the previous layers connected through the inbound nodes
 - `_outbound_nodes` (attribute of layer)
 - Each time the output of a layer is used by another layer, a node is added to `layer._outbound_nodes`
 - `outbound_layer` (attribute of node)
 - References to the next layers connected through the outbound nodes

2-2. Call the `layer()` Functor (1)

- ❖ Understanding the `layer()` functor

- `tensors_from_layer = layer(tensors_to_layer)`
 - Input parameter: Keras tensors that are inputs to the layer being created
 - Return value: Keras tensors that are outputs from the layer being created
 - The `__call__` method of the `layer` object is invoked

2-2. Call the `layer()` Functor (2)

❖ In `layer()`

1. Reads input tensor's Keras history to identify its origin layer
2. Creates a new node
 - Set `Node.outbound_layer` = current layer
 - Set `Node.inbound_layers` = origin layers of input tensors
3. Record `Node.input_tensors` and `Node.output_tensors`
4. Append the node to
 - `inbound_nodes` list of the current layer
 - `outbound_nodes` list of each input tensor's origin layer

Contents

- I. Slicing TensorFlow Model in HDF5: Big Picture
- II. Prior Knowledge to Understand Model Slicer
- III. The Slicing Algorithm**
- IV. Step-by-Step Model Slicer Walkthrough
- V. Hands-On Exercise: Run Model Slicer

Constructing Algorithm by Induction

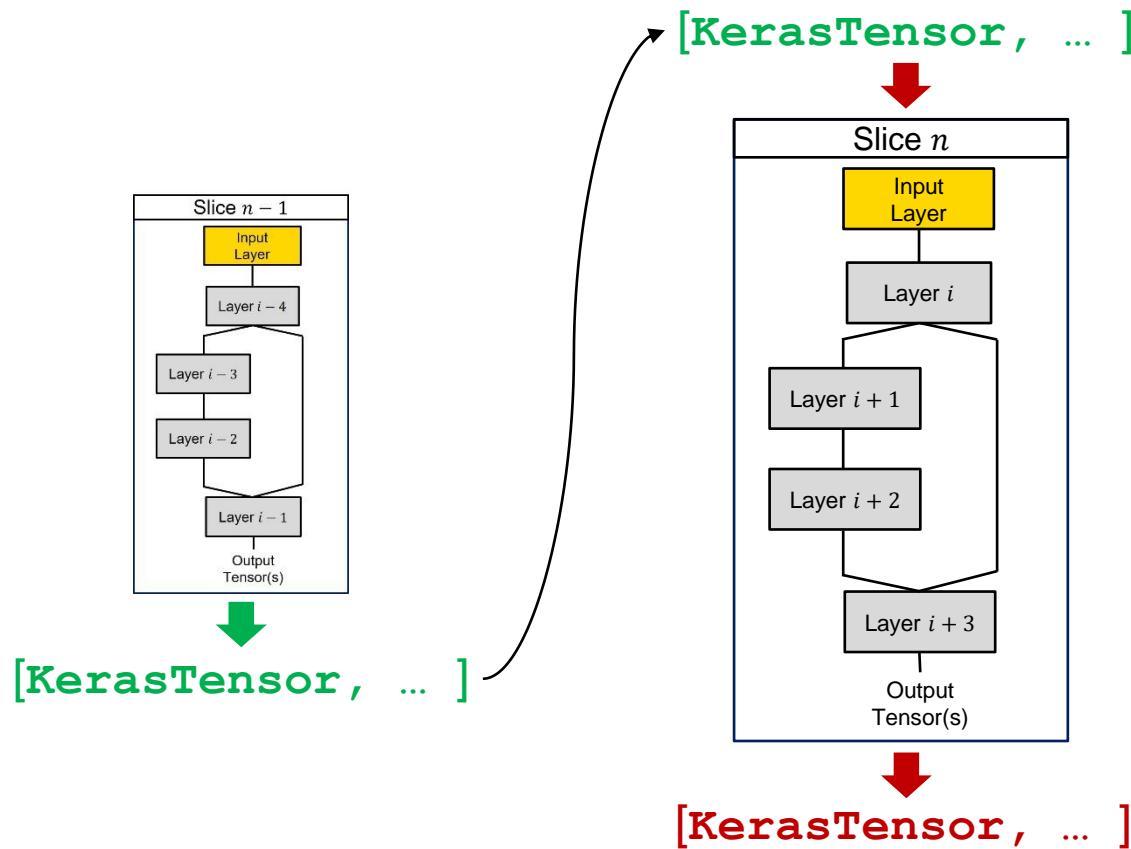
- ❖ A slice's *output tensors* are fed to the next slice as *input*
 1. We correctly hand-craft the input tensors to the first slice
 2. We generate a slice in the middle correctly
 - Its internal structure matches the corresponding part of the model
 - It produces the correct number and shape of output tensors
 3. We repeatedly generate slices from the first to the last

7 Cases for the Model Slicer

- ❖ For a given slice, perform
 1. Input layer generation
 2. Input tensor initialization of start layer
 3. Inside-ending skip connection initialization
 4. Outside-ending skip connection initialization
 5. Per-layer processing
 6. Inside-ending skip connection update
 7. Outside-ending skip connection update

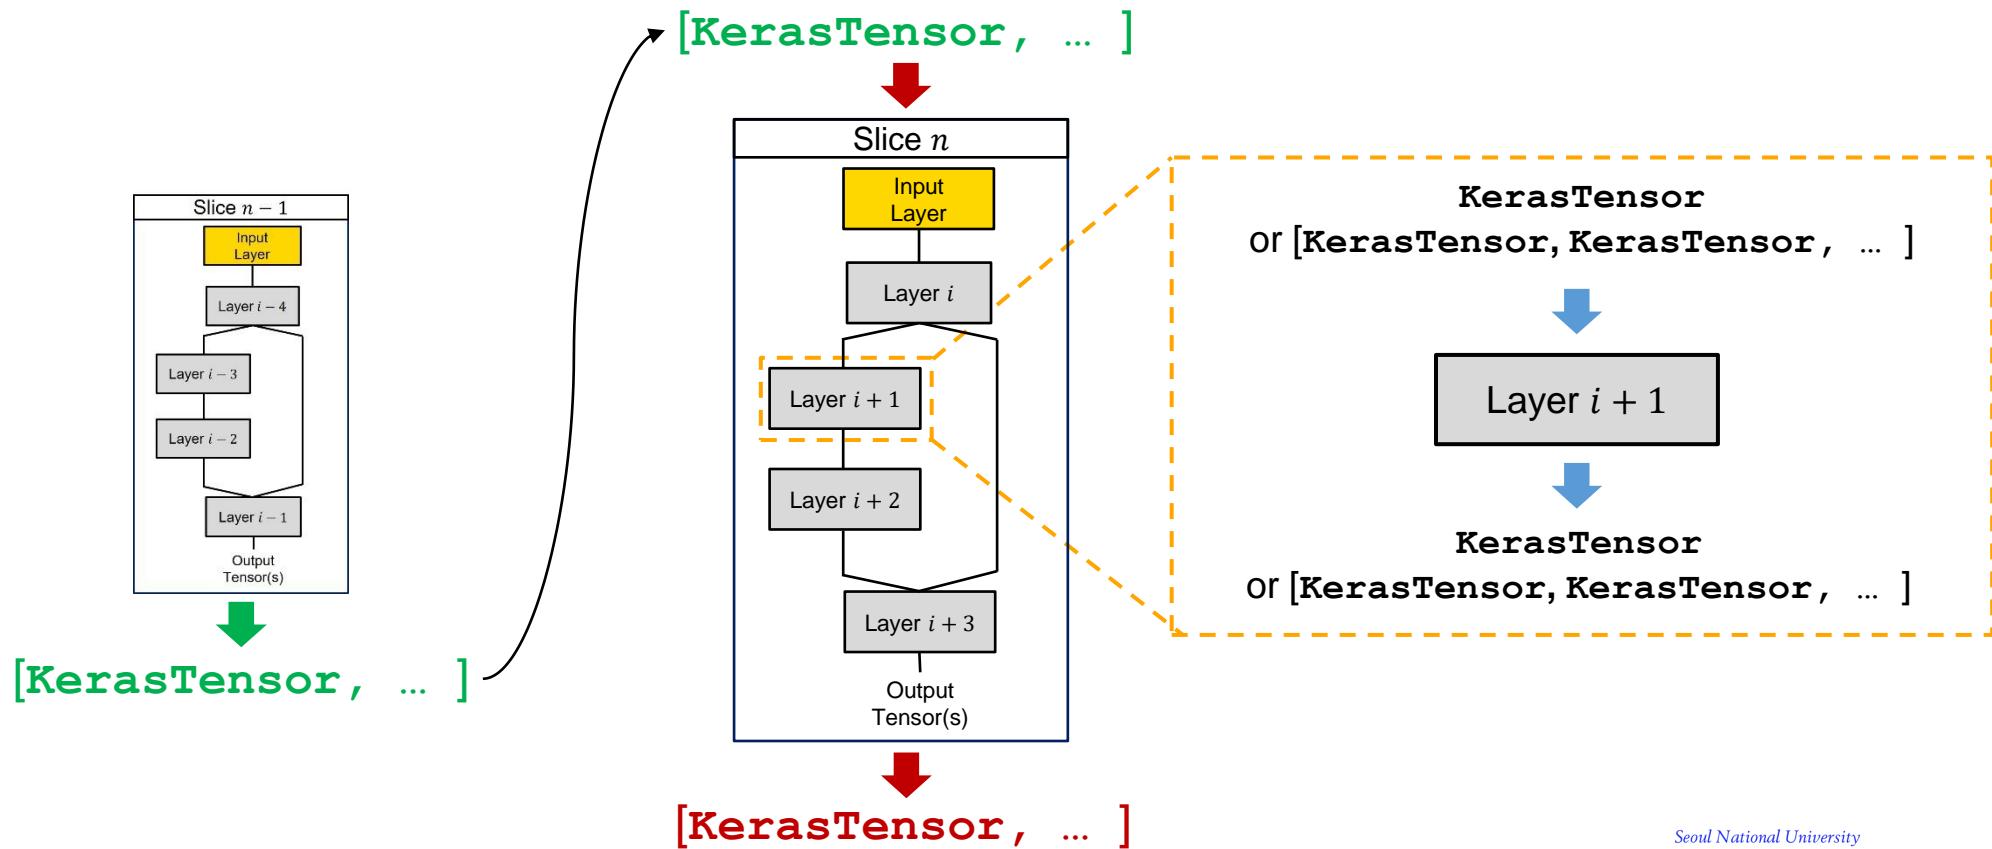
Priori Knowledge: 1. Input and Output of Slice

- ❖ A list containing one or more Keras tensors



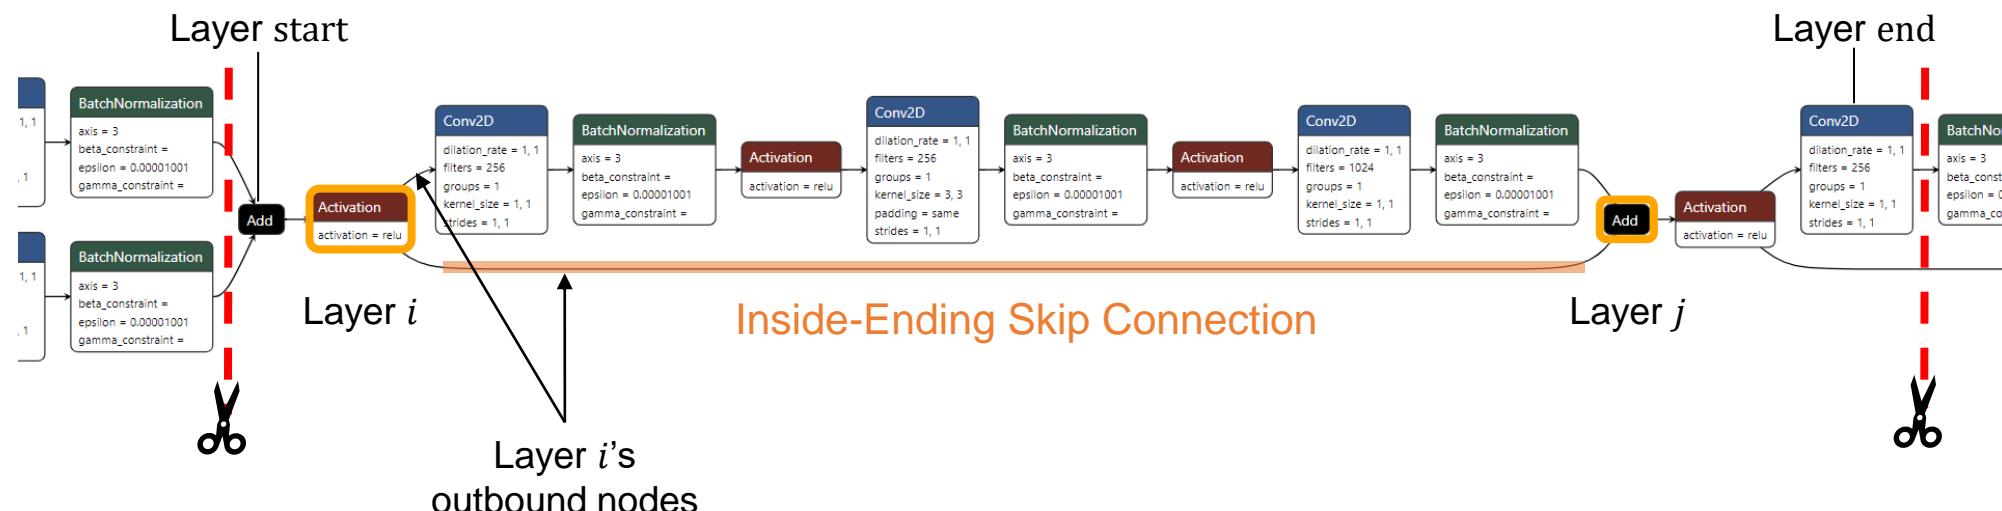
Priori Knowledge: 2. Input and Output of Layer

- Either a single Keras tensor or a list containing multiple Keras tensors



Priori Knowledge: 3. Inside-Ending Skip Connections

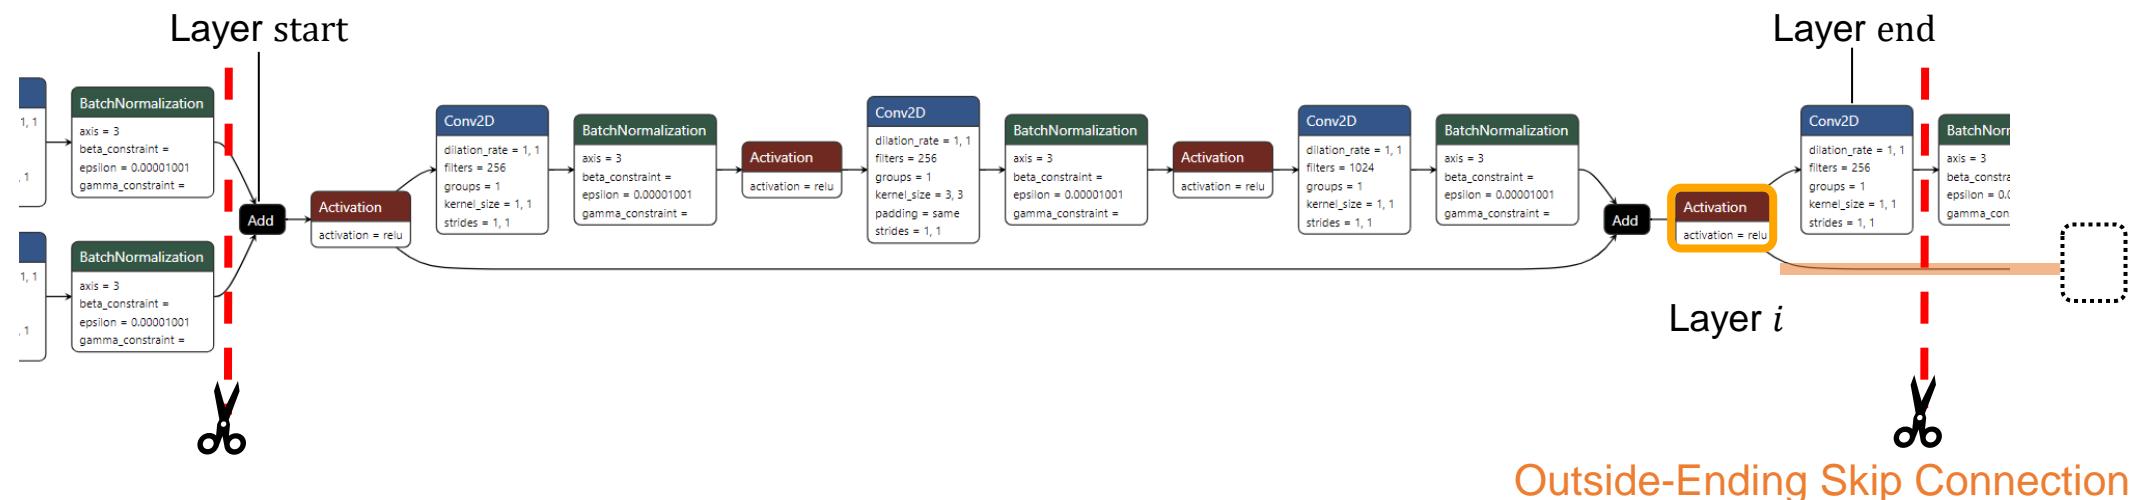
- ❖ Inside-ending skip connection of a slice with range of $[start, end]$ is
 - Among the *outbound nodes of newly created input layers*, a node corresponding to a layer with index j such that $j > start$ and $j \leq end$
 - And among the *outbound nodes of layer i* where $i \in [start, end)$, a node corresponding to a layer with index j such that $j > i + 1$ and $j \leq end$



Priori Knowledge:

4. Outside-Ending Skip Connections

- ❖ Outside-ending skip connection of a slice with range of [start, end] is
 - Among the *outbound nodes of newly created input layers*, a node corresponding to a layer with index j such that $j > \text{end}$
 - And among the *outbound nodes of layer i* where $i \in [\text{start}, \text{end}]$, a node corresponding to a layer with index j such that $j > \text{end}$



Case 1. Input Layer Generation

- ❖ Create input layers to the slice from the provided input tensors
 - `input_layers = {} # dictionary {layer name, KerasTensor}`
`input_layers[name] =`
 `tf.keras.layers.Input(shape=tensor.shape[1:], name=name)`
 - `name`: Name of the layer that generated the tensor
 - `tensor.shape[1:]`: Height, width, and channel of `tensor` as a list

Case 2. Input Tensor Initialization of Start Layer

- ❖ Identify the outbound nodes of the input layers
- ❖ Among those, select the nodes that target the start layer
- ❖ Assemble their Keras tensors into the start layer's input
 - `tensors_to_start_layer = []`
 - `tensors_to_start_layer.append(input_layers[name])`

Case 3. Inside-Ending Skip Connection Initialization

- ❖ Identify the outbound nodes of the input layers
- ❖ Among those, select the nodes that target the layers in the same slice, other than the start layer
- ❖ Obtain Keras tensors and register them as inside-ending skip connections
 - `inside_ending_skips = {}
dictionary {layer name, KerasTensor}`
 - `inside_ending_skips[name] = input_layers[name]`

Case 4. Outside-Ending Skip Connection Initialization

- ❖ Identify the outbound nodes of the input layers
- ❖ Among those, select the nodes that target the layers outside the current slice
- ❖ Obtain Keras tensors and register them as outside-ending skip connections
 - `outside_ending_skips = {}
dictionary {layer name, KerasTensor}`
 - `outside_ending_skips[name] = input_layers[name]`

Case 5. Per-Layer Processing

- ❖ Layer i ($i \in [\text{start}, \text{end}]$)
 - Call the layer's functor with
 - $i = \text{start}$: Tensor(s) initialized in Case 2
 - $i > \text{start}$: Tensor(s) from layer $i - 1$ and any required inside-ending skip tensors
 - Identify required tensors by inspecting the inbound nodes of layer i in the model

Case 6. Inside-Ending Skip Connection Update

- ❖ After processing layer i ($i \in [\text{start}, \text{end}]$), find its outbound nodes
- ❖ If an outbound node targets a layer in this slice (target $\in (i + 1, \text{end}]$), register the output tensor as an inside-ending skip connection
 - `inside_endings_skips[layer.name] = tensors_from_layer`

Case 7. Outside-Ending Skip Connection Update

- ❖ After processing layer i ($i \in [\text{start}, \text{end}]$), find its outbound nodes
- ❖ If an outbound node targets a layer outside the current slice ($\text{target} > \text{end}$), register the output tensor as an outside-ending skip connection
 - `outside_ending_skips[layer.name] = tensors_from_layer`

Code for Cases 1, 2, 3, and 4

```
28 v def slice_dnn(model, start, end, input_tensors):
51     # Case 1, 2, 3, and 4
52     for tensor in input_tensors:
53         # Input layers are created
54         name = tensor._keras_history[0].name
55         input_layers[name] = tf.keras.layers.Input(shape=tensor.shape[1:], name=name)
56
57         # Inspect an input layer's outbound nodes to see where the model consumes it
58         origin_layer = model.get_layer(name)
59         if len(origin_layer._outbound_nodes) > 1: # If an input layer has multiple outbound layers, its output feeds multiple layers
60             for origin_outbound_node in origin_layer._outbound_nodes:
61                 origin_outbound_layer = origin_outbound_node.outbound_layer
62                 target_idx = model.layers.index(origin_outbound_layer)
63                 if target_idx == start:
64                     tensors_to_start_layer.append(input_layers[name])
65                 elif target_idx <= end and target_idx > start:
66                     inside_ending_skips[name] = input_layers[name]
67                 elif target_idx > end:
68                     # Any outside-ending skip connections that are made here are not used in this slice
69                     outside_ending_skips[name] = input_layers[name]
70             else:
71                 origin_outbound_layer = origin_layer._outbound_nodes[0].outbound_layer
72                 target_idx = model.layers.index(origin_outbound_layer)
73                 if target_idx == start:
74                     tensors_to_start_layer.append(input_layers[name])
75                 elif target_idx <= end and target_idx > start:
76                     inside_ending_skips[name] = input_layers[name]
77                 elif target_idx > end:
78                     # Any outside-ending skip connections that are made here are not used in this slice
79                     outside_ending_skips[name] = input_layers[name]
```

Code for Cases 5, 6, and 7 (1)

```
28 ✓ def slice_dnn(model, start, end, input_tensors):
  1 # Case 5, 6, and 7
  2 # Set the start layer's inputs from tensors_to_start_layer
  3 if(len(tensors_to_start_layer) == 1):
  4     tensors_to_layer = tensors_to_start_layer[0]
  5 else:
  6     tensors_to_layer = tensors_to_start_layer
  7
  8 for i in range(start, end+1):
  9     layer = model.layers[i]
 10     origin_inbound_layers = layer._inbound_nodes[0].inbound_layers
 11
 12     # Skip the layer if it is an InputLayer
 13     if isinstance(layer, tf.keras.layers.InputLayer):
 14         continue
 15
 16     # Build current (i-th) layer
 17     if isinstance(origin_inbound_layers, list): # When current layer expects multiple inputs (list of KerasTensors)
 18         if(i == start):
 19             tensors_from_layer = layer(tensors_to_layer)
 20         else:
 21             # Check if layer i's output tensor(s) is used in layer i+1
 22             if model.layers[i-1] in origin_inbound_layers:
 23                 tensors_to_layer = [tensors_to_layer] if not isinstance(tensors_to_layer, list) else tensors_to_layer
 24             else:
 25                 tensors_to_layer = []
```

Code for Cases 5, 6, and 7 (2)

```
107     # From inbound layers, collect the required inside-ending skip tensors
108     for origin_inbound_layer in origin_inbound_layers:
109         inside Ending_skip = model.get_layer(origin_inbound_layer.name)
110         if inside Ending_skip.name not in [t.name.split('/')[0] for t in tensors_to_layer]:
111             tensors_to_layer.append(inside Ending_skips[inside Ending_skip.name])
112
113     # Call the functor of the current layer to build a new layer
114     try:
115         tensors_from_layer = \
116             layer(tensors_to_layer)
117     except: # When a custom layer's call signature deviates from Keras expectations
118         raise ValueError(f"Failed to call layer {layer.name} with tensors {tensors_to_layer}. "
119                         | "Please check the layer's call function and the input tensors.")
120     else: # When current layer expects a single input (KerasTensor)
121         if origin_inbound_layers.name in inside Ending_skips:
122             tensors_from_layer = layer(inside Ending_skips[origin_inbound_layers.name])
123         else:
124             tensors_from_layer = layer(tensors_to_layer)
125
126     # Update inside Ending_skips and outside Ending_skips based on the current layer's outbound nodes
127     if i < end: # Ensure the current layer is not the slice's last layer
128         for origin_outbound_node in layer._outbound_nodes:
129             skip_target_idx = model.layers.index(origin_outbound_node.outbound_layer)
130             if skip_target_idx <= end and skip_target_idx > i + 1:
131                 inside Ending_skips[layer.name] = tensors_from_layer
132             elif skip_target_idx > end:
133                 outside Ending_skips[layer.name] = tensors_from_layer
134
135     tensors_to_layer = tensors_from_layer # End of for i in range(start, end+1)
```

Contents

- I. Slicing TensorFlow Model in HDF5: Big Picture
- II. Prior Knowledge to Understand Model Slicer
- III. The Slicing Algorithm
- IV. Step-by-Step Model Slicer Walkthrough**
- V. Hands-On Exercise: Run Model Slicer

The Model Slicer

- ❖ A Python script tool
 - Developed at RTOSLab. of Seoul National University
- ❖ Derived from the original implementation of **DNNPipe**

*DNNPipe: Dynamic Programming-based
Optimal DNN Partitioning for Pipelined
Inference on IoT Networks*

Woobean Seo¹⁾, Saehwa Kim²⁾, and Seongsoo Hong¹⁾

¹⁾Department of Electrical and Computer Engineering, Seoul National University, Seoul 08826, Korea
²⁾Department of Information Communications Engineering, Hankuk University of Foreign Studies, Yongin 17035, Korea

ABSTRACT— Pipeline parallelization is an effective technique that enables the efficient execution of deep neural network (DNN) inference on resource-constrained IoT devices. To enable pipeline parallelization across computing nodes with asymmetric performance profiles, interconnected via low-latency, high-bandwidth networks, we propose DNNPipe, a DNN partitioning algorithm that constructs a pipeline plan for a given DNN. The primary objective of DNNPipe is to minimize the total latency of the pipeline plan. To achieve this, DNNPipe uses dynamic programming, which is repeatedly executed online in dynamically changing IoT environments. To achieve this, DNNPipe uses dynamic programming (DP) with pruning techniques that preserve optimality to explore the search space and find the optimal solution. In addition, DNNPipe uses upper-bound-based pruning to reduce the search space. Specifically, it aggressively prunes suboptimal pipeline plans using two pruning techniques: *upper-bound-based pruning* and *underutilized-stage pruning*. Our experimental results demonstrate that pipelined inference using an obtained optimal partitioning algorithm achieves up to 1.7 times lower latency compared to the original DNN inference. In addition, DNNPipe achieves up to 98.26% lower runtime overhead compared to PipeEdge, the fastest known optimal DNN partitioning algorithm.

KEYWORDS: IoT system, embedded AI, DNN partitioning, pipeline parallelization, inference, throughput optimization.

1. INTRODUCTION

The proliferation of Internet of Things (IoT) technology and the adoption of AI in various IoT applications [2, 3] has led to an increased demand for real-time processing of streaming data on deep neural networks (DNNs). However, the limited memory size of IoT devices often poses a challenge to the feasibility of executing even simple AI models. Even when a model can be loaded, the low processing power of such devices results in limited throughput, which causes a bottleneck in DNN inference. In addition, the limited bandwidth between the device and the network causes input data accumulates in a queue, leading to increasing queuing delays and eventually unbounded end-to-end (E2E) latency. A common workaround is to reduce the input data to prevent system overload, but this leads to data loss, which is unacceptable for most applications that require full-datum inference.

To address these challenges, pipeline parallelization has emerged as a promising solution. It not only reduces per-device memory requirements, but also enables the efficient execution of DNN inference on resource-constrained IoT devices.

Pipeline parallelization relies on DNN partitioning, which divides a DNN model into multiple stages, each consisting of a non-empty, contiguous subset of layers, and distributes their workload across computing nodes with asymmetric performance profiles. It can accelerate DNN inference through parallel and pipelined execution.

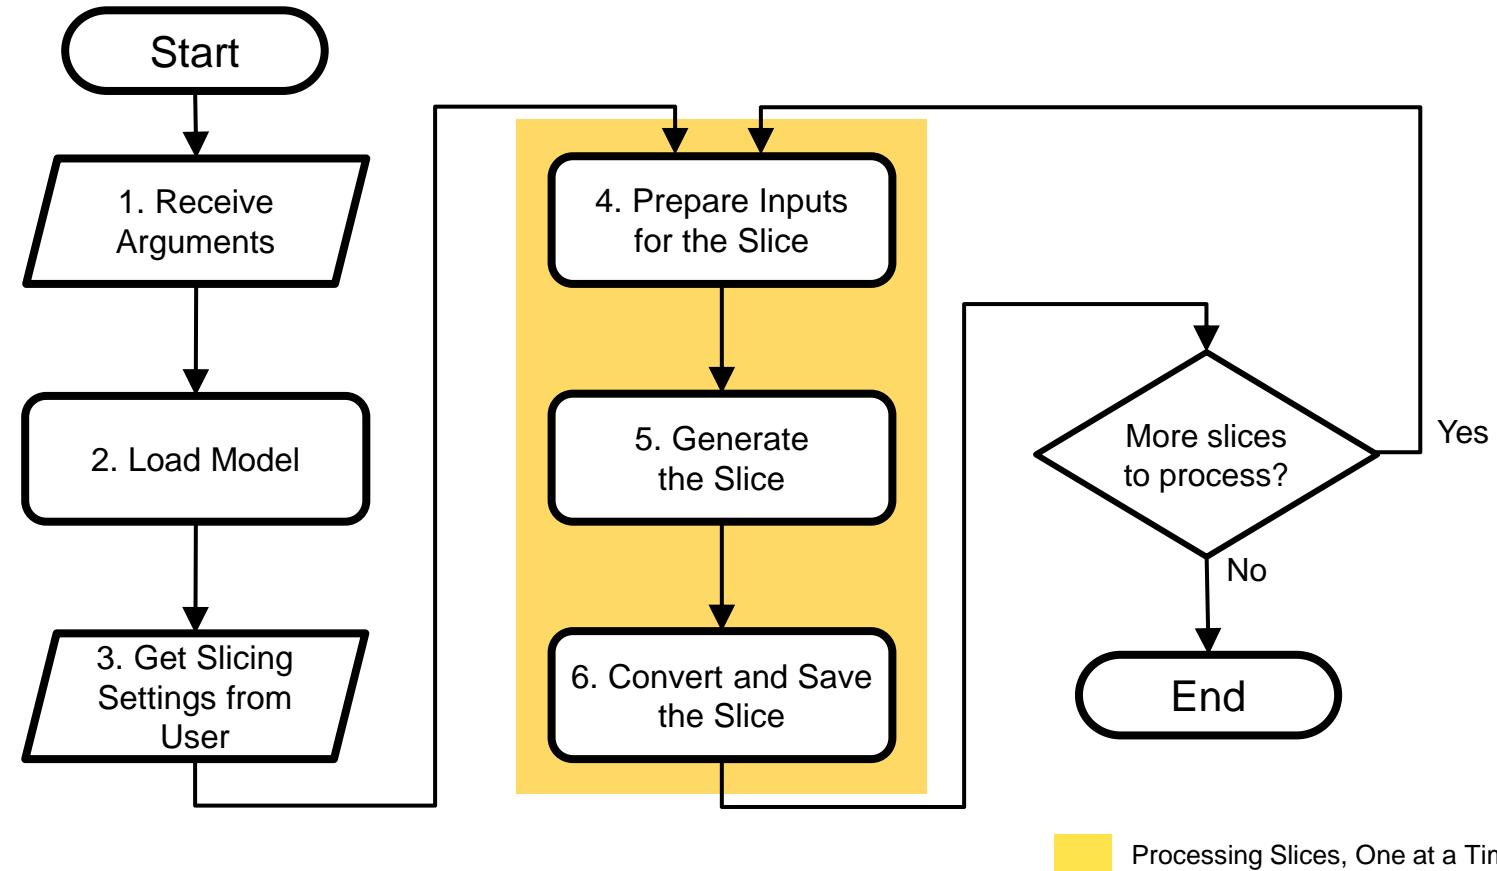
¹⁾Corresponding author.

E-mail address: kseob@kufs.ac.kr (S. Kim).

This paper is an extended version of the paper that appeared in the Ph.D. Symposium at the 2024 IEEE 44th International Conference on Distributed Computing Systems (ICDCS 2024) [1].

*Woobean Seo, Saehwa Kim, and Seongsoo Hong,
<https://doi.org/10.1016/j.sysarc.2025.103462>. "DNNPipe:
Dynamic Programming-based Optimal DNN
Partitioning for Pipelined Inference on IoT
Networks," Journal of Systems Architecture, vol.
166, 2025, 103462, ISSN 1383–7621,*

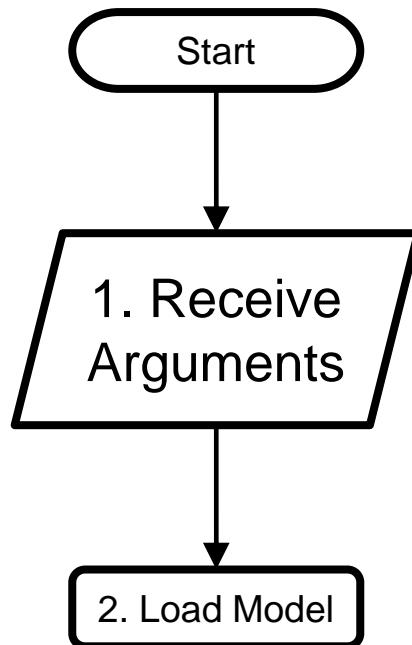
Model Slicer Code Flow



Processing Slices, One at a Time

1. Receive Arguments

- ❖ Parse command-line arguments for the model path and output directory

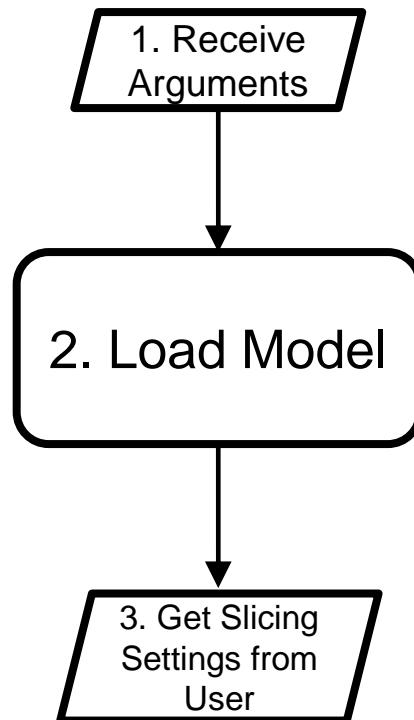


```
239 def main():  
240     # Parse command line arguments  
241     args = parse_arguments()
```

```
163 # Parse command line arguments  
164 def parse_arguments():  
165     parser = argparse.ArgumentParser()  
166     parser.add_argument('--model-path', type=str, required=True, help='Path to a model file (.h5)')  
167     parser.add_argument('--output-dir', type=str, default='./models')  
168     return parser.parse_args()
```

2. Load Model

- ❖ Load .h5 model file into memory as Keras model object
 - Slicing is performed on model object, not on .h5 model file

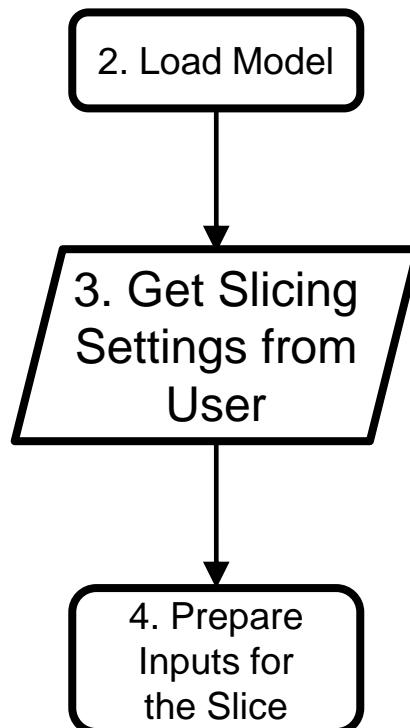


```
16 import tensorflow as tf
17 import numpy as np
18 import os
19 from tensorflow.keras.models import load_model
20 import argparse

243 # Load the model from the given path without compilation (for inference/slicing only)
244 model = load_model(args.model_path, compile=False)
```

3. Get Slicing Settings from User

- ❖ Get the number of submodels and the index of the first layer in each submodel

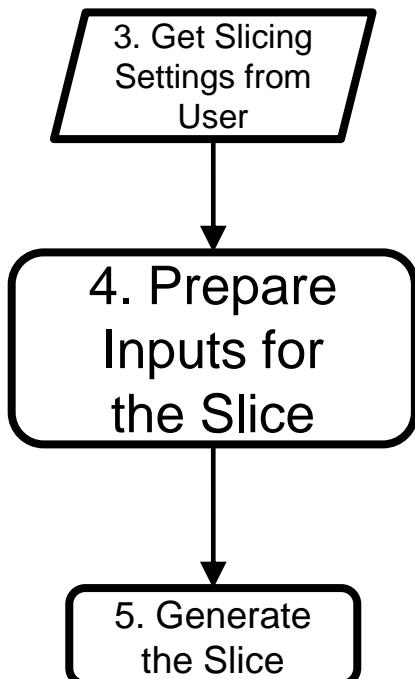


```
246     # Ask the user for the number of slices and the index of the last layer in each slice
247     num_layers = len(model.layers)
248     num_slices, starts = get_slice_starts(num_layers)
```

```
(.venv) rtoslab@RUBIKPi:~/DNNPipe-Tutorial$ python model_slicer.py --model-path ./models/resnet50.h5
How many submodels? 2
Generated submodels look like: (1, x1-1), (x1, 176)
Enter x1: 89
Layer index ranges for each submodel: [(1, 88), (89, 176)]
Saved LiteRT model to: ./models/submodel_0.tflite
Saved LiteRT model to: ./models/submodel_1.tflite
```

4. Prepare Inputs for the Slice (1)

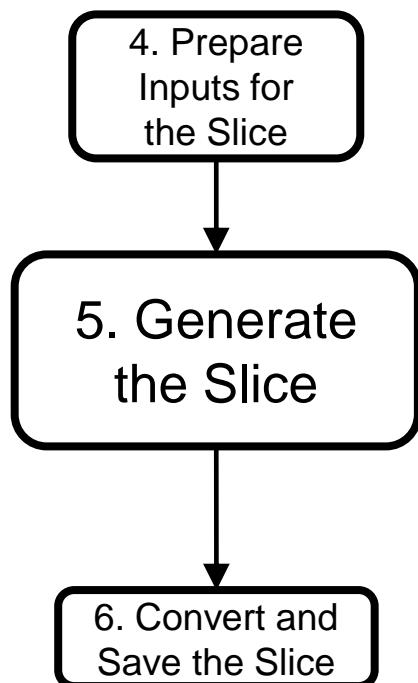
- ❖ List of Keras tensor(s) is created for the input to the slice



```
250 # Perform slicing and conversion
251 slices = []
252 slice_inputs = []
253 for i in range(num_slices):
254     # Prepare inputs for each slice
255     if i == 0:
256         slice_inputs = model.inputs
257     else:
258         slice_inputs = slices[i-1].outputs
259
260     # Slice the model using slice_dnn
261     slice = slice_dnn(model, starts[i], starts[i+1]-1, slice_inputs)
262     slices.append(slice)
263
264     # Convert and save the slice to a LiteRT model
265     convert_save_slice(args.output_dir, slice, i)
```

5. Generate the Slice

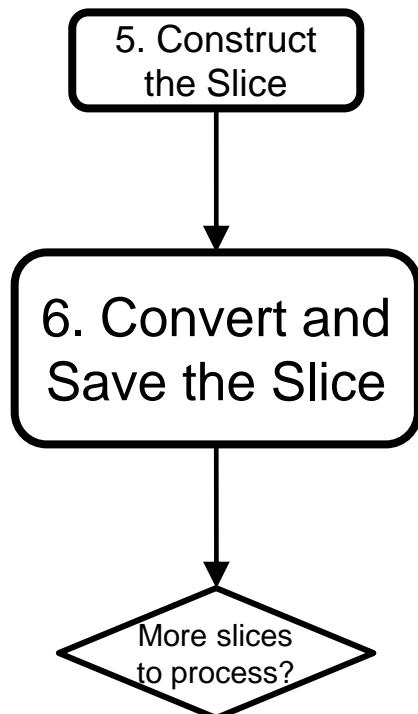
- ❖ For the given model, create a slice, one at a time
 - According to user-provided slicing settings



```
250 # Perform slicing and conversion
251 slices = []
252 slice_inputs = []
253 for i in range(num_slices):
254     # Prepare inputs for each slice
255     if i == 0:
256         slice_inputs = model.inputs
257     else:
258         slice_inputs = slices[i-1].outputs
259
260     # Slice the model using slice_dnn
261     slice = slice_dnn(model, starts[i], starts[i+1]-1, slice_inputs)
262     slices.append(slice)
263
264     # Convert and save the slice to a LiteRT model
265     convert_save_slice(args.output_dir, slice, i)
```

6. Convert and Save the Slice

- ❖ Convert and save the slice into a LiteRT model



```
250 # Perform slicing and conversion
251 slices = []
252 slice_inputs = []
253 for i in range(num_slices):
254     # Prepare inputs for each slice
255     if i == 0:
256         slice_inputs = model.inputs
257     else:
258         slice_inputs = slices[i-1].outputs
259
260     # Slice the model using slice_dnn
261     slice = slice_dnn(model, starts[i], starts[i+1]-1, slice_inputs)
262     slices.append(slice)
263
264     # Convert and save the slice to a LiteRT model
265     convert_save_slice(args.output_dir, slice, i)
```

Contents

-
- I. Slicing TensorFlow Model in HDF5: Big Picture
 - II. Prior Knowledge to Understand Model Slicer
 - III. The Slicing Algorithm
 - IV. Step-by-Step Model Slicer Walkthrough
 - V. **Hands-On Exercise: Run Model Slicer**

Run the Model Slicer

❖ Objective

- Run model slicer python script and see the outputs

❖ Do

- Slice ResNet50 into two submodels

❖ Verify

- Observe submodel structure using Netron

❖ Time

- 5 minutes

Slicing ResNet50 (1)

- ❖ Command to slice the target DNN model (ResNet50)

- `$ python model_slicer.py --model-path ./models/resnet50.h5`

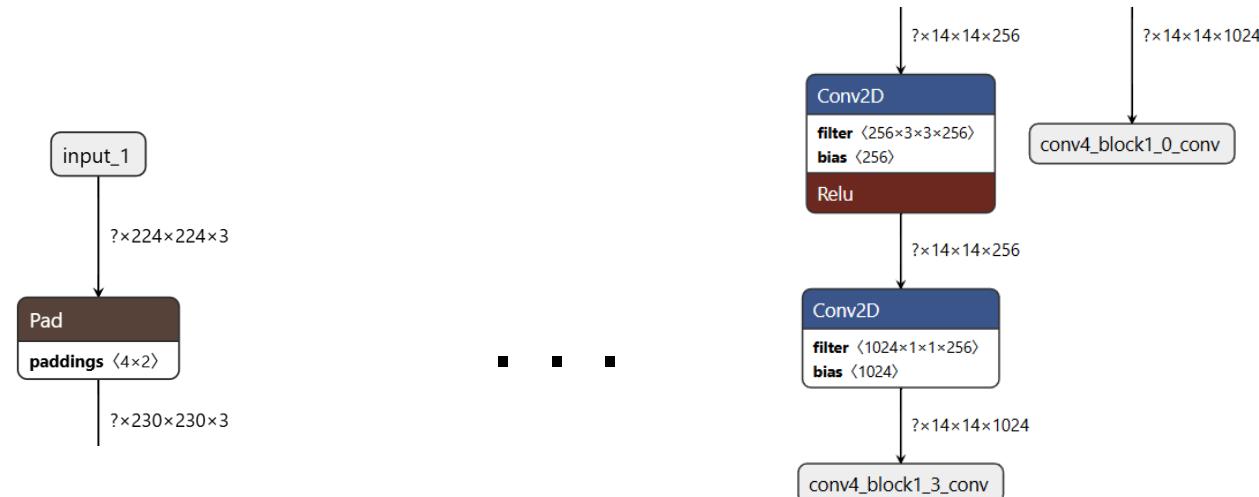
- ❖ For lecture 4, slice ResNet50 into two submodels

- `submodel_0`: From layer 1 to layer $x_1 - 1$
 - `submodel_1`: From layer x_1 to layer 176
 - Choose x_1 freely

```
(.venv) rtoslab@RUBIKPi:~/DNNPipe-Tutorial$ python model_slicer.py --model-path ./models/resnet50.h5
How many submodels? 2
Generated submodels look like: (1, x1-1), (x1, 176)
Enter x1: 89
Layer index ranges for each submodel: [(1, 88), (89, 176)]
Saved LiteRT model to: ./models/submodel_0.tflite
Saved LiteRT model to: ./models/submodel_1.tflite
```

Slicing ResNet50 (2)

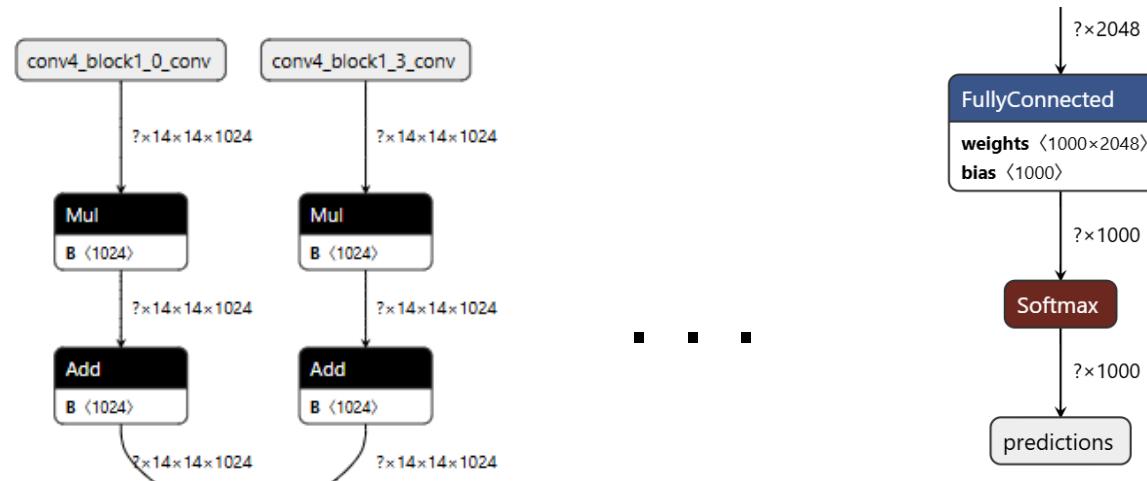
❖ `netron ./models/submodel_0.tflite`



submodel	Input	Output	Saved As
0	<code>input_1</code>	<code>conv4_block1_3_conv, conv4_block1_0_conv</code>	<code>submodel_0.tflite</code>
1	<code>conv4_block1_3_conv, conv4_block1_0_conv</code>	Predictions (final output)	<code>submodel_1.tflite</code>

Slicing ResNet50 (3)

❖ `netron ./models/submodel_1.tflite`



submodel	Input	Output	Saved As
0	input_1	<code>conv4_block1_3_conv, conv4_block1_0_conv</code>	submodel_0.tflite
1	<code>conv4_block1_3_conv, conv4_block1_0_conv</code>	Predictions (final output)	submodel_1.tflite

Tutorial Outline

Lecture 1: *Exercise Overview and Setup* (1 h 30 min)

Lecturer Seongsoo Hong

Topics Motivating Example

Development Environment Setup

Coffee Break (30 min)

Lecture 2: *From Inference Driver to Inference Runtime* (1 h 30 min)

Lecturer Seongsoo Hong and Namcheol Lee

Topics Step-by-Step Inference Driver Walkthrough

Internals of LiteRT

Lunch Break (1 h)

Lecture 3: *Model Slicer* (1 h 30 min)

Lecturer Seongsoo Hong

Topic Model Slicer: Slicing and Conversion Tool for LiteRT

Coffee Break (30 min)

Lecture 4: *Throughput Enhancement on Heterogeneous Accelerators* (1 h 30 min)

Lecturer Namcheol Lee

Topic Implementing a Pipelined Inference Driver for Heterogeneous Processors

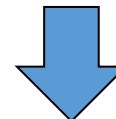
Contents

- I. Overview
- II. Main Thread
- III. Stage Threads
- IV. Throughput Comparison
- V. Hands-on Exercise: Balance Pipeline Stages

Goal

- ❖ In this lecture, we will
 - Implement a pipelined inference driver
 - Measure throughput and compare it with the throughput of the inference driver

[INFO] Throughput: 20.312 items/sec (500 items in 24616 ms)



[INFO] Throughput: [] ? (500 items in [] ?)

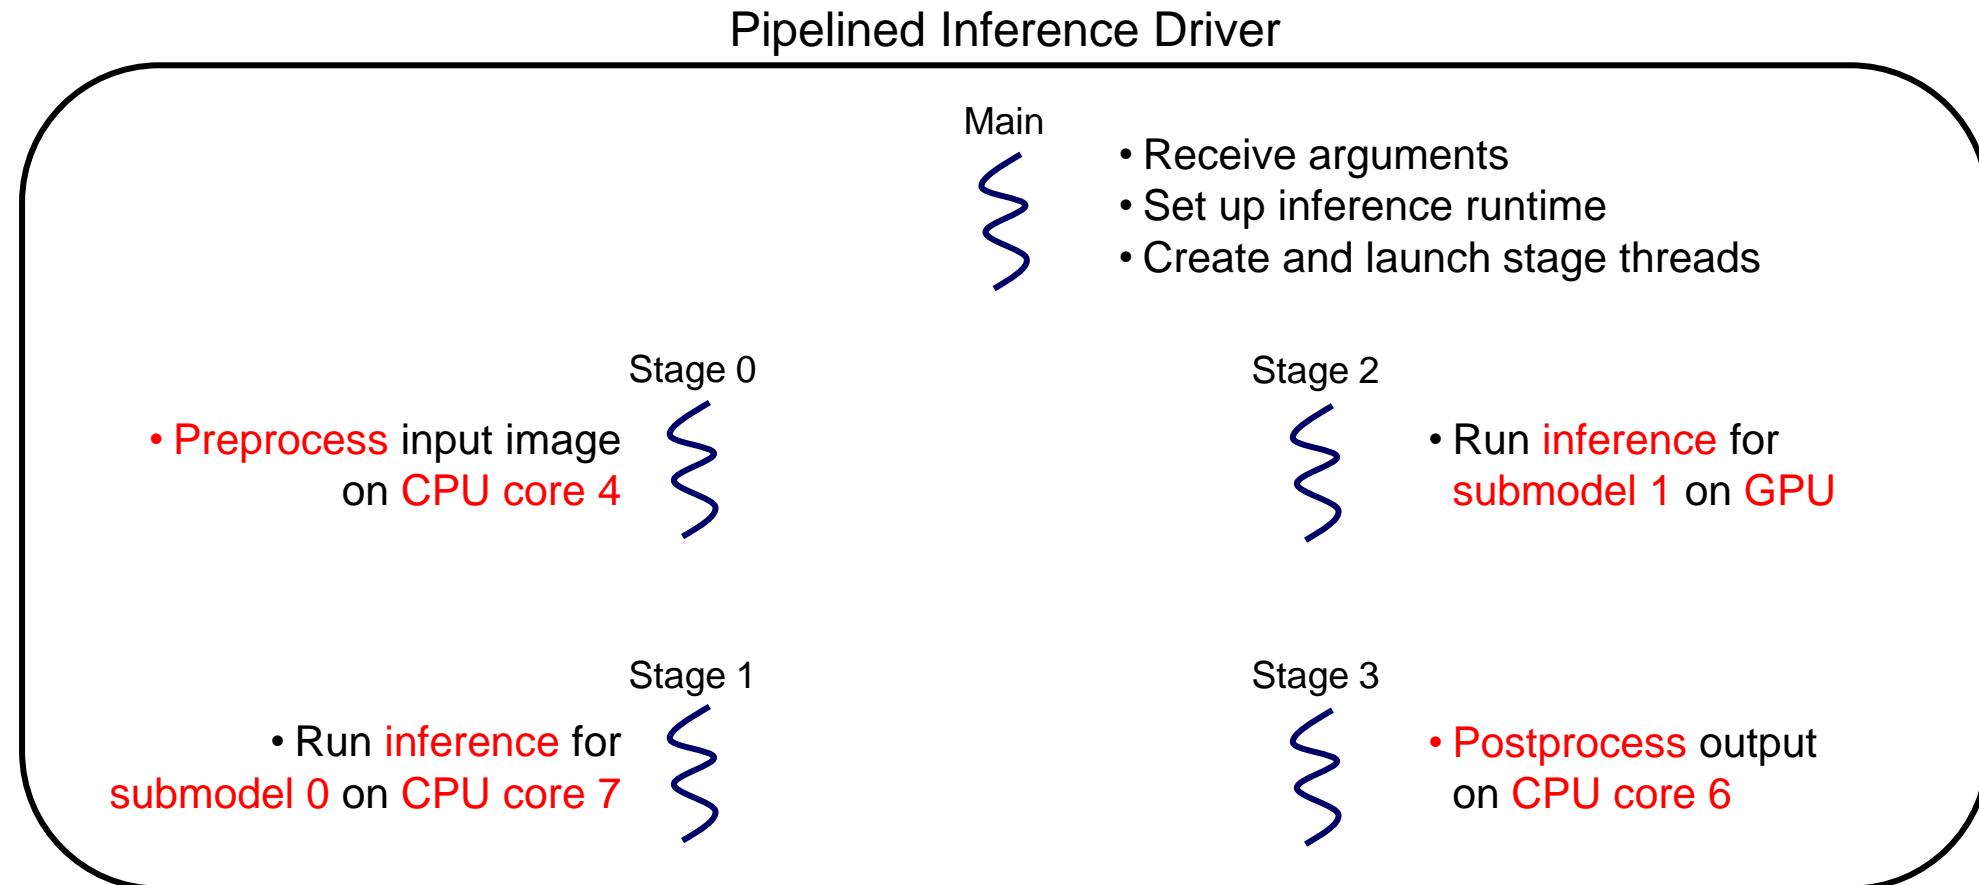
Materials

- ❖ Skeleton code
 - `src/pipelined_inference_driver.cpp`
- ❖ Submodels
 - `models/submodel_0.tflite`
 - Layers from 1 to $x_1 - 1$ of `models/resnet50.h5`
 - `models/submodel_1.tflite`
 - Layers from x_1 to 176 of `models/resnet50.h5`

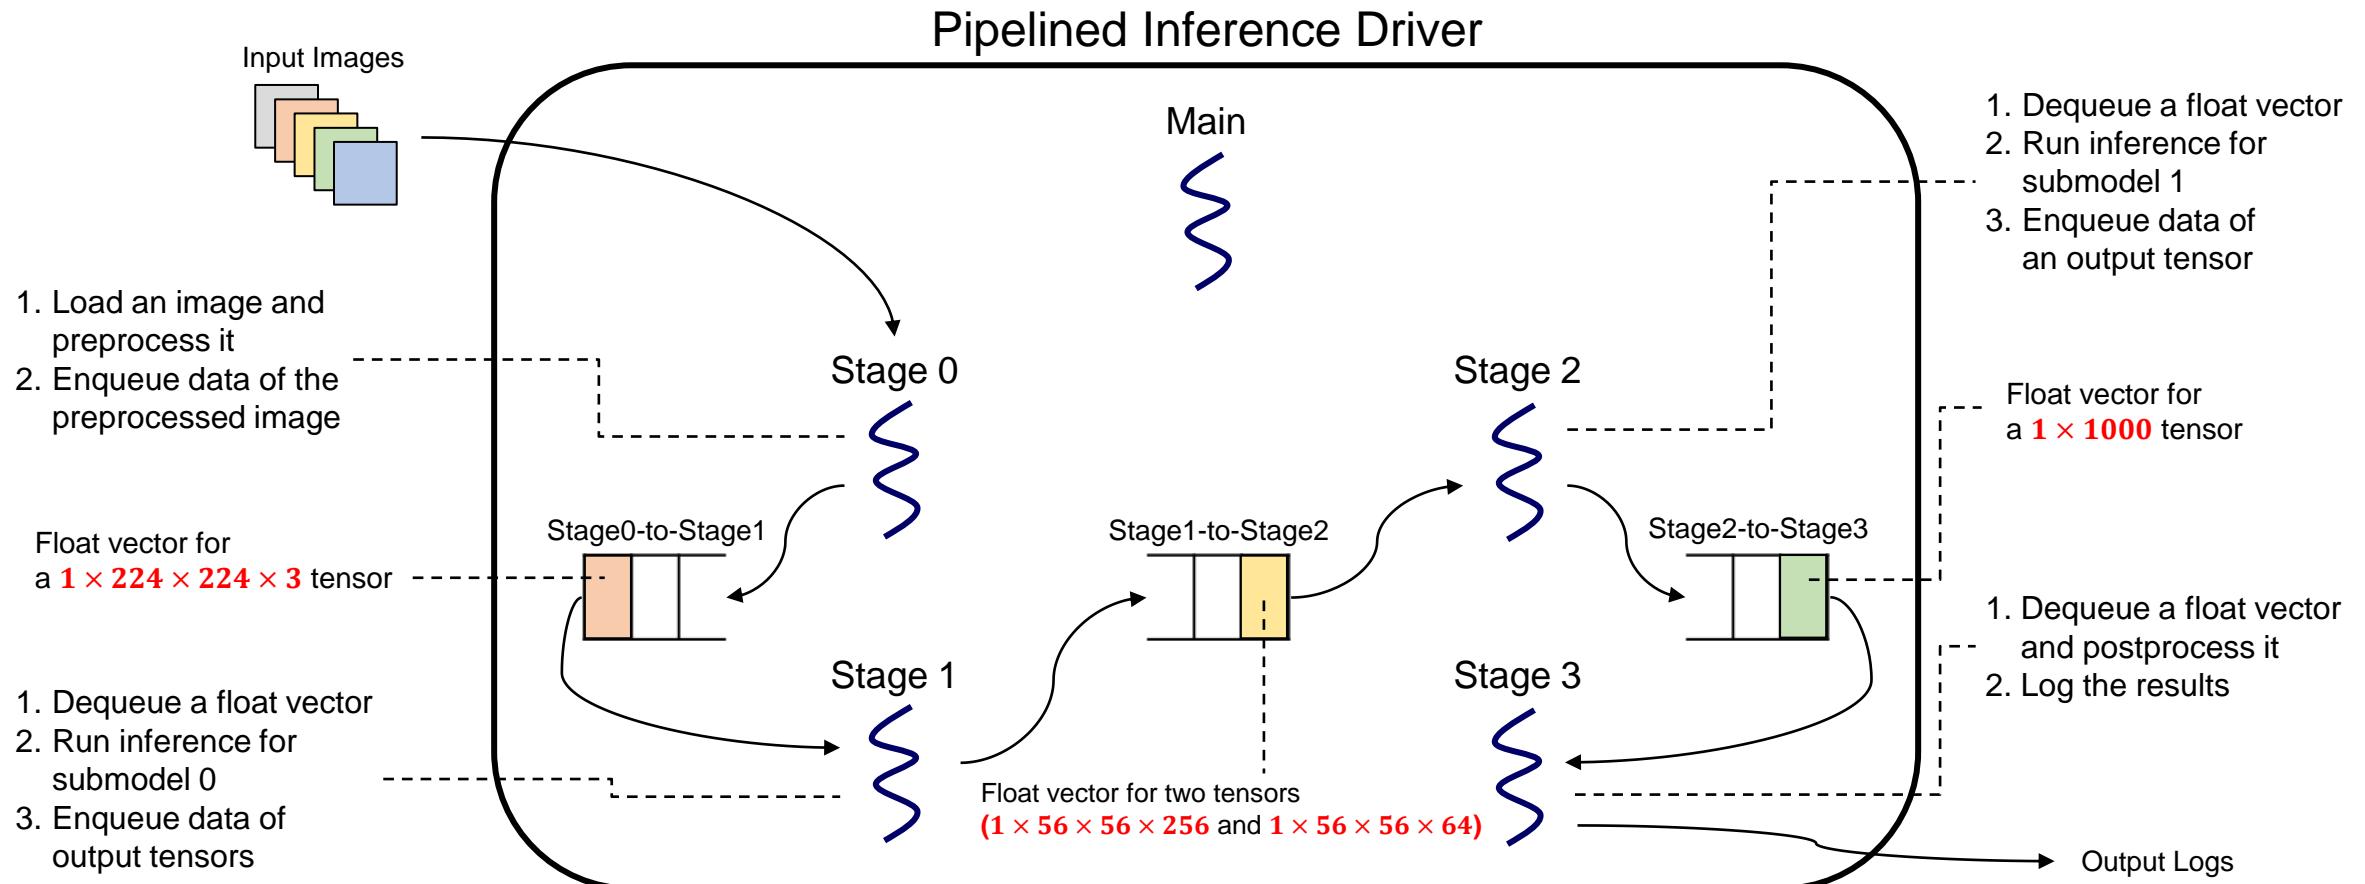
Multithreaded Structure (1)

- ❖ Pipelined inference driver with two submodels consists of
 - One main thread
 - Four worker threads
 - One for preprocessing, another for postprocessing
 - Two threads for two submodels, one for each
 - Equivalently, four stages
 - A worker thread becomes a *stage* when it is allocated onto a dedicated accelerator

Multithreaded Structure (2)



Multithreaded Structure (3): Operational Flow



Data Structure (1)

❖ InterStageQueue

- Class implementing a thread-safe queue for transferring **StagePayload** variables between stages

```
template <typename T>
class InterStageQueue
{
public:
    void push(T item) // Push an item to the queue
    {
        ...
    }

    bool pop(T &item) // Pop an item from the queue
    {
        ...
    }

    void signal_shutdown() // Signal that no more items will be pushed
    {
        ...
    }

    size_t size() // Get the number of items in the queue
    {
        ...
    }
}
```

Data Structure (2)

❖ **InterStageQueue** (cont'd)

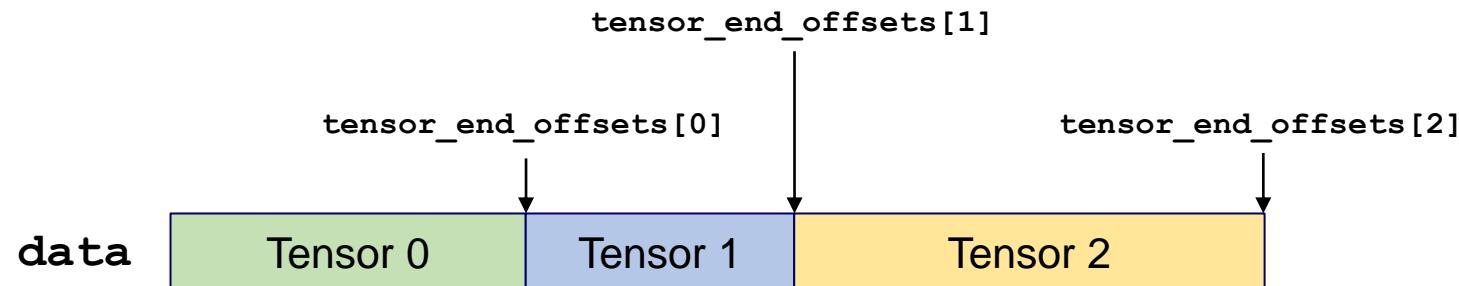
- In the pipelined inference driver
 - Three instances of **InterStageQueue** are defined as global variables

```
34 // === Queues for inter-stage communication ===
35 // stageX_to_stageY_queue: from stageX to stageY
36 InterStageQueue<StagePayload> stage0_to_stage1_queue;
37 InterStageQueue<StagePayload> stage1_to_stage2_queue;
38 InterStageQueue<StagePayload> stage2_to_stage3_queue;
```

Queue Item Structure

❖ StagePayload

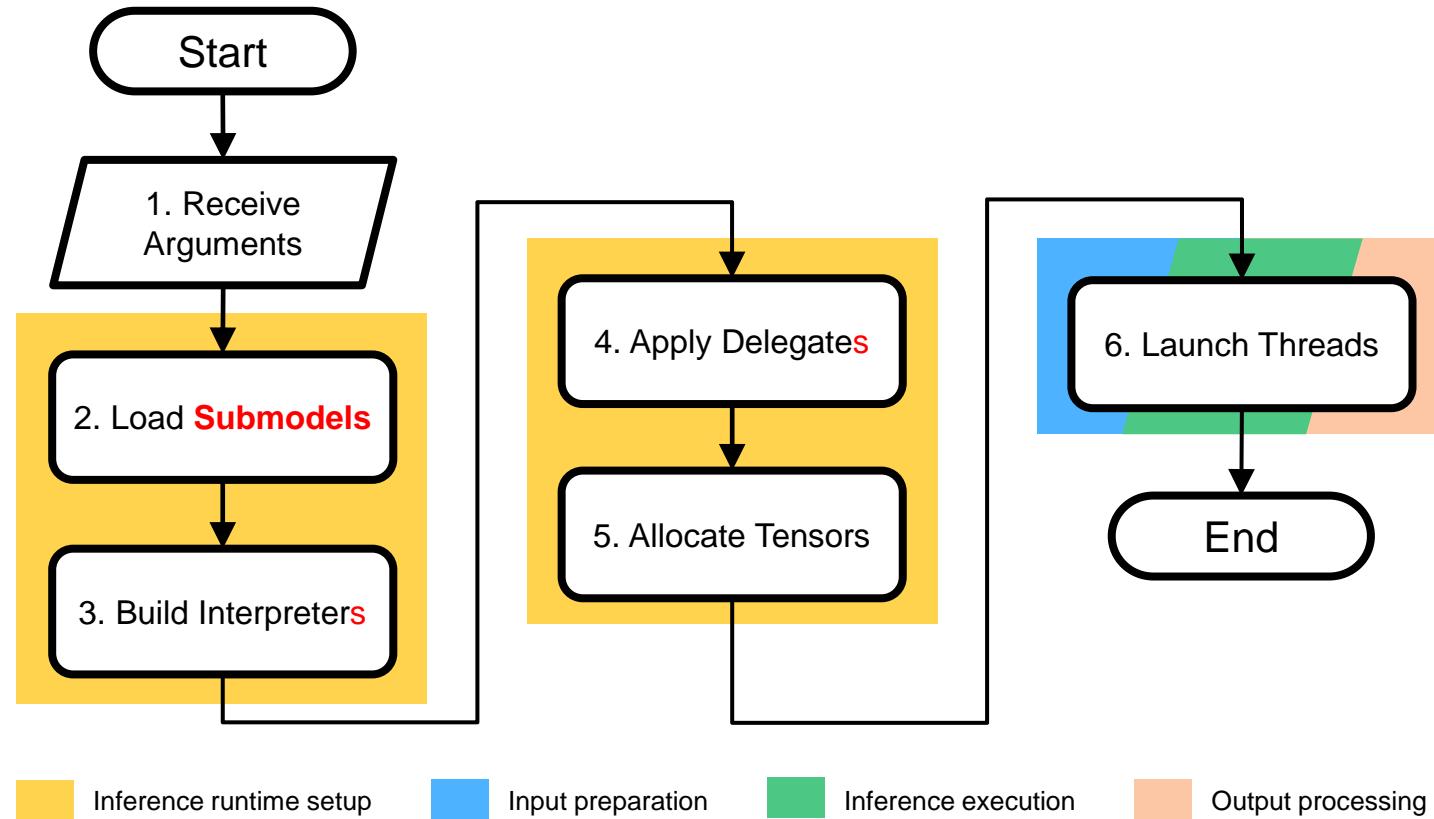
```
struct StagePayload {  
    int index; // Index of the input image (used for tracking)  
    std::vector<float> data; // Flattened data of input/output tensors  
    std::vector<int> tensor_end_offsets; // Offsets marking the end of each tensor in the flattened data  
};
```



Contents

- I. Overview
- II. Main Thread**
- III. Stage Threads
- IV. Throughput Comparison
- V. Hands-on Exercise: Balance Pipeline Stages

Code Flow



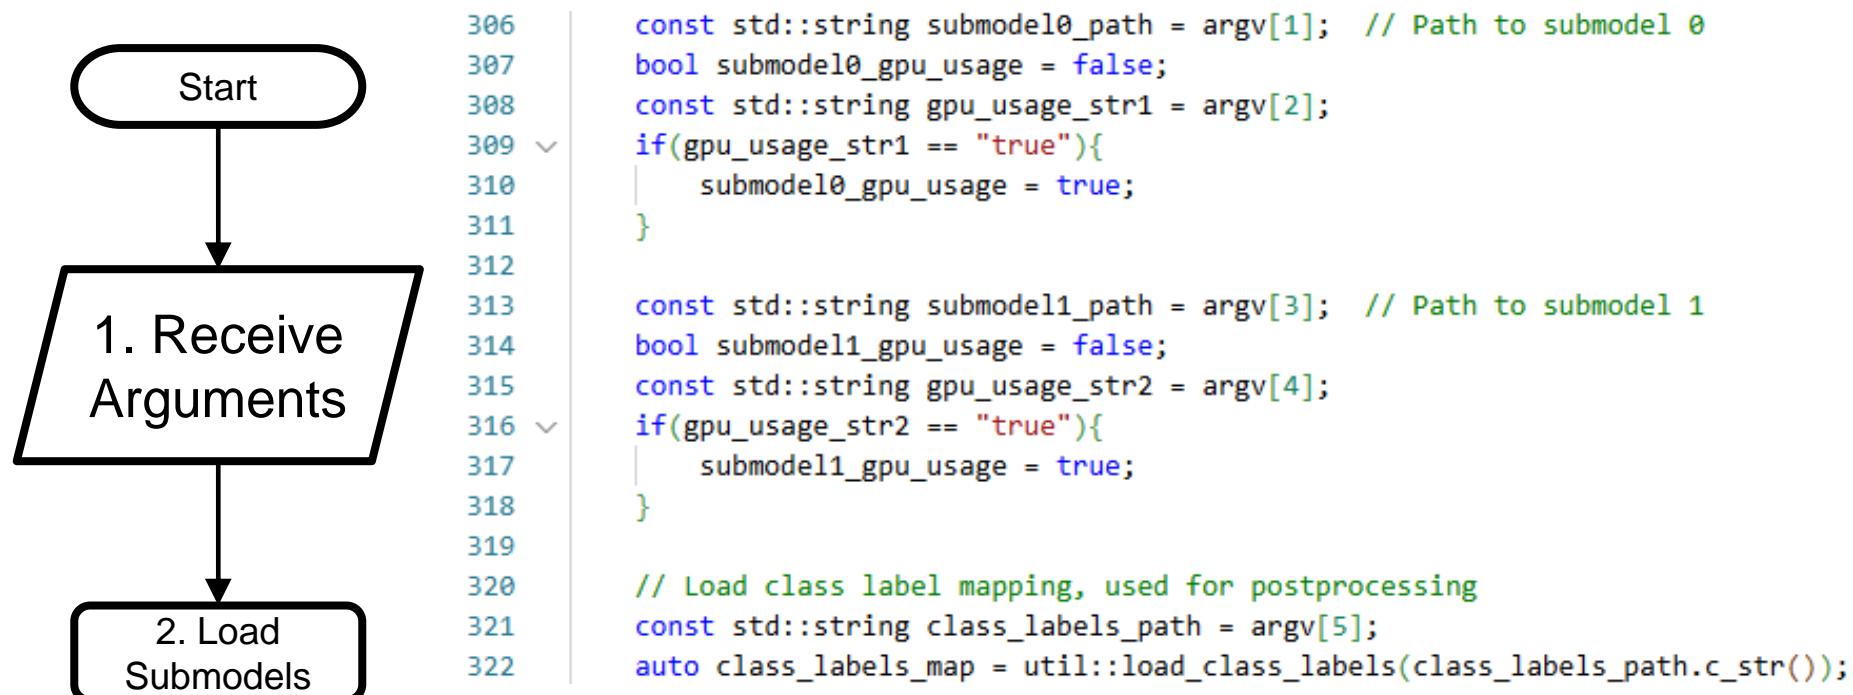
1. Receive Arguments (1)

- ❖ Check the number of arguments



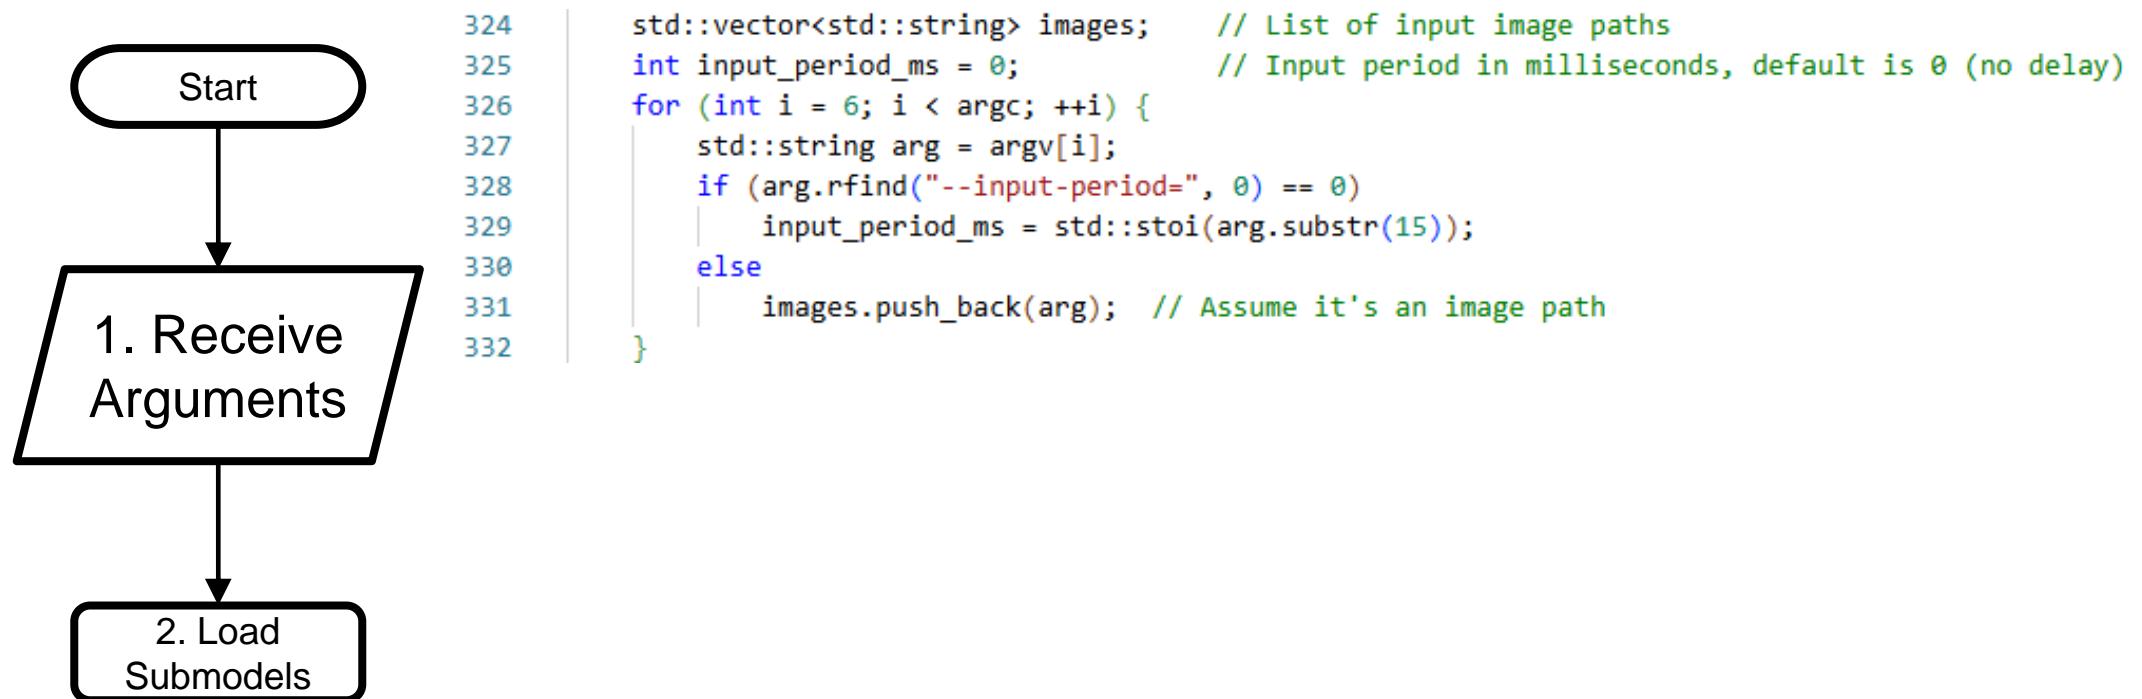
1. Receive Arguments (2)

- ❖ Set variables based on the arguments



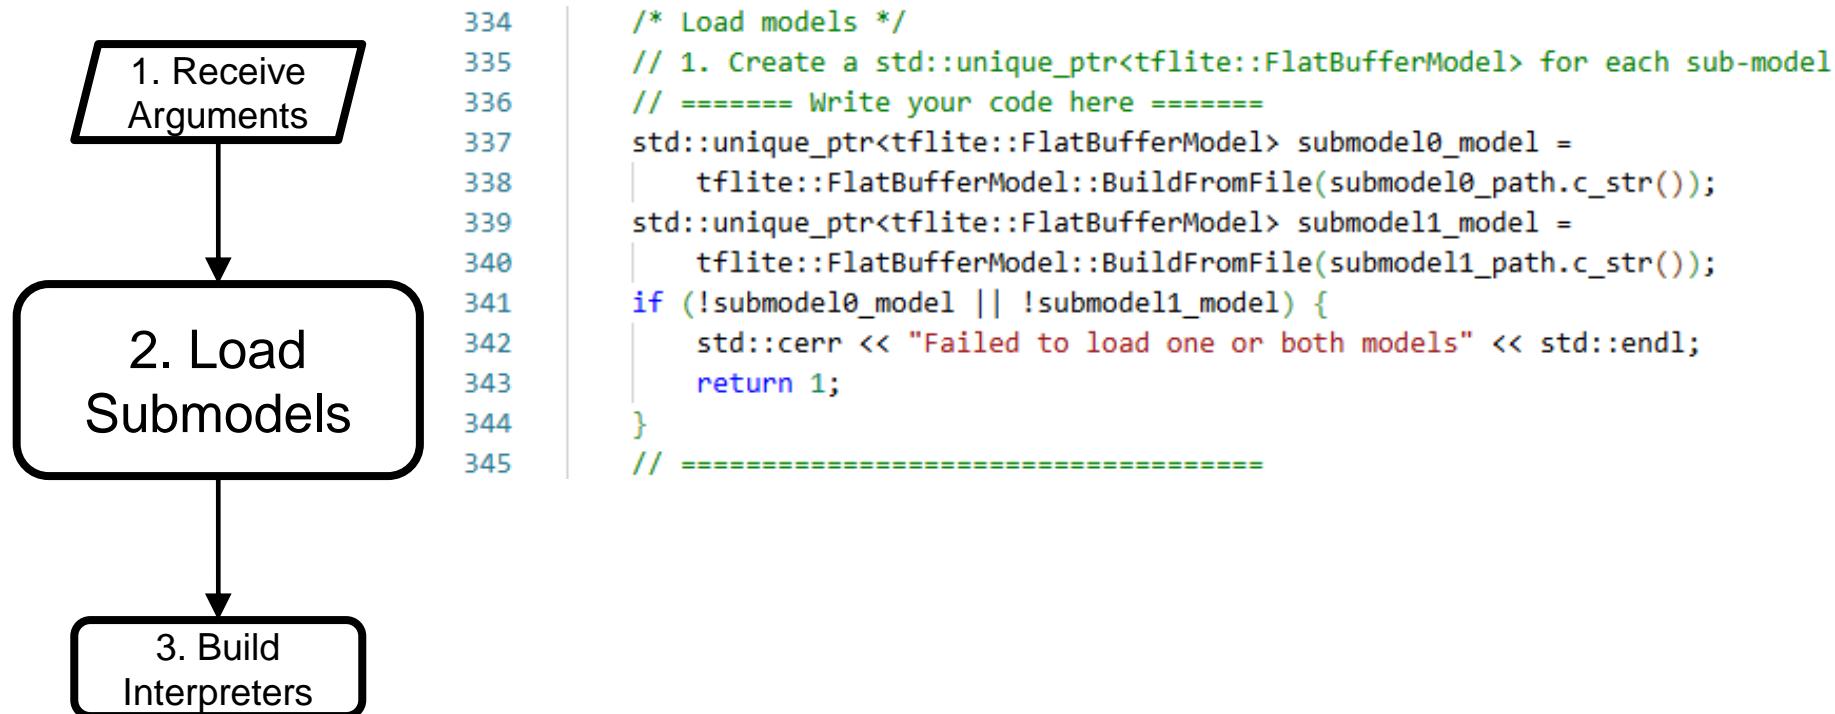
1. Receive Arguments (3)

- ❖ Set variables based on the arguments (cont'd)



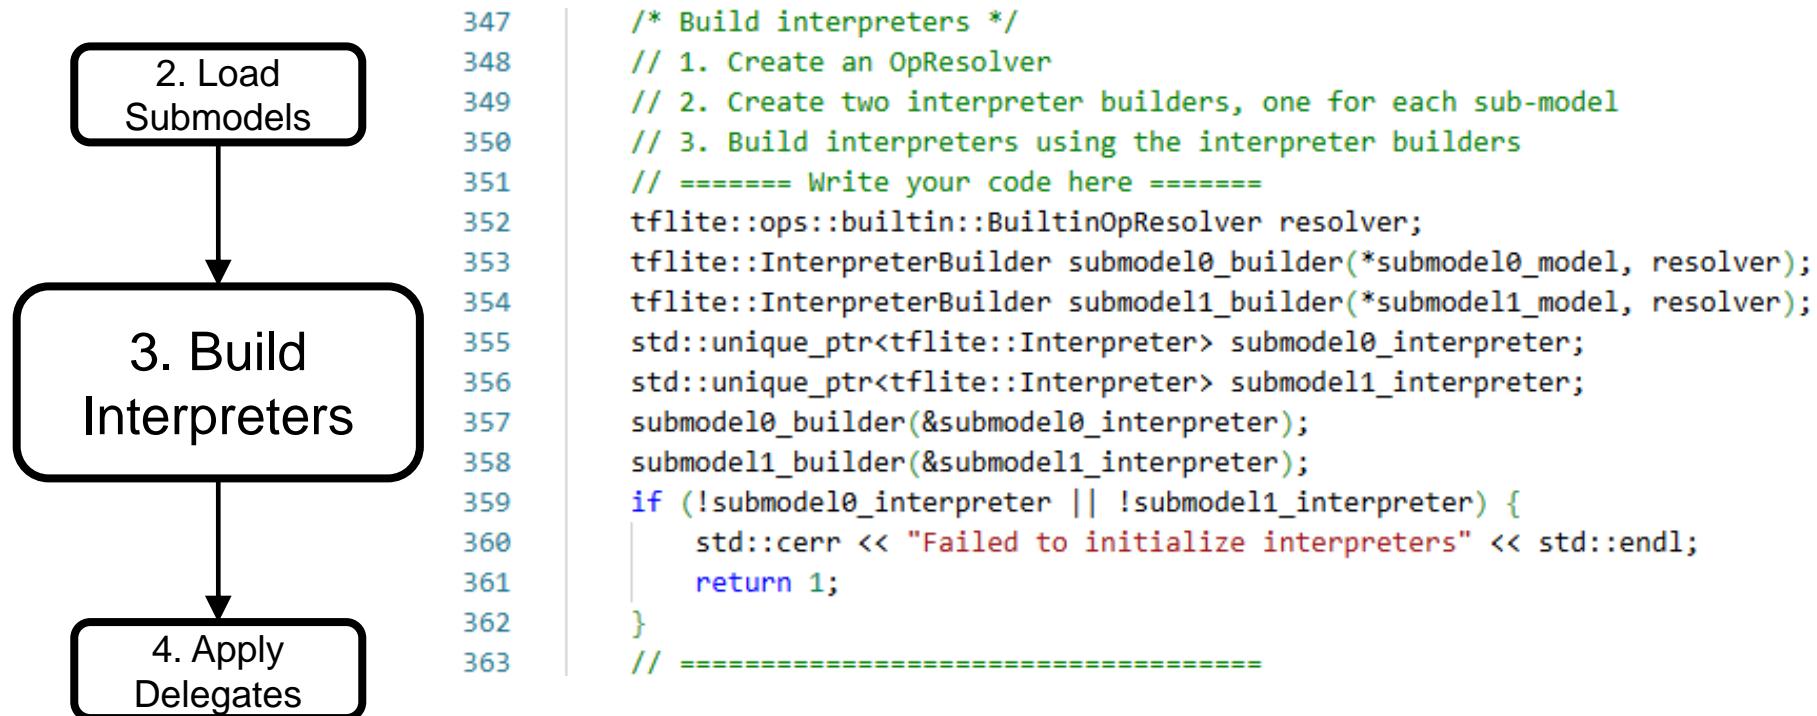
2. Load Submodels

- ❖ Load submodel 0 and 1



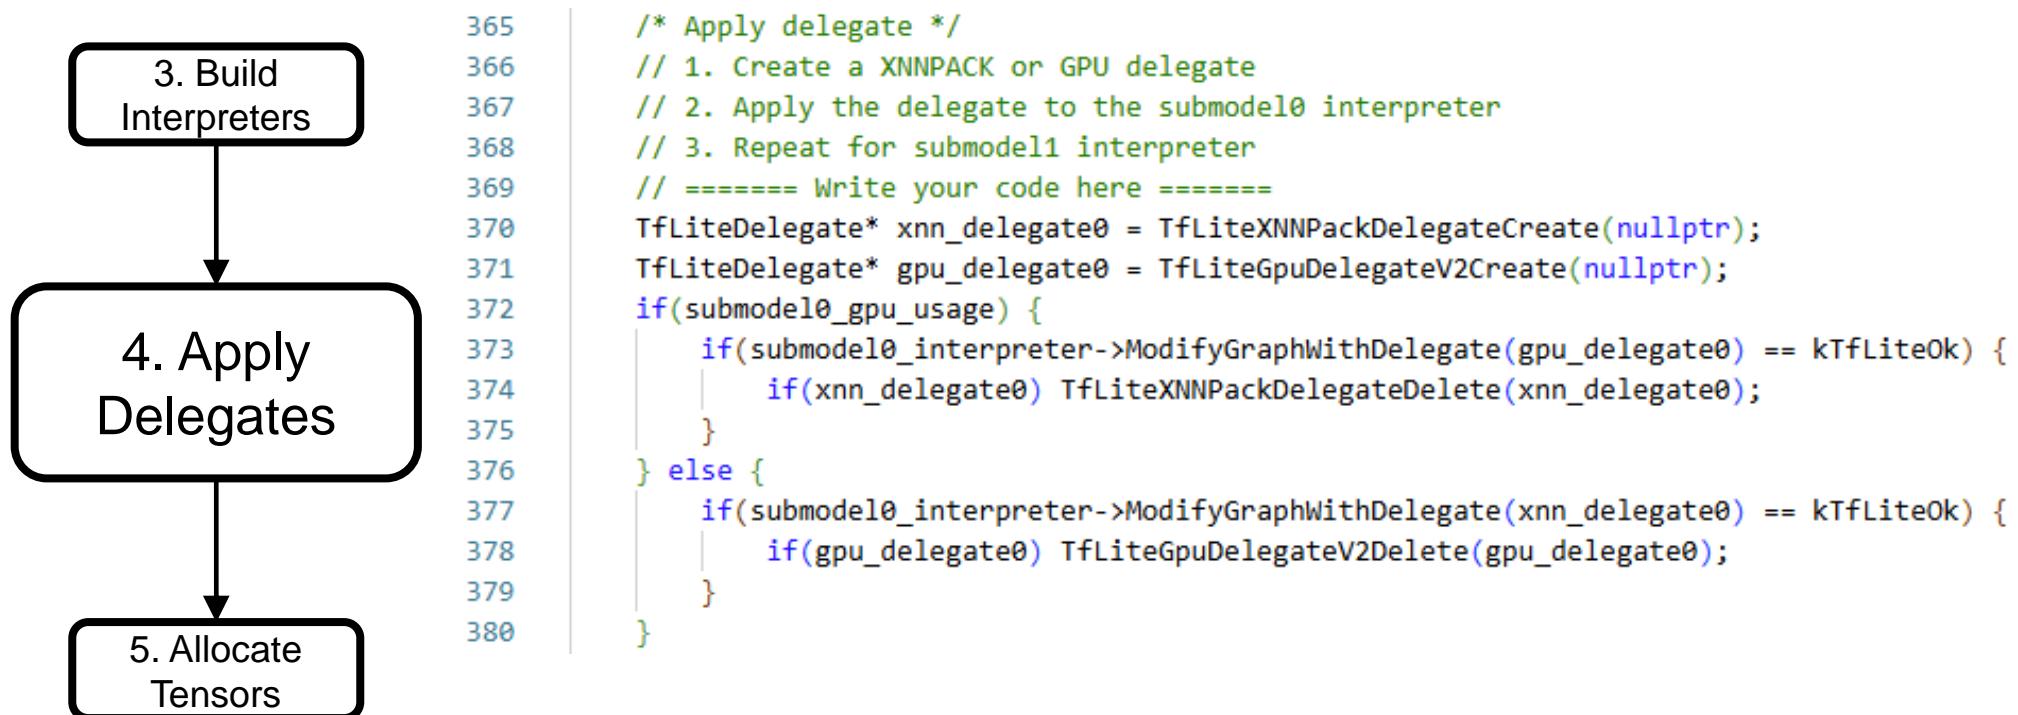
3. Build Interpreters

- ❖ Build interpreters for submodel 0 and submodel 1



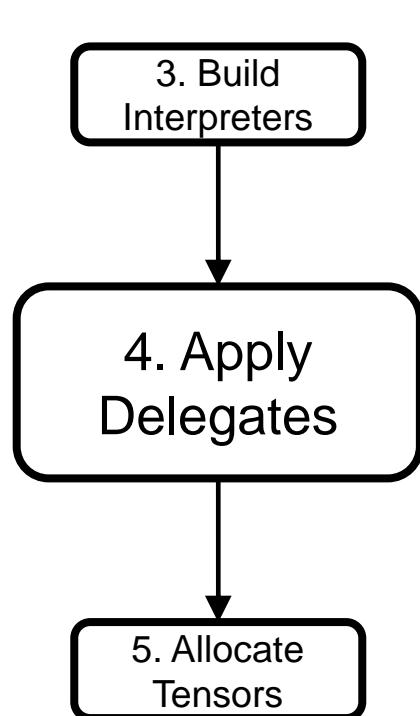
4. Apply Delegates (1)

- ❖ Create and apply delegates to the interpreters



4. Apply Delegates (2)

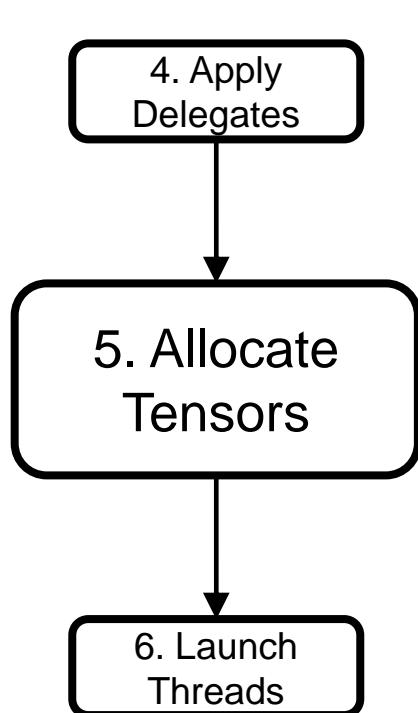
- ❖ Create and apply delegates to the interpreters (cont'd)



```
382     TfLiteDelegate* xnn_delegate1 = TfLiteXNNPackDelegateCreate(nullptr);
383     TfLiteDelegate* gpu_delegate1 = TfLiteGpuDelegateV2Create(nullptr);
384     if(submodel1_gpu_usage) {
385         if(submodel1_interpreter->ModifyGraphWithDelegate(gpu_delegate1) == kTfLiteOk) {
386             if(xnn_delegate1) TfLiteXNNPackDelegateDelete(xnn_delegate1);
387         }
388     } else {
389         if(submodel1_interpreter->ModifyGraphWithDelegate(xnn_delegate1) == kTfLiteOk) {
390             if(gpu_delegate1) TfLiteGpuDelegateV2Delete(gpu_delegate1);
391         }
392     }
393 // =====
```

5. Allocate Tensors

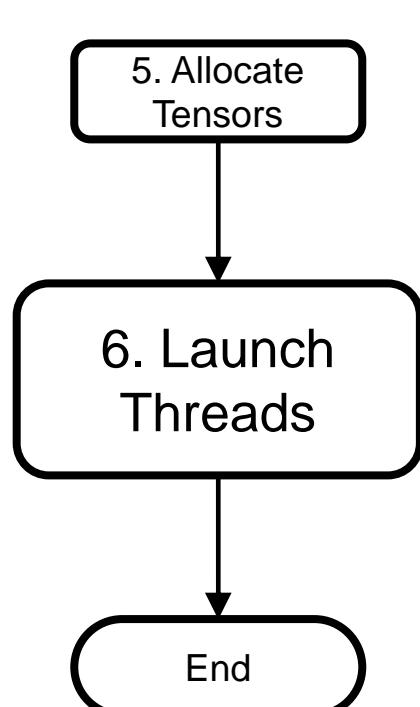
- ❖ Allocate tensors for both interpreters



```
395     /* Allocate tensors */  
396     // 1. Allocate tensors for both interpreters  
397     // ===== Write your code here =====  
398     if (submodel0_interpreter->AllocateTensors() != kTfLiteOk) {  
399         std::cerr << "Failed to allocate tensors for submodel0" << std::endl;  
400         return 1;  
401     }  
402     if (submodel1_interpreter->AllocateTensors() != kTfLiteOk) {  
403         std::cerr << "Failed to allocate tensors for submodel1" << std::endl;  
404         return 1;  
405     }  
406     // =====
```

6. Launch Threads

- ❖ Create and launch threads for each stage



```
411  /* Create and launch threads */
412  // 1. Launch stage0_worker in a new thread with images and input_period_ms
413  // 2. Launch stage1_worker in a new thread with submodel0 interpreter
414  // 3. Launch stage2_worker in a new thread with submodel1 interpreter
415  // 4. Launch stage3_worker in a new thread with class_labels_map
416  // ===== Write your code here =====
417  std::thread stage0_thread(stage0_worker, images, input_period_ms);
418  std::thread stage1_thread(stage1_worker, submodel0_interpreter.get());
419  std::thread stage2_thread(stage2_worker, submodel1_interpreter.get());
420  std::thread stage3_thread(stage3_worker, class_labels_map);
421  // ======
```

Contents

- I. Overview
- II. Main Thread
- III. Stage Threads**
- IV. Throughput Comparison
- V. Hands-on Exercise: Balance Pipeline Stages

Four Stage Workers

- ❖ Functions executed by stage threads

1. `void stage0_worker(...);`
2. `void stage1_worker(...);`
3. `void stage2_worker(...);`
4. `void stage3_worker(...);`

1. Stage0 Worker (1)

❖ Specification

- Inputs
 - Vector of image file paths (`images`)
 - Input period in milliseconds (`input_period_ms`)
- Output
 - `StagePayload` containing a preprocessed image of shape **$224 \times 224 \times 3$**
- Behavior
 1. Load and preprocess an image from the image file path vector
 2. Convert the preprocessed image into an `StagePayload`
 3. Push the `StagePayload` into `stage0_to_stage1_queue`
 4. Sleep until the next input period, unless already behind schedule
 - If no images remain, signal shutdown and return

1. Stage0 Worker (2)

1. Load and preprocess an image from the image file path vector

```
40  void stage0_worker(const std::vector<std::string>& images, int input_period_ms) {
41      auto next_wakeup_time = std::chrono::high_resolution_clock::now();
42      size_t idx = 0;
43      do {
44          std::string label = "Stage0 " + std::to_string(idx);
45          util::timer_start(label);
46          /* Preprocessing */
47          // Load image
48          cv::Mat image = cv::imread(images[idx]);
49          if (image.empty()) {
50              std::cerr << "[Stage0] Failed to load image: " << images[idx] << "\n";
51              util::timer_stop(label);
52              continue;
53          }
54
55          // Preprocess image
56          cv::Mat preprocessed_image = util::preprocess_image_resnet(image, 224, 224);
57          if (preprocessed_image.empty()) {
58              std::cerr << "[Stage0] Preprocessing failed: " << images[idx] << "\n";
59              util::timer_stop(label);
60              continue;
61      }
```

1. Stage0 Worker (3)

2. Convert the preprocessed image into an **StagePayload**

```
63     /* Create an StagePayload, copy preprocessed_image data into it,
64      * and push it into stage0_to_stage1_queue */
65      // Hint: std::memcpy(destination_ptr, source_ptr, num_bytes);
66      StagePayload stage_payload;
67      // ===== Write your code here =====
68      stage_payload.index = idx;
69      stage_payload.data.resize(
70          preprocessed_image.total() * preprocessed_image.channels());
71      std::memcpy(stage_payload.data.data(), preprocessed_image.ptr<float>(),
72                  stage_payload.data.size() * sizeof(float));
73      stage_payload.tensor_end_offsets =
74          {static_cast<int>(stage_payload.data.size())};
75      // =====
```

1. Stage0 Worker (4)

3. Push the **StagePayload** into **stage0_to_stage1_queue**

```
76     stage0_to_stage1_queue.push(std::move(stage_payload));  
77     ++idx;
```

4. Sleep until the next input period, unless already behind schedule

- If no images remain, signal shutdown and return

```
81     // Sleep to control the input rate  
82     // If next_wakeup_time is in the past, it will not sleep  
83     next_wakeup_time += std::chrono::milliseconds(input_period_ms);  
84     std::this_thread::sleep_until(next_wakeup_time);  
85 } while (idx < images.size());  
86  
87 // Notify stage1_thread that no more data will be sent  
88 stage0_to_stage1_queue.signal_shutdown();  
89 } // end of stage0_worker
```

2. Stage1 Worker (1)

❖ Specification

- Input
 - **StagePayload** containing one tensor of shape **$1 \times 224 \times 224 \times 3$**
- Outputs
 - **StagePayload** containing two tensors of shape
 $1 \times 56 \times 56 \times 256$ and **$1 \times 56 \times 56 \times 64$**

2. Stage1 Worker (2)

❖ Specification (cont'd)

▪ Behavior

1. Pop a **StagePayload** from **stage0_to_stage1_queue**
2. Copy its data into the input tensor of the **interpreter**
3. Run inference using the **interpreter**
4. Extract data from the **interpreter**'s output tensors and copy into a **StagePayload**
5. Push the **StagePayload** into **stage1_to_stage2_queue**
6. Repeat until **stage0_to_stage1_queue** signals shutdown and is empty
 - Then signal shutdown **stage1_to_stage2_queue** and return

2. Stage1 Worker (3)

1. Pop an StagePayload from stage0_to_stage1_queue

```
91 void stage1_worker(tflite::Interpreter* interpreter) {  
92     StagePayload stage_payload;  
93     bool print_tensor_shape = true;  
94  
95     while (stage0_to_stage1_queue.pop(stage_payload)) {
```

2. Stage1 Worker (4)

2. Copy its data into the input tensor of the **interpreter**

```
99     /* Extract tensor data from stage_payload.data and copy into
100    * the corresponding input tensors of the interpreter */
101   // ===== Write your code here =====
102   if(print_tensor_shape) std::cout << "Input Tensor(s) of Stage 1" << std::endl;
103   for (size_t i = 0; i < interpreter->inputs().size(); ++i) {
104     // Get i-th input tensor object
105     TfLiteTensor* input_tensor = interpreter->input_tensor(i);
```

2. Stage1 Worker (5)

2. Copy its data into the input tensor of the **interpreter** (cont'd)

```
107     // Calculate the required number of elements for the i-th input tensor based on its dimensions
108     int num_elements = 1;
109     if(print_tensor_shape) std::cout << "- Shape of input tensor " << i << ": [";
110     for (int d = 0; d < input_tensor->dims->size; ++d) {
111         num_elements *= input_tensor->dims->data[d];
112         if(print_tensor_shape) {
113             std::cout << input_tensor->dims->data[d];
114             if(d != input_tensor->dims->size - 1) std::cout << ", ";
115         }
116     }
117     if(print tensor shape) std::cout << "], # of elements = " << num_elements << std::endl;
```

```
Input Tensor(s) of Stage 1
- Shape of input tensor 0: [1, 224, 224, 3], # of elements = 150528
```

- Shape of a **TfLiteTensor*** is stored under **TfLiteTensor->dims->data**
 - It is an integer array with one entry per dimension in order

2. Stage1 Worker (6)

2. Copy its data into the input tensor of the **interpreter** (cont'd)

```
119     // Copy data from stage_payload to i-th input tensor
120     int start_idx = (i == 0) ? 0 : stage_payload.tensor_end_offsets[i-1];
121     int end_idx = stage_payload.tensor_end_offsets[i];
122     int given_elements = end_idx - start_idx;
123
124     if(num_elements != given_elements){
125         std::cerr << "[Stage 1] Input tensor " << i << " size mismatch! " << std::endl;
126         std::cerr << "Expected " << num_elements << ", got " << given_elements << std::endl;
127     } else {
128         std::memcpy(input_tensor->data.f, stage_payload.data.data() + start_idx,
129                     (given_elements) * sizeof(float));
130     }
131 } // end of for loop
132 // =====
```

2. Stage1 Worker (7)

3. Run inference using the **interpreter**

```
134     /* Inference */  
135     // ===== Write your code here =====  
136     interpreter->Invoke();  
137     // =====
```

2. Stage1 Worker (8)

-
4. Extract data from the `interpreter`'s output tensors
and copy into a `StagePayload`

```
139     /* Extract data from the interpreter's output tensors and copy into a StagePayload */
140     // Clear data in it for reuse
141     stage_payload.data.clear();
142     stage_payload.tensor_end_offsets.clear();
143     if(print_tensor_shape) std::cout << "Output Tensor(s) of Stage 1" << std::endl;
144     // ===== Write your code here =====
145     for (size_t i = 0; i < interpreter->outputs().size(); ++i) {
146         // Get i-th output tensor object
147         TfLiteTensor* output_tensor = interpreter->output_tensor(i);
```

2. Stage1 Worker (9)

4. Extract data from the `interpreter`'s output tensors and copy into a `StagePayload` (cont'd)

```
149     // Calculate the number of elements in the i-th output tensor
150     if(print_tensor_shape) std::cout << "- Shape of output tensor " << i << ": [";
151     int num_elements = 1;
152     for (int d = 0; d < output_tensor->dims->size; ++d) {
153         num_elements *= output_tensor->dims->data[d];
154         if(print_tensor_shape) {
155             std::cout << output_tensor->dims->data[d];
156             if(d != output_tensor->dims->size - 1) std::cout << ", ";
157         }
158     }
159     if(print_tensor_shape) std::cout << "], # of elements = " << num_elements << std::endl;
```

Output Tensor(s) of Stage 1
- Shape of output tensor 0: [1, 56, 56, 256], # of elements = 802816
- Shape of output tensor 1: [1, 56, 56, 64], # of elements = 200704

2. Stage1 Worker (10)

4. Extract data from the **interpreter**'s output tensors
and copy into a **StagePayload** (cont'd)

```
161     // Resize stage_payload.data and copy output tensor's data into it
162     int current_data_length = stage_payload.data.size();
163     stage_payload.data.resize(current_data_length + num_elements);
164     std::memcpy(stage_payload.data.data() + current_data_length,
165                 output_tensor->data.f,
166                 num_elements * sizeof(float));
167     stage_payload.tensor_end_offsets.push_back(current_data_length + num_elements);
168 } // end of for loop
169 print_tensor_shape = false;
170 // =====
```

2. Stage1 Worker (11)

5. Push the **StagePayload** into **stage1_to_stage2_queue**

```
172     stage1_to_stage2_queue.push(std::move(stage_payload));  
173  
174     util::timer_stop(label);  
175 } // end of while loop
```

6. Repeat until **stage0_to_stage1_queue**
signals shutdown and is empty

- Then signal shutdown **stage1_to_stage2_queue** and return

```
177     // Notify stage2_thread that no more data will be sent  
178     stage1_to_stage2_queue.signal_shutdown();  
179 } // end of stage1_worker
```

3. Stage2 Worker (1)

❖ Specification

- Inputs
 - **StagePayload** containing two tensors of shape
 $1 \times 56 \times 56 \times 256$ and **$1 \times 56 \times 56 \times 64$**
- Output
 - **StagePayload** containing one tensor of shape **1×1000**

3. Stage2 Worker (2)

❖ Specification (cont'd)

▪ Behavior

1. Pop a **StagePayload** from **stage1_to_stage2_queue**
2. Copy its data into the input tensors of the **interpreter**
3. Run inference using the **interpreter**
4. Extract data from the **interpreter**'s output tensors and copy into a **StagePayload**
5. Push the **StagePayload** into **stage2_to_stage3_queue**
6. Repeat until **stage1_to_stage2_queue** signals shutdown and is empty
 - Then signal shutdown **stage2_to_stage3_queue** and return

3. Stage2 Worker (3)

1. Pop a StagePayload from stage1_to_stage2_queue

```
181 < void stage2_worker(tfLite::Interpreter* interpreter) {  
182     StagePayload stage_payload;  
183     bool print_tensor_shape = true;  
184  
185     while (stage1_to_stage2_queue.pop(stage_payload)) {
```

3. Stage2 Worker (4)

2. Copy its data into the input tensors of the **interpreter**

```
189     /* Extract tensor data from stage_payload.data and copy into
190      * the corresponding input tensors of the interpreter */
191      if(print_tensor_shape) std::cout << "Input Tensor(s) of Stage 2" << std::endl;
192      // ===== Write your code here =====
193      for (size_t i = 0; i < interpreter->inputs().size(); ++i) {
194          // Get i-th input tensor object
195          TfLiteTensor* input_tensor = interpreter->input_tensor(i);
```

3. Stage2 Worker (5)

2. Copy its data into the input tensors of the **interpreter** (cont'd)

```
197     // Calculate the required number of elements for the i-th input tensor based on its dimensions
198     int num_elements = 1;
199     if(print_tensor_shape) std::cout << "- Shape of input tensor " << i << ": [";
200     for (int d = 0; d < input_tensor->dims->size; ++d) {
201         num_elements *= input_tensor->dims->data[d];
202         if(print_tensor_shape) {
203             std::cout << input_tensor->dims->data[d];
204             if(d != input_tensor->dims->size - 1) std::cout << ", ";
205         }
206     }
207     if(print_tensor_shape) std::cout << "], # of elements = " << num_elements << std::endl;
```

Input Tensor(s) of Stage 2

- Shape of input tensor 0: [1, 56, 56, 256], # of elements = 802816
- Shape of input tensor 1: [1, 56, 56, 64], # of elements = 200704

3. Stage2 Worker (6)

2. Copy its data into the input tensors of the **interpreter** (cont'd)

```
209     // Copy data from stage_payload to i-th input tensor
210     int start_idx = (i == 0) ? 0 : stage_payload.tensor_end_offsets[i-1];
211     int end_idx = stage_payload.tensor_end_offsets[i];
212     int given_elements = end_idx - start_idx;
213
214     if(num_elements != given_elements){
215         std::cerr << "[Stage 2] Input tensor " << i << " size mismatch! " << std::endl;
216         std::cerr << "Expected " << num_elements << ", got " << given_elements << std::endl;
217     } else {
218         std::memcpy(input_tensor->data.f, stage_payload.data.data() + start_idx,
219                     (given_elements) * sizeof(float));
220     }
221 } // end of for loop
// =====
```

3. Stage2 Worker (7)

3. Run inference using the **interpreter**

```
224     /* Inference */  
225     // ===== Write your code here =====  
226     interpreter->Invoke();  
227     // =====
```

3. Stage2 Worker (8)

-
4. Extract data from the `interpreter`'s output tensors and copy into a `StagePayload`

```
229     /* Extract data from the interpreter's output tensors and copy into a StagePayload */
230     // Clear data in it for reuse
231     stage_payload.data.clear();
232     stage_payload.tensor_end_offsets.clear();
233     if(print_tensor_shape) std::cout << "Output Tensor(s) of Stage 2" << std::endl;
234     // ===== Write your code here =====
235     for (size_t i = 0; i < interpreter->outputs().size(); ++i) {
236         // Get i-th output tensor object
237         TfLiteTensor* output_tensor = interpreter->output_tensor(i);
```

3. Stage2 Worker (9)

4. Extract data from the **interpreter**'s output tensors and copy into a **StagePayload** (cont'd)

```
239     // Calculate the number of elements in the tensor
240     if(print_tensor_shape) std::cout << "- Shape of output tensor " << i << ": [";
241     int num_elements = 1;
242     for (int d = 0; d < output_tensor->dims->size; ++d) {
243         num_elements *= output_tensor->dims->data[d];
244         if(print_tensor_shape) {
245             std::cout << output_tensor->dims->data[d];
246             if(d != output_tensor->dims->size - 1) std::cout << ", ";
247         }
248     }
249     if(print_tensor_shape) std::cout << "], # of elements = " << num_elements << std::endl;
```

Output Tensor(s) of Stage 2
- Shape of output tensor 0: [1, 1000], # of elements = 1000

3. Stage2 Worker (10)

4. Extract data from the `interpreter`'s output tensors and copy into a `StagePayload` (cont'd)

```
251     // Resize stage_payload.data and copy output tensor data into it
252     int current_data_length = stage_payload.data.size();
253     stage_payload.data.resize(current_data_length + num_elements);
254     std::memcpy(stage_payload.data.data() + current_data_length,
255                 output_tensor->data.f,
256                 num_elements * sizeof(float));
257     stage_payload.tensor_end_offsets.push_back(current_data_length + num_elements);
258 } // end of for loop
259 print_tensor_shape = false;
260 // =====
```

3. Stage2 Worker (11)

5. Push the **StagePayload** into **stage2_to_stage3_queue**

```
261     stage2_to_stage3_queue.push(std::move(stage_payload));
262     util::timer_stop(label);
263 } // end of while loop
```

6. Repeat until **stage1_to_stage2_queue**
signals shutdown and is empty

- Then signal shutdown **stage2_to_stage3_queue** and return

```
265     stage2_to_stage3_queue.signal_shutdown();
266 } // end of stage2_worker
```

4. Stage3 Worker (1)

❖ Specification

- Inputs
 - Map of class labels (`class_labels_map`)
 - `StagePayload` containing one tensor of shape 1×1000
- Output
 - Prediction result
- Behavior
 1. Pop a `StagePayload` from `stage2_to_stage3_queue`
 2. Postprocess the output tensor data in the `StagePayload`
 3. Repeat until `stage2_to_stage3_queue` signals shutdown and is empty

4. Stage3 Worker (2)

1. Pop an StagePayload from stage2_to_stage3_queue

```
268 void stage3_worker(std::unordered_map<int, std::string> class_labels_map) {  
269     StagePayload stage_payload;  
270  
271     while (stage2_to_stage3_queue.pop(stage_payload)) {
```

4. Stage3 Worker (3)

2. Postprocess the output tensor data in the **StagePayload**

```
275     const std::vector<float>& probs = stage_payload.data;  
276  
277     if ((stage_payload.index + 1) % 10 == 0) {  
278         auto top_k_indices = util::get_topK_indices(probs, 3);  
279         std::cout << "\n[stage3] Top-3 prediction for image index "  
280             << stage_payload.index << ":\n";  
281         for (int idx : top_k_indices) {  
282             std::string label = class_labels_map.count(idx)  
283                 ? class_labels_map.at(idx)  
284                 : "unknown";  
285             std::cout << "- Class " << idx << " (" << label  
286                 << "): " << probs[idx] << std::endl;  
287         }  
288     }
```

3. Repeat until **stage2_to_stage3_queue** signals shutdown and is empty

```
291     } // end of while loop  
292 } // end of stage3_worker
```

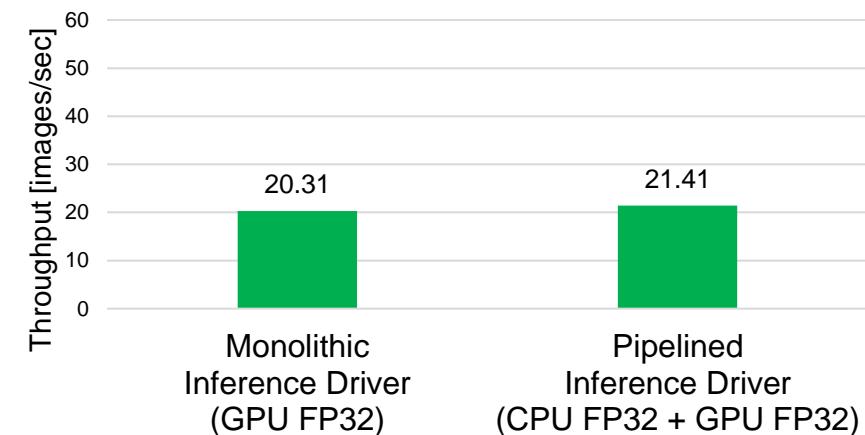
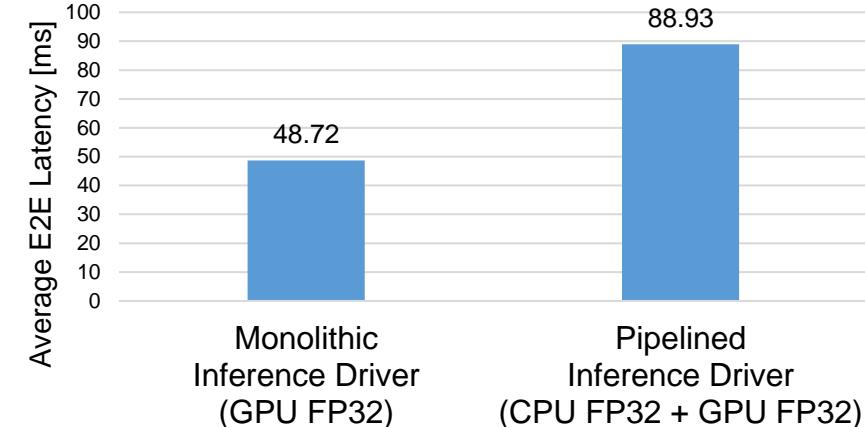
Contents

- I. Overview
- II. Main Thread
- III. Stage Threads
- IV. Throughput Comparison**
- V. Hands-on Exercise: Balance Pipeline Stages

Monolithic vs. Pipelined Inference Drivers (1)

- ❖ End-to-End latency
 - Travel time of a single sample
 - Increased by about 1.83 times

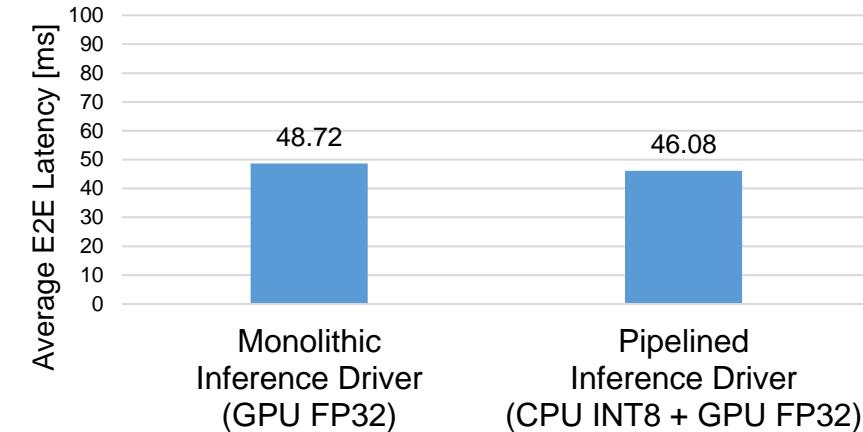
- ❖ Throughput
 - Frames per second
 - Increased by about 5.4%



Monolithic vs. Pipelined Inference Drivers (2)

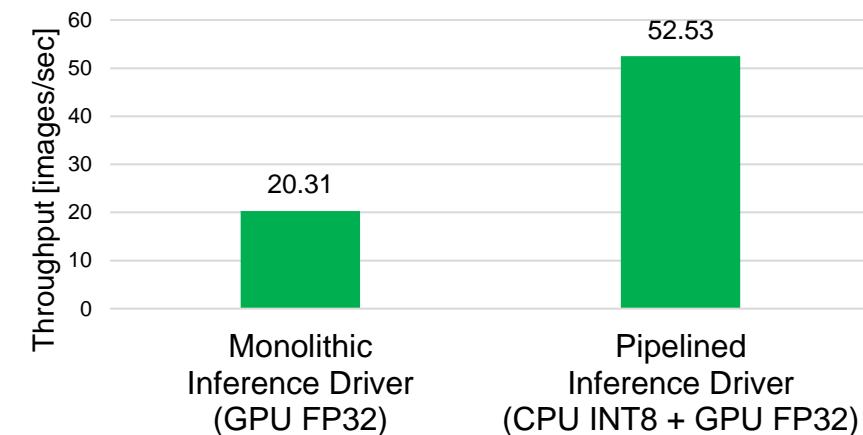
❖ End-to-End latency

- Travel time of a single sample
- Decreased by about 0.95 times



❖ Throughput

- Frames per second
- Increased by about 158.64%
 - When 6 out of 8 cores utilized



Contents

- I. Overview
- II. Main Thread
- III. Stage Threads
- IV. Throughput Comparison
- V. **Hands-on Exercise: Balance Pipeline Stages**

Balance Pipeline Stages

❖ Objective

- Balance latency between pipeline stages, especially stage 1 and stage 2

❖ Do

- Follow the instructions in the next slides

❖ Verify

- Check the logs printed in the terminal

❖ Time

- 15 minutes

Parameters to Adjust

- ❖ Number of layers in each submodel
 - Decide how many layers go into submodel 0 and submodel 1
- ❖ Accelerator assignment
 - Select which accelerator (CPU or GPU) runs stage 1 and stage 2
 - Revise configurations in

[~/DNNPipe-Tutorial/run_pipelined_inference_driver.sh](#)

```
19 submodel0="./models/submodel_0.tflite"
20 submodel0_gpu_usage="false" # adjust as needed
21 submodel1="./models/submodel_1.tflite"
22 submodel1_gpu_usage="true" # adjust as needed
```

Run Pipelined Inference Driver

- ❖ In `~/DNNPipe-Tutorial`, run
 - `make`
 - `./run_pipelined_inference_driver.sh`
- ❖ Expected output

```
[INFO] Top 3 predictions for image index 499:  
- Class 867 (trailer_truck): 0.531833  
- Class 675 (moving范): 0.459434  
- Class 569 (garbage_truck): 0.00453667  
[INFO] Average Stage0 latency (500 runs): 2.136 ms  
[INFO] Average Stage1 latency (500 runs): 46.2 ms  
[INFO] Average Stage2 latency (500 runs): 40.594 ms  
[INFO] Average Stage3 latency (500 runs): 0 ms  
[INFO] Throughput: 21.4096 items/sec (500 items in 23354 ms)
```